



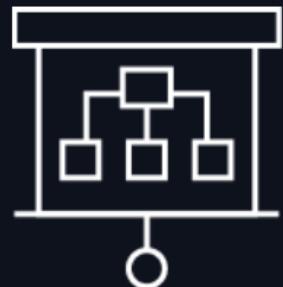
Istio in Practice - Day 1



Introductions

Each person say:

- Hello & where they are from
- How much k8s/Istio experience they have
- The main thing they want to learn
- Favourite something computer related
- Favourite something not computer related



Training

Learn and engage directly alongside our team, with courses for all stages of your Kubernetes journey



Workshops

Enhancing Kubernetes knowledge with our workshops. Whichever stage of your cloud native journey you are on we can help you do it better



Consulting

Consulting and engineering to make the most of Kubernetes and move you to production quickly



Subscription

Reference architecture, online training and SLA support 24x7 for your production Kubernetes deployment



Agenda

- Introduction
- Traffic Management
- Security
- Observability



What is a Service Mesh?

An architectural pattern that provides common network services as a feature of the infrastructure



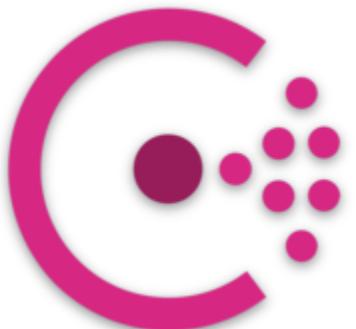
Why Do We Need a Service Mesh?

The scale and complexity of microservice architectures and distributed systems benefits greatly from a dedicated layer to orchestrate cross cutting concerns

- Facilitates service-to-service communication between microservices
- Abstracts the network and allows operators to think in terms of services rather than specific IPs and ports
- Functionality includes discovery, load balancing, failure recovery, policy enforcement, monitoring and telemetry

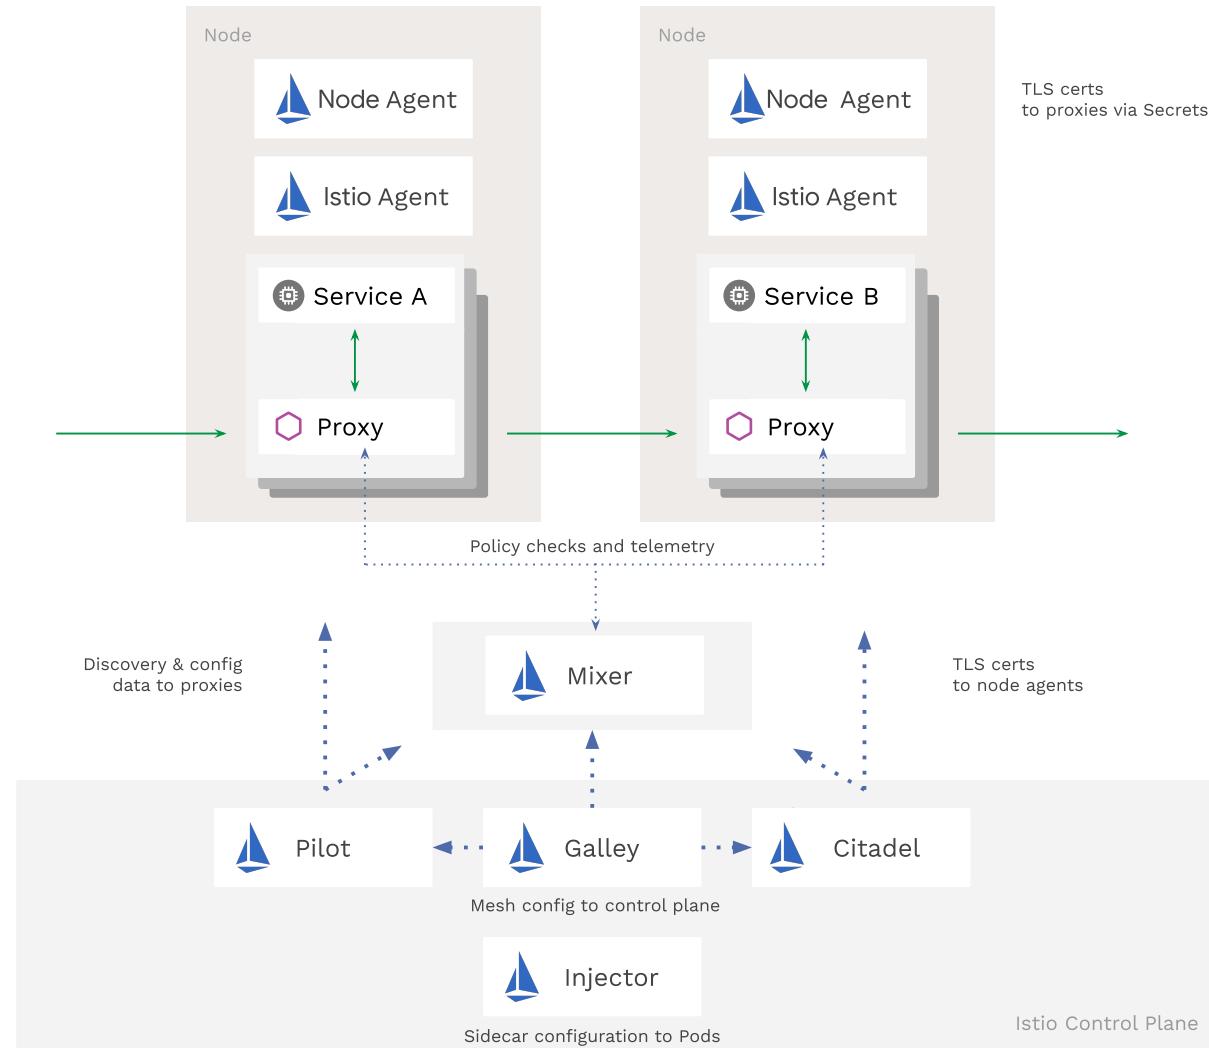


Service Mesh Implementations



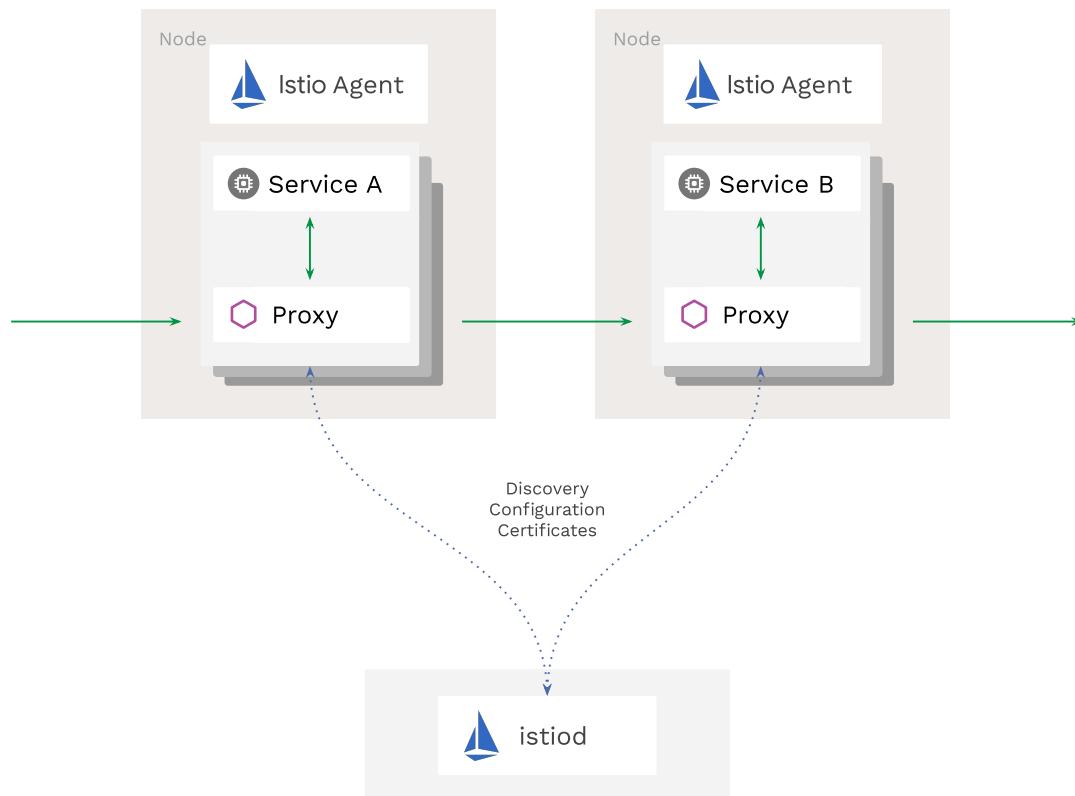


Istio Architecture v1.4.X



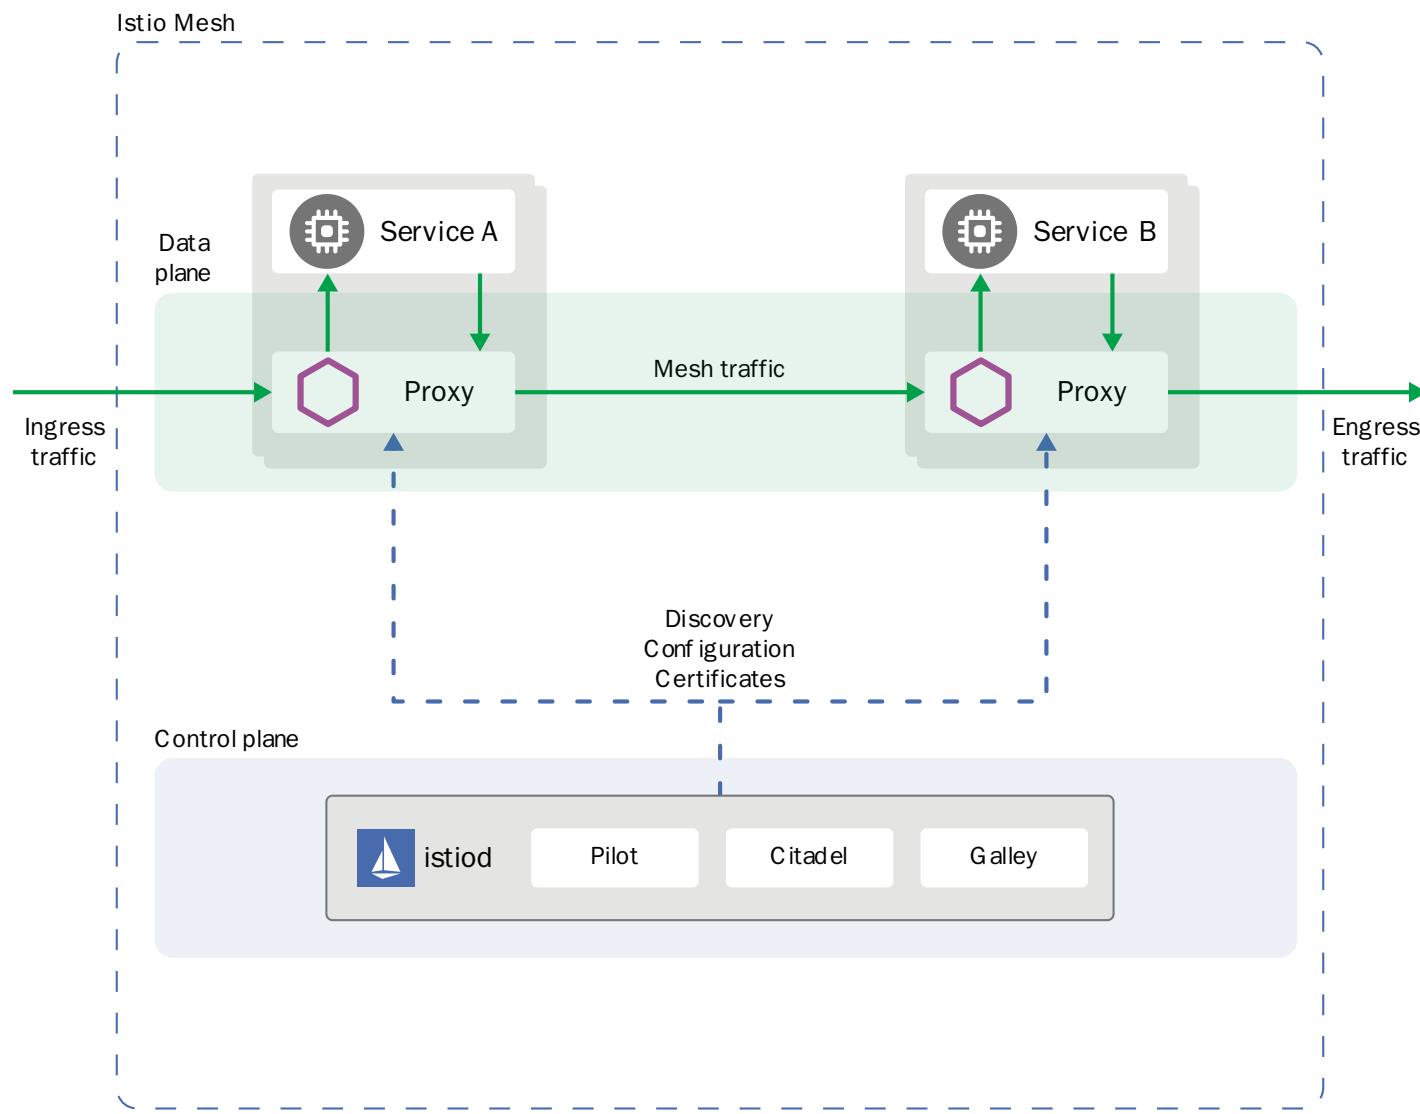


Istio Architecture v1.5.X+





Istio Architecture v1.5.X+



<https://istio.io/docs/ops/deployment/architecture/arch.svg>



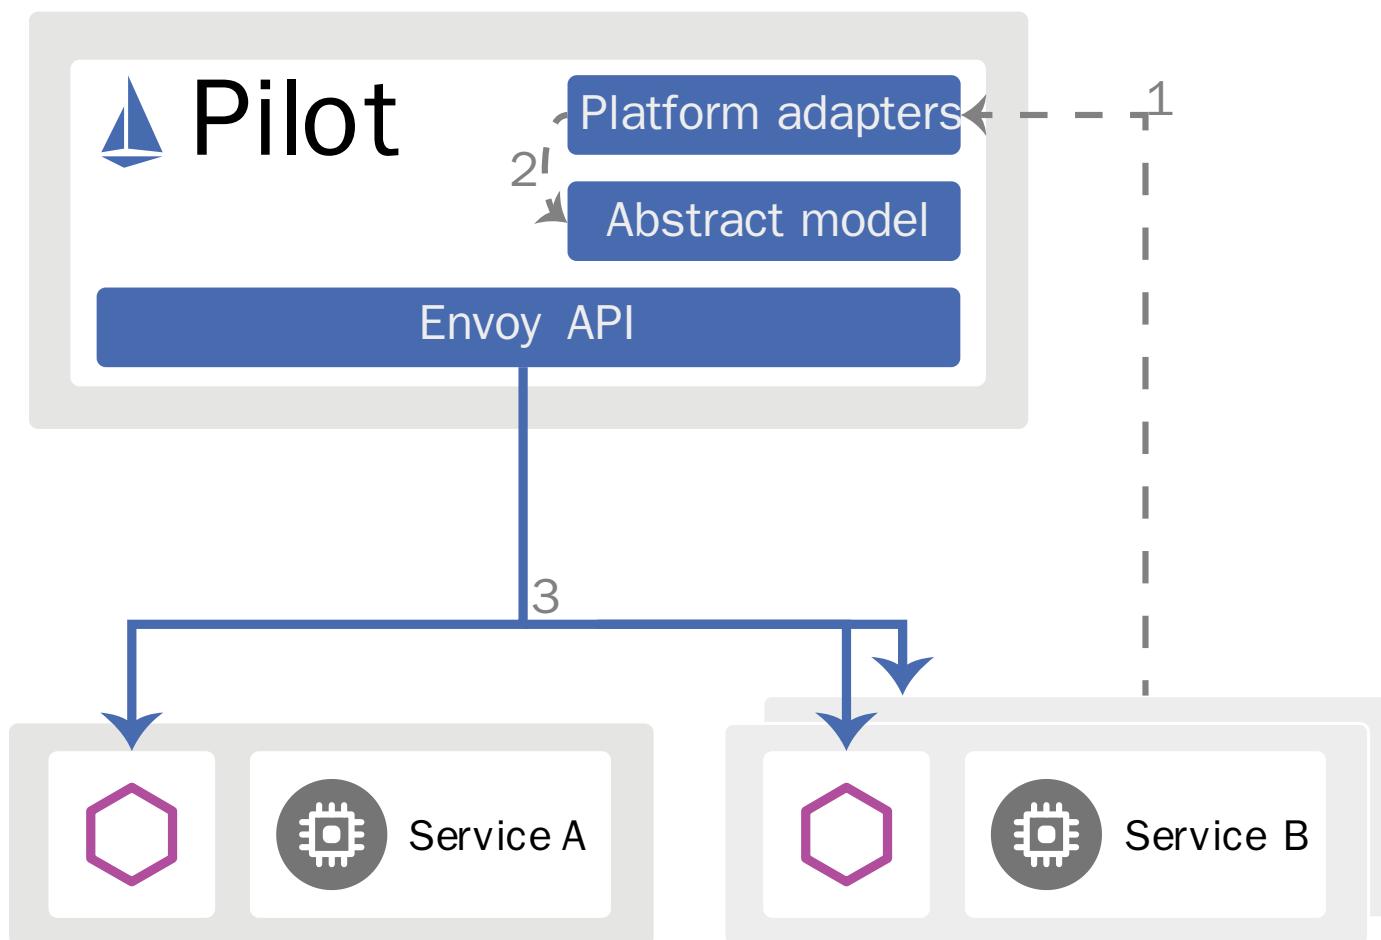
Control Plane / istiod

- Pilot
- Citadel
- Galley
- Sidecar Injection



Pilot

Serves xDS APIs to configure service discovery, traffic management capabilities and security for the Envoy sidecars

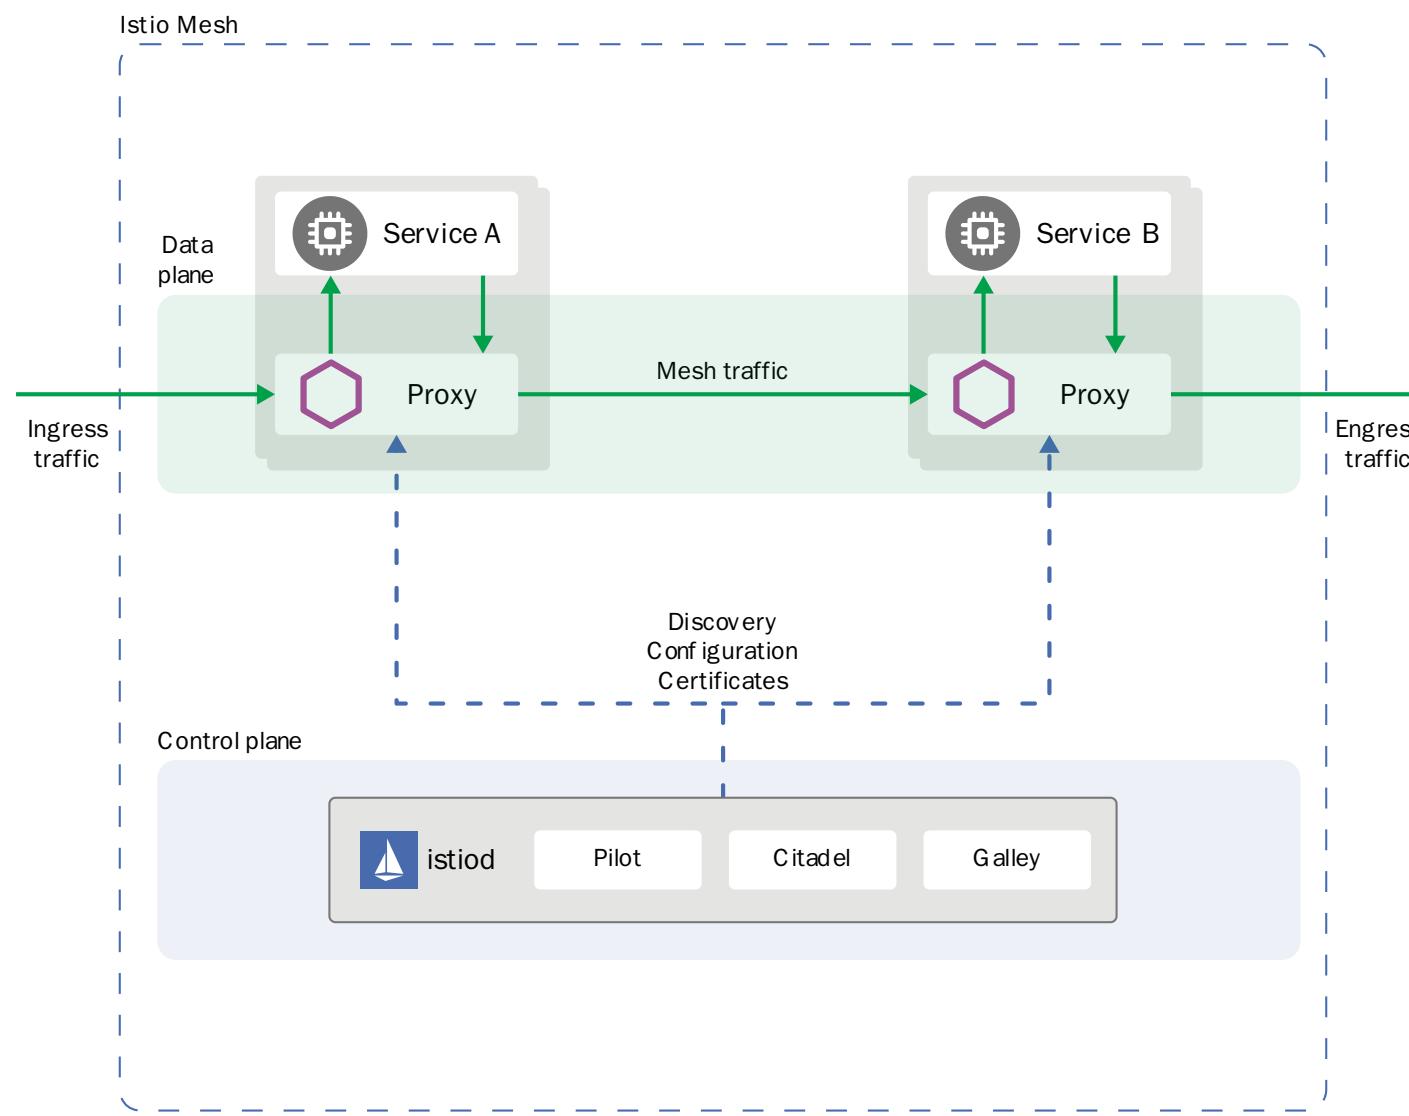


<https://istio.io/latest/docs/ops/deployment/architecture/discovery.svg>



Citadel

Enables strong service-to-service and end-user authentication with built-in identity and credential management

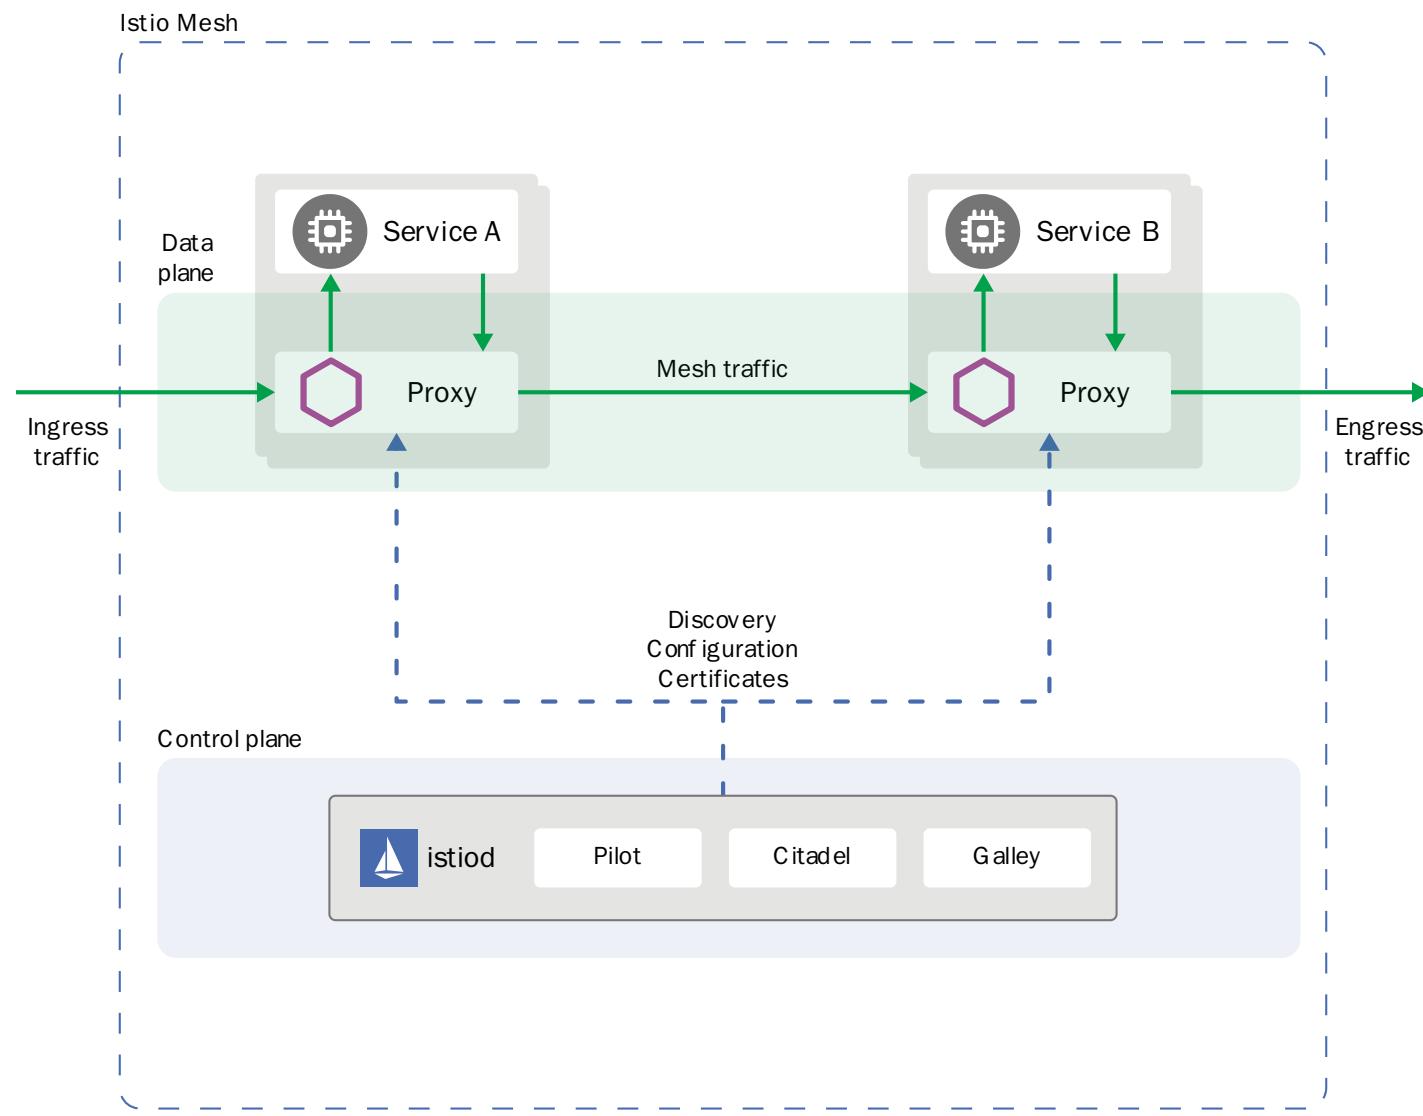


<https://istio.io/latest/docs/ops/deployment/architecture/arch.svg>



Galley

Converts platform specific CRDs into component configuration which can be consumed (e.g. by Pilot) over the Mesh Configuration Protocol



<https://istio.io/latest/docs/ops/deployment/architecture/arch.svg>



Sidecar Injection

Manual

```
istioctl kube-inject -f deployment.yaml | \  
    kubectl apply -f -
```

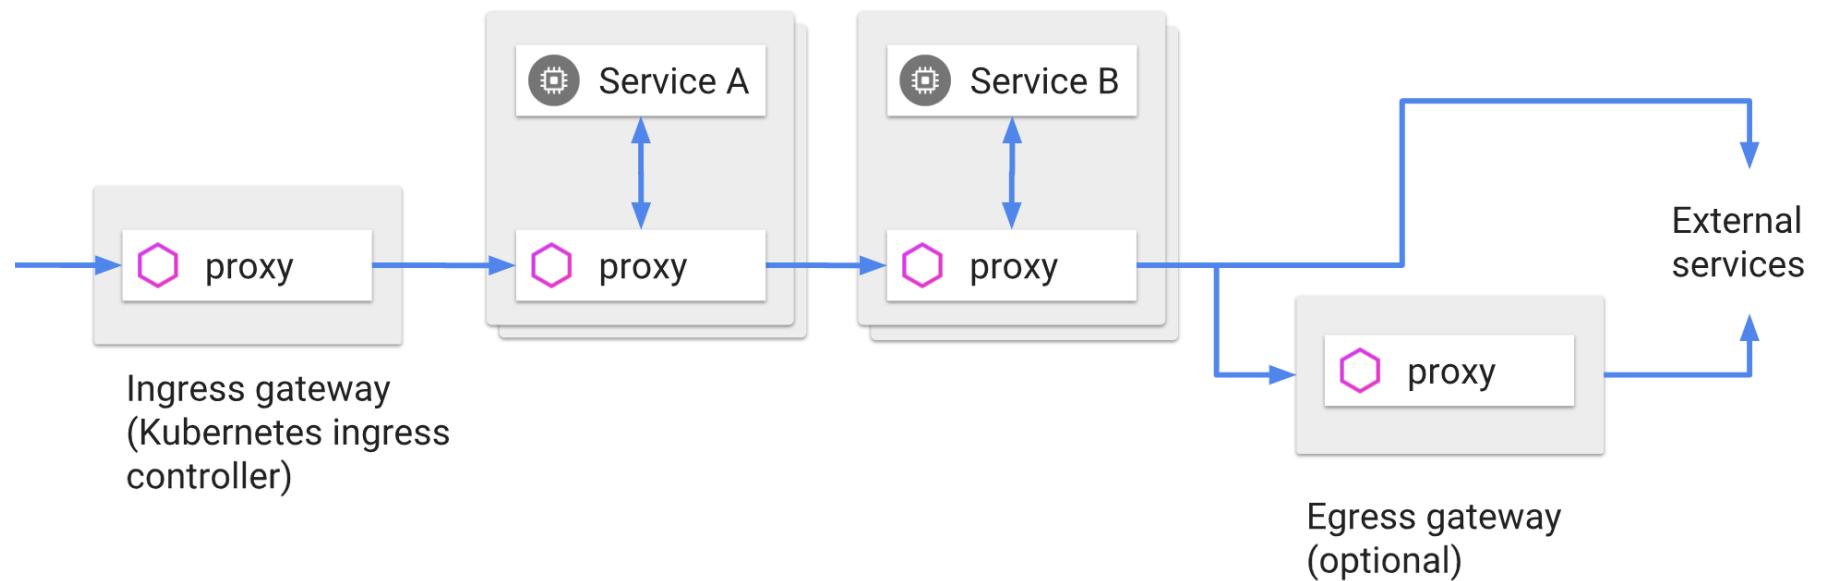
Automatic

```
apiVersion: admissionregistration.k8s.io/v1  
kind: MutatingWebhookConfiguration  
metadata:  
  name: istio-sidecar-injector  
webhooks:  
- rules:  
  - apiGroups:  
    - ""  
    apiVersions:  
    - v1  
    operations:  
    - CREATE  
    resources:  
    - pods  
    scope: '*'  
  ...
```



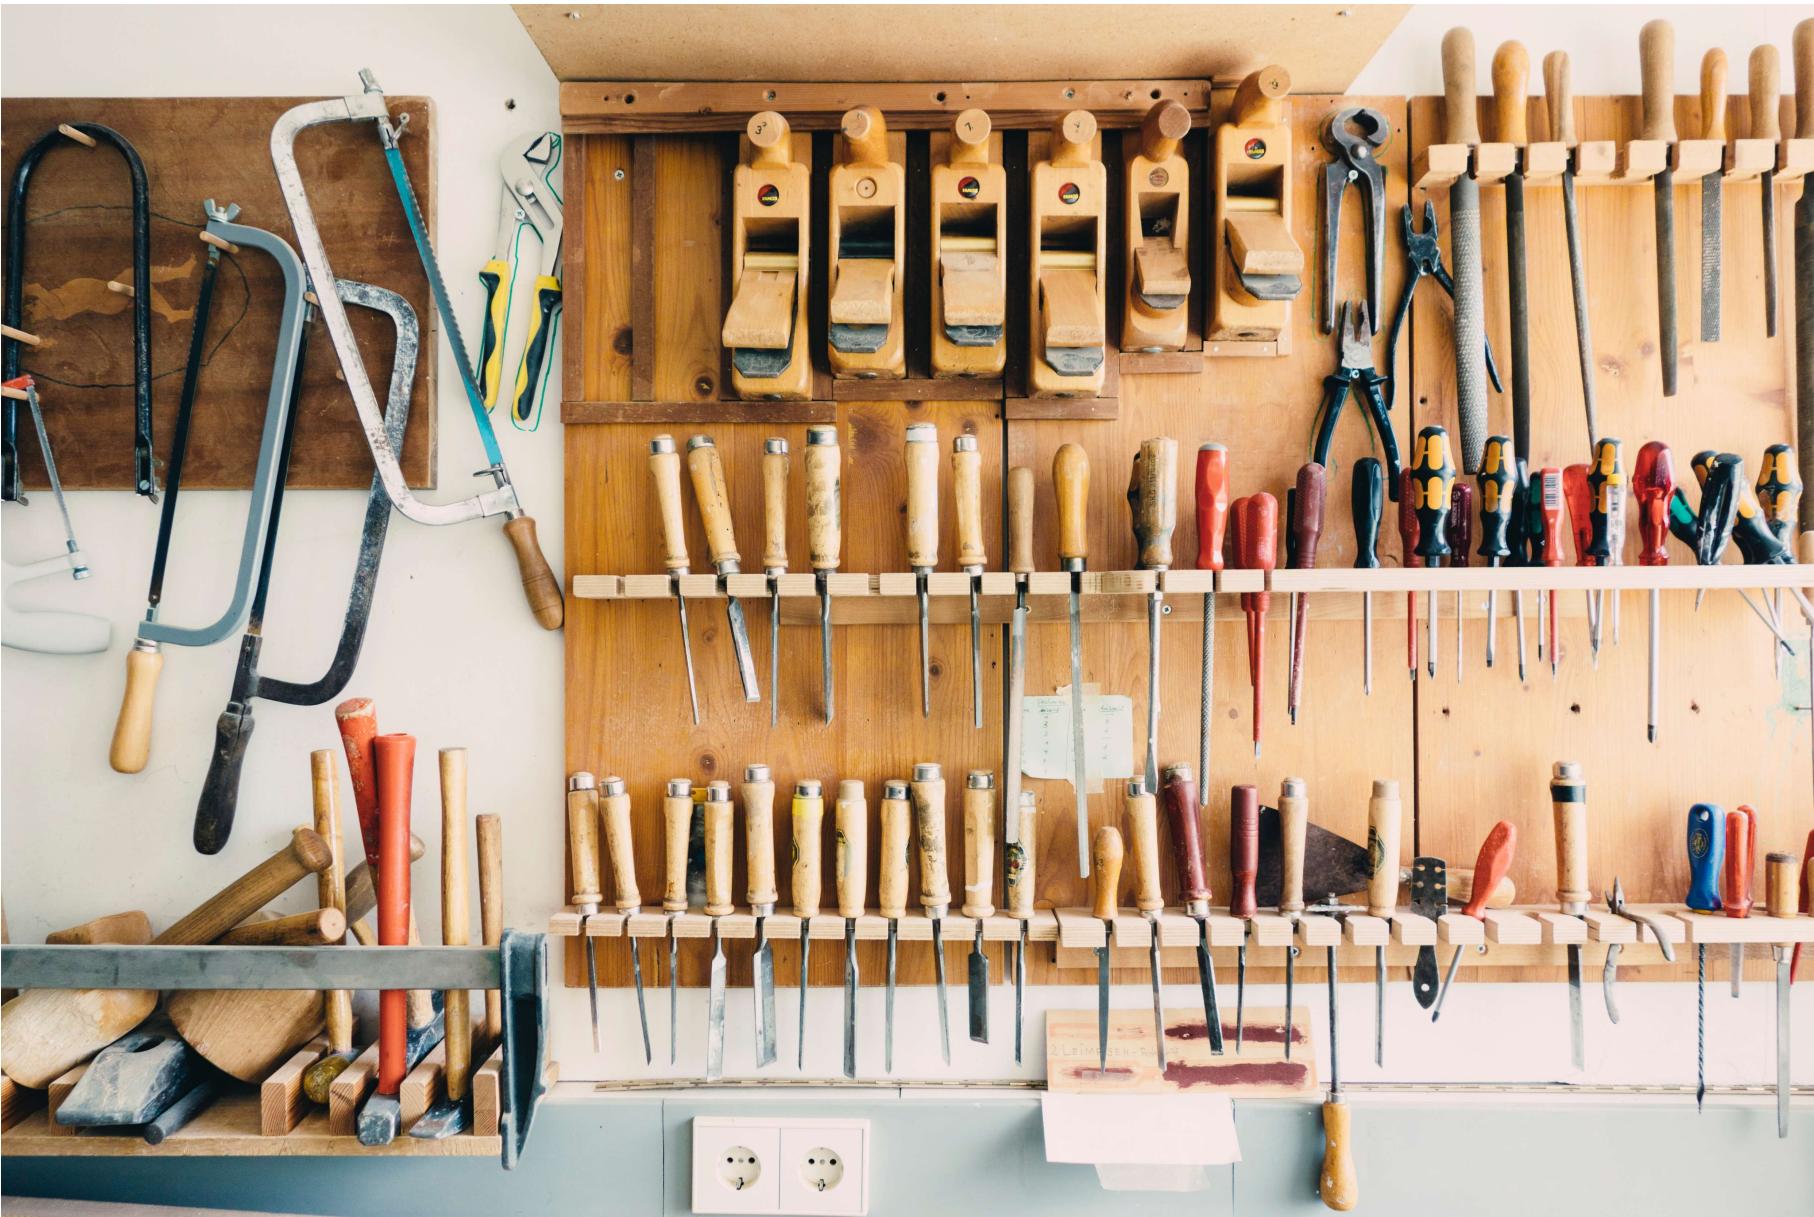
Data Plane

Envoy





Workshop: Deploy GKE Cluster and Install Istio





IstioOperator API

```
apiVersion: install.istio.io/v1alpha1
kind: IstioOperator
spec:
  profile: default
```

```
istioctl install -f istio-operator.yaml
```



Traffic Management

Istio's traffic routing rules are used to control the flow of traffic and API calls between services.

Istio simplifies configuration of service-level properties like circuit breakers, timeouts, and retries, and makes it easy to set up important tasks like A/B testing, canary rollouts, and staged rollouts with percentage-based traffic splits.

`VirtualServices` and `DestinationRules` are key resources for configuring Istio's traffic routing functionality.



VirtualService

A VirtualService is used to configure how requests are routed to a service.

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
  - reviews
  http:
  - match:
    - headers:
        end-user:
          exact: jason
  route:
  - destination:
      host: reviews
      subset: v2
  - route:
      destination:
        host: reviews
        subset: v3
```



The hosts field

The `hosts` field lists the destination hosts to which traffic routing rules are applied.

```
hosts:  
- reviews # short name expanded to FQDN
```



Routing rules

The `http` section contains an ordered list of route rules for HTTP traffic.

- match conditions and actions for routing HTTP/1.1, HTTP2, and gRPC traffic sent to the destination(s) specified in the `hosts` field.
- A routing rule consists of the destination of the traffic and zero or more match conditions, depending on the use case.



Match condition

In this case, the routing should apply to all requests from the user `jason`, so the `end-user` header, and `exact` field are used to select the appropriate requests.

```
- match:  
  - headers:  
    end-user:  
      exact: jason
```



Destination

The `route` section's `destination` field specifies the actual destination for traffic that matches this condition.

```
route:  
- destination:  
  host: reviews  
  subset: v2
```



DestinationRule

VirtualServices define how traffic is routed to a given destination.

DestinationRules define policies that apply to traffic intended for a service after routing has occurred.

A DestinationRule is applied after VirtualService routing rules are evaluated, so they apply to the traffic's "real" destination.



my-destination-rule.yaml

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: my-destination-rule
spec:
  host: my-svc
  trafficPolicy:
    loadBalancer:
      simple: RANDOM
  subsets:
  - name: v1
    labels:
      version: v1
  - name: v2
    labels:
      version: v2
    trafficPolicy:
      loadBalancer:
        simple: ROUND_ROBIN
  - name: v3
    labels:
      version: v3
```



Load balancing options

By default, Istio uses a round-robin load balancing policy but can also apply the following models:

- Random: Requests are forwarded at random to instances in the pool.
- Weighted: Requests are forwarded to instances in the pool according to a specific percentage.
- Least requests: Requests are forwarded to instances with the least number of requests.



Workshop - VirtualServices and DestinationRules





ServiceEntry



ServiceEntry

In order to direct traffic within your mesh, Istio needs to know where all your endpoints are and which services they belong to.



ServiceEntry

A ServiceEntry is used to add an entry into Istio's internal service registry.

After a ServiceEntry is added, the Envoy proxies can send traffic to the specified host as if it was a service in the mesh.



ServiceEntry

Configuring service entries allows traffic to be managed for services running outside of the mesh, including the following tasks:

- Redirect and forward traffic for external destinations, such as APIs consumed from the web, or traffic to services in legacy infrastructure.
- Define retry, timeout, and fault injection policies for external destinations.
- Logically add services from a different cluster to the mesh to configure a multicloud Istio mesh on Kubernetes.



The following example MESH_EXTERNAL service entry adds the ext-svc external dependency to Istio's service registry:

svc-entry.yaml

```
apiVersion: networking.istio.io/v1alpha3
kind: ServiceEntry
metadata:
  name: svc-entry
spec:
  hosts:
  - ext-svc.example.com      # FQDN of external resource or wildcard pref
  ports:
  - number: 443
    name: https
    protocol: HTTPS
  location: MESH_EXTERNAL
  resolution: DNS
```



Multi cluster services

```
apiVersion: networking.istio.io/v1alpha3
kind: ServiceEntry
metadata:
  name: httpbin-bar
spec:
  hosts:
    # must be of form name.namespace.global
    - httpbin.bar.global
    # Treat remote cluster services as part of the service mesh
    # as all clusters in the service mesh share the same root of trust.
    location: MESH_INTERNAL
  ports:
    - name: http1
      number: 8000
      protocol: http
  resolution: DNS
  addresses:
    # the IP address to which httpbin.bar.global will resolve to
    # must be unique for each remote service, within a given cluster.
    # This address need not be routable. Traffic for this IP will be
    # captured by the sidecar and routed appropriately.
    - 240.0.0.2
  endpoints:
    # This is the routable address of the ingress gateway in cluster2
    # that sits in front of sleep.foo service. Traffic from the sidecar
    # will be routed to this address.
    - address: ${CLUSTER2_GW_ADDR}
      ports:
        http1: 15443
```



Gateway

Istio can control inbound and outbound traffic to the mesh by defining ingress and egress Gateways.

Gateway configurations are applied to standalone Envoy proxies that are running at the edge of the mesh, rather than sidecar Envoy proxies running alongside the service workloads.



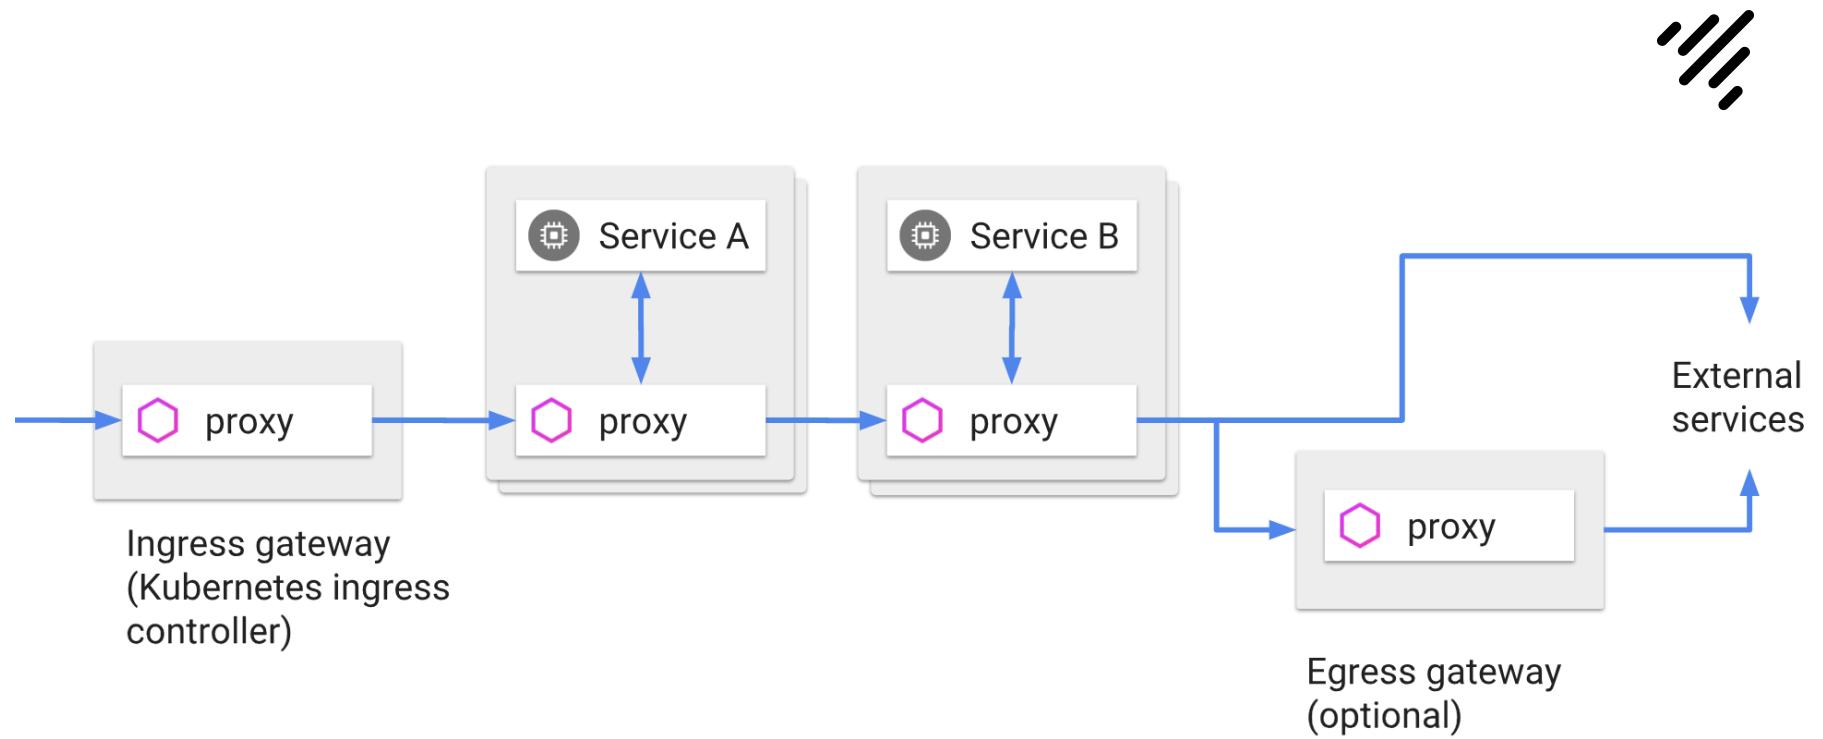
Ingress Gateway

An ingress Gateway describes a load balancer operating at the edge of the mesh that receives incoming HTTP/TCP connections.



Egress Gateway

An egress gateway configures a dedicated exit node for the traffic leaving the mesh, limiting which services can or should access external networks, or to enable secure control of egress traffic to add security to the mesh.



<https://blogs.sap.com/wp-content/uploads/2019/02/Screenshot-2019-02-08-at-16.02.10.png>



httpbin-gateway.yaml

```
apiVersion: networking.istio.io/v1alpha3
kind: Gateway
metadata:
  name: httpbin-gateway
spec:
  selector:
    istio: ingressgateway # use Istio default gateway implementation
  servers:
  - port:
      number: 80
      name: http
      protocol: HTTP
    hosts:
    - "httpbin.example.com"
```



Ingress and Egress Gateway installation

```
istioctl manifest install -f <(cat <<EOF
apiVersion: install.istio.io/v1alpha1
kind: IstioOperator
spec:
  profile: default
  components:
    ingressGateways:
    - enabled: true # default is true for default profile
    egressGateways:
    - enabled: true
EOF
)
```

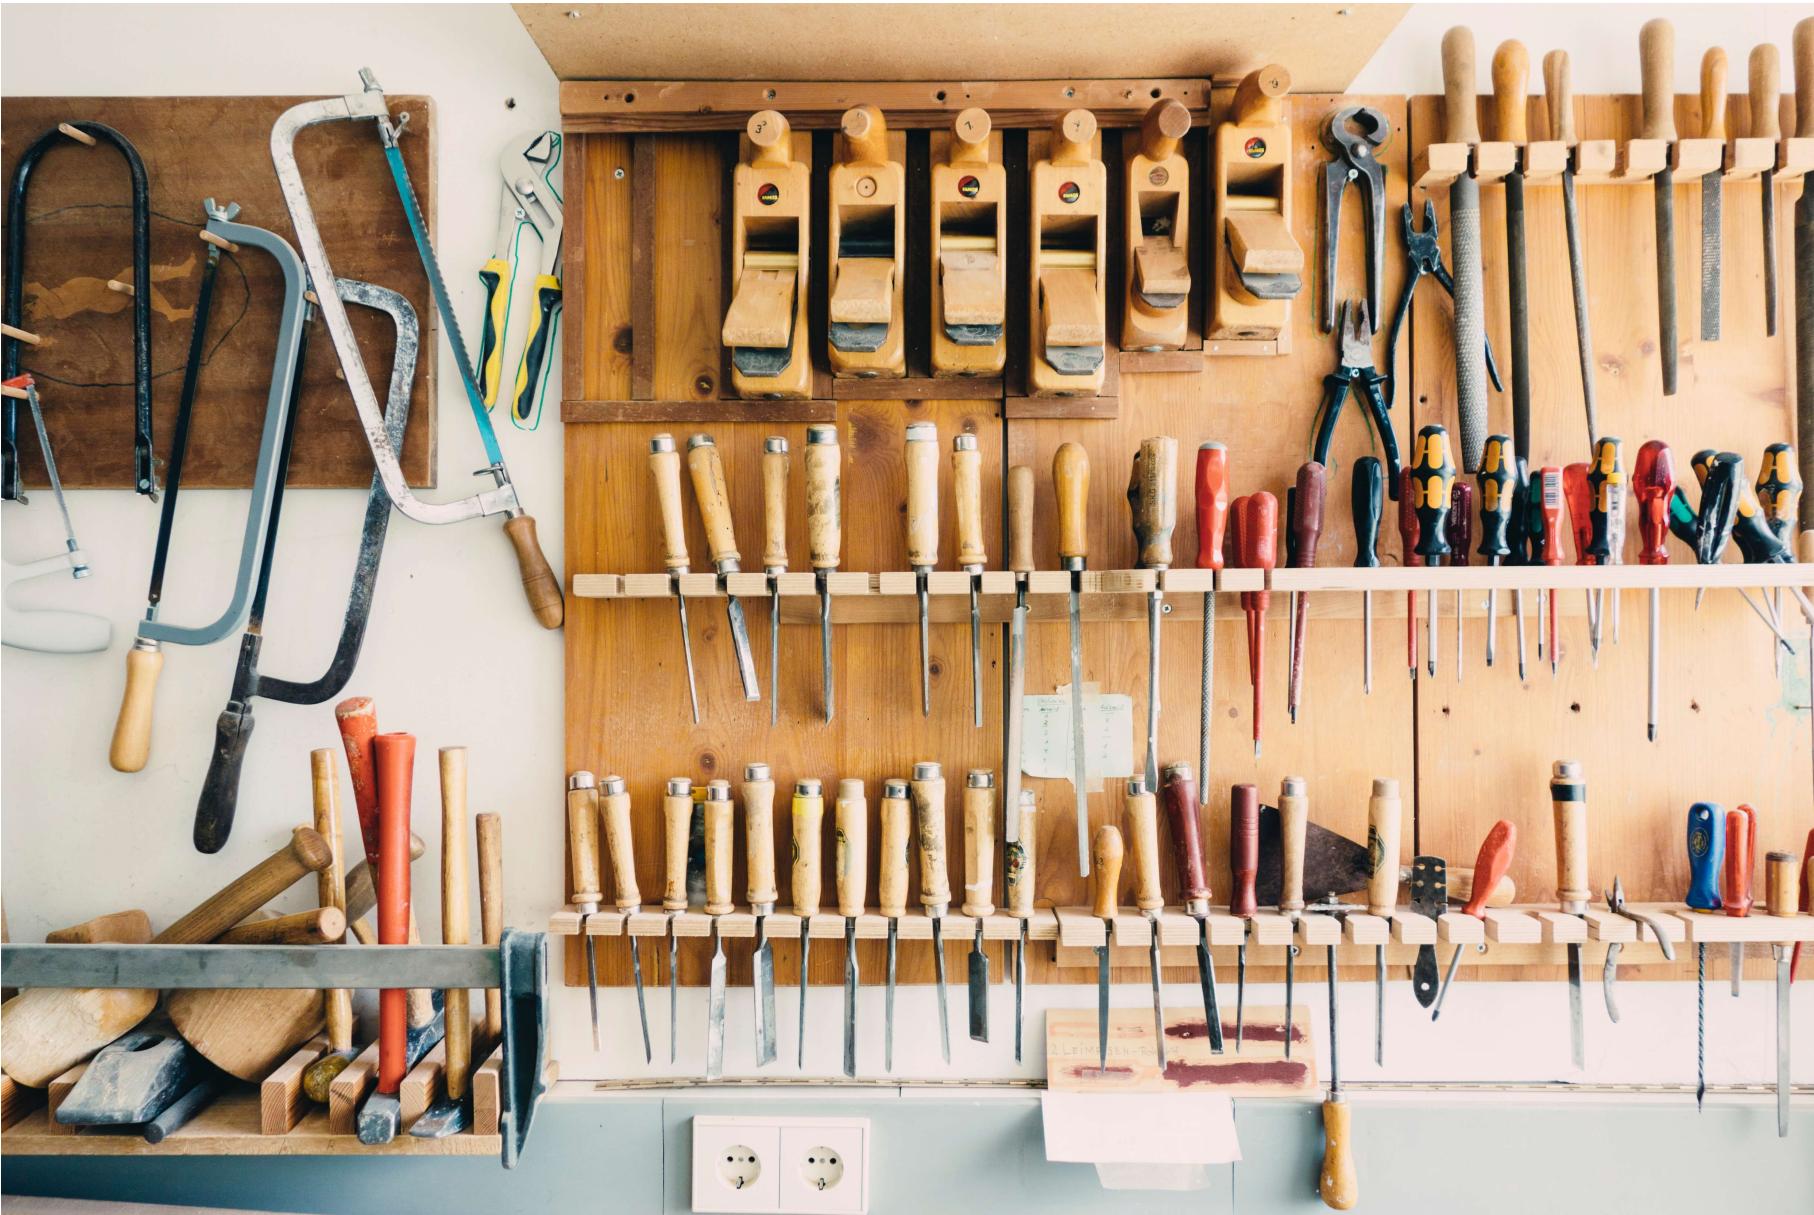


The Gateway must be bound to a VirtualService and specify routing as show below:

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: httpbin
spec:
  hosts:
    - "httpbin.example.com"
  gateways:
    - httpbin-gateway
  http:
    - match:
        - uri:
            prefix: /status
        - uri:
            prefix: /delay
    route:
      - destination:
          port:
            number: 8000
          host: httpbin
```



Workshop - ServiceEntry and Gateway





Security



Introduction

Istio implements a comprehensive array of security features to address common concerns:

- To defend against man-in-the-middle attacks, we need traffic encryption.
- To provide flexible service access control, we need mutual TLS and fine-grained access policies.
- To audit who did what at what time, we need auditing tools.

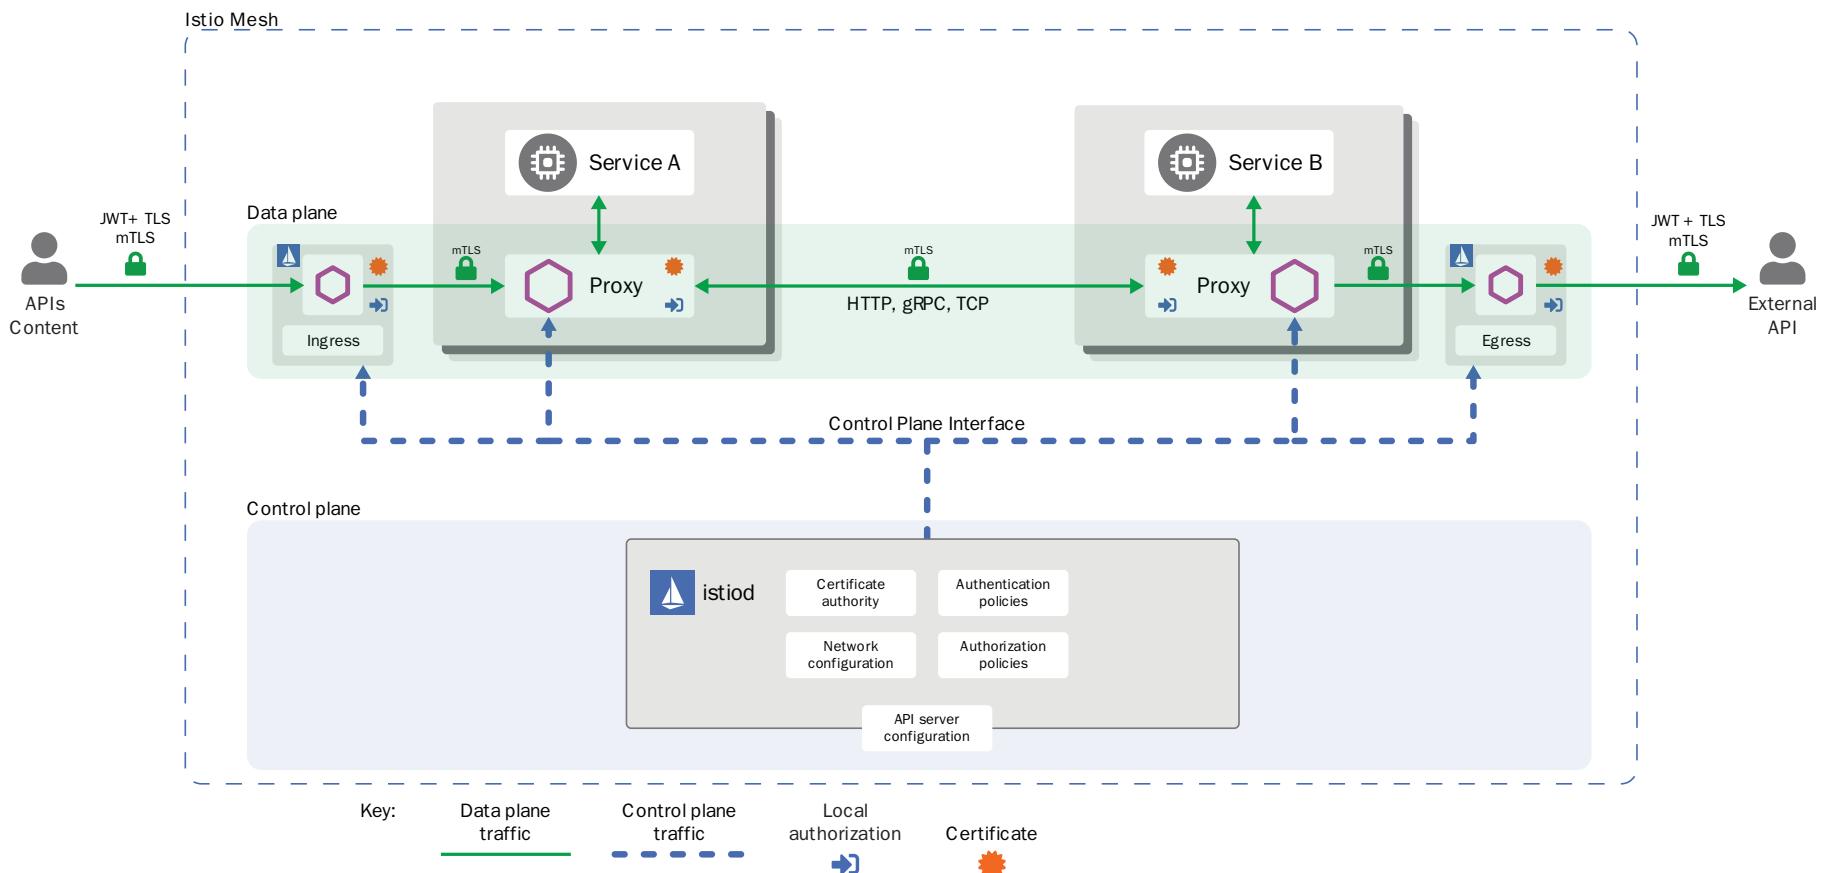


The goals of Istio security are:

- Security by default: no changes needed for application code and infrastructure
- Defence in depth: integrate with existing security systems to provide multiple layers of defence
- Zero-trust network: build security solutions on untrusted networks



High-level architecture





Istio identity

Identity is a fundamental concept of any security infrastructure.

In the Istio identity model, Istio uses the first-class service identity to determine the identity of a service.

This gives great flexibility and granularity to represent a human user, an individual service, or a group of services.

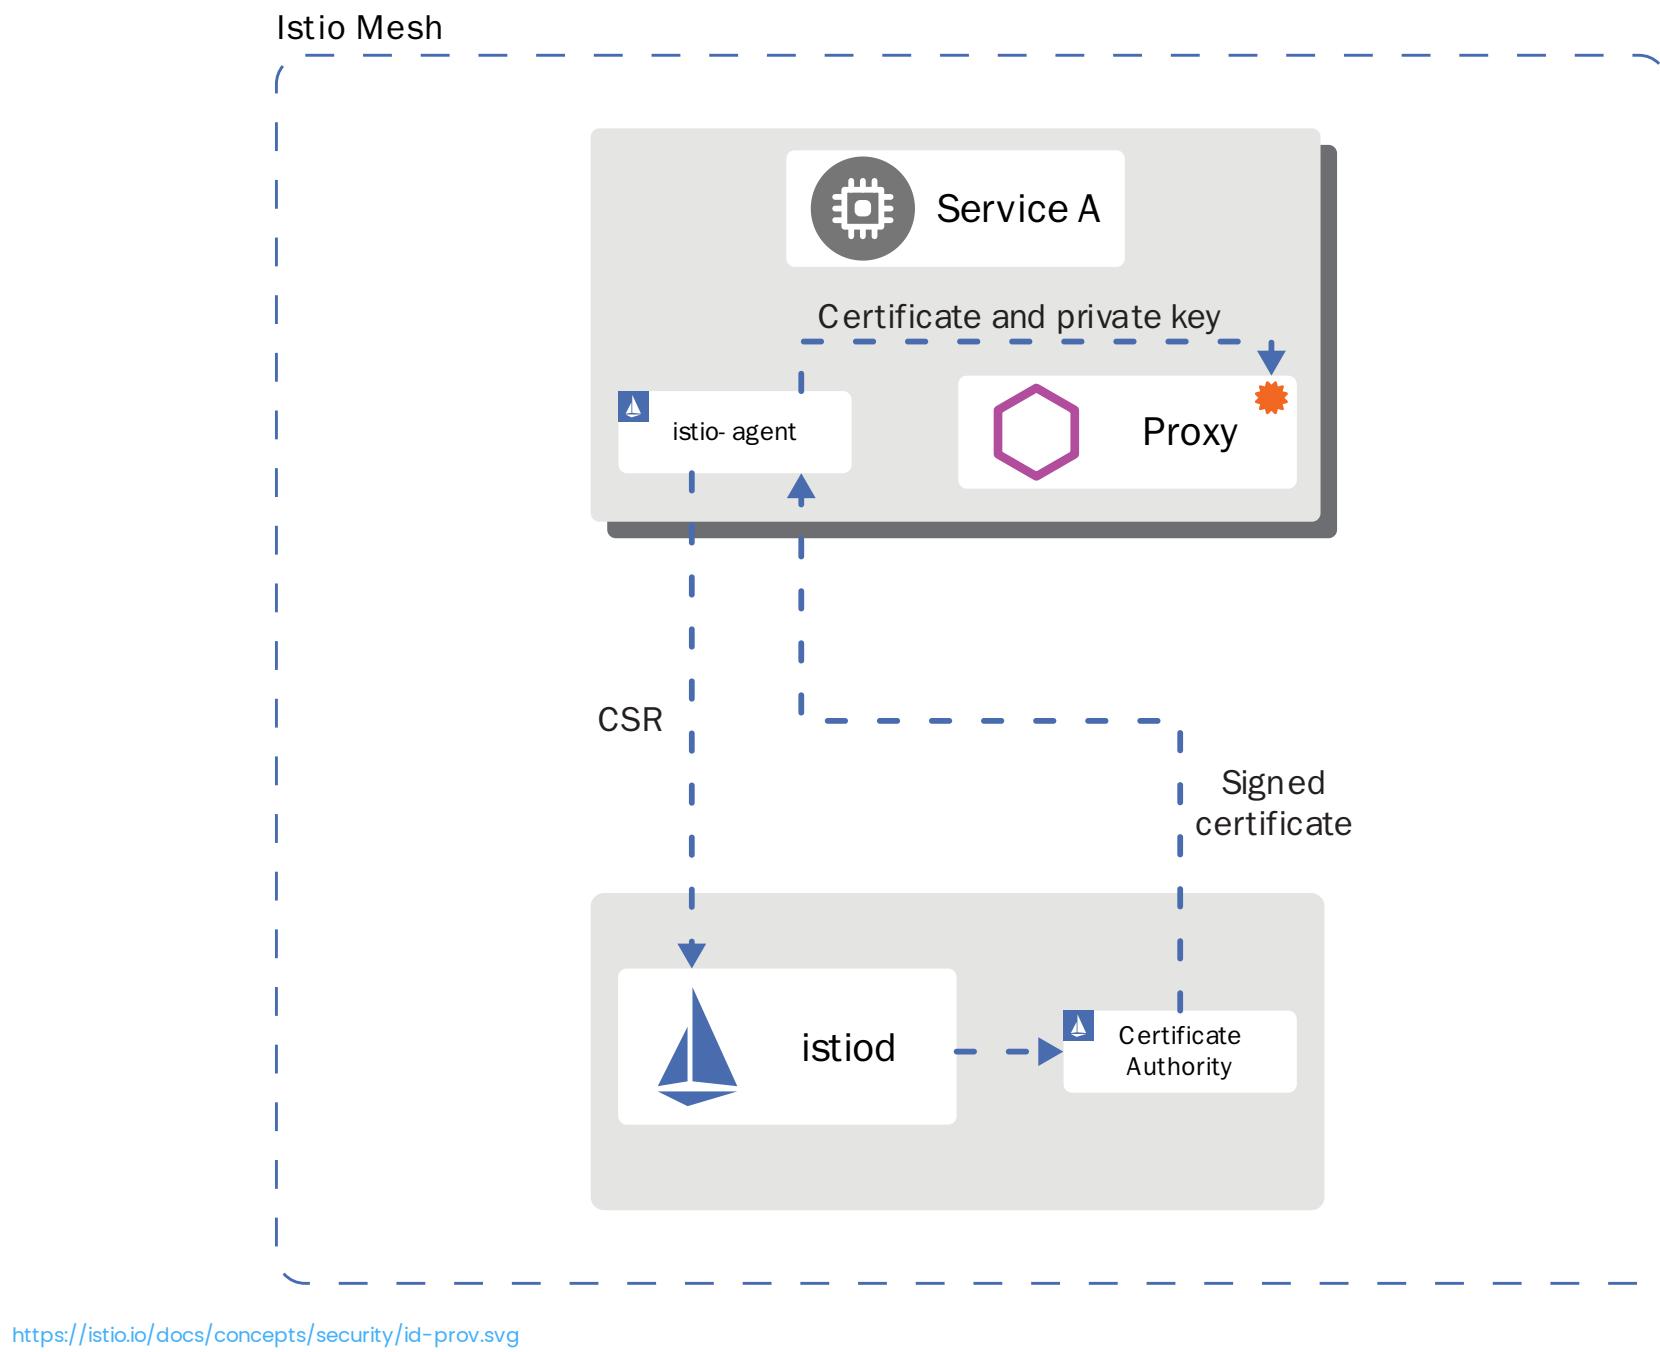
On platforms that do not have such identity available, Istio can use other identities that can group service instances, such as service names.



PKI

The Istio PKI securely provisions strong identities to every workload with X.509 certificates.

To automate key and certificate rotation at scale, the PKI runs an Istio agent alongside each Envoy proxy for certificate and key provisioning.





Mutual TLS

Istio tunnels service-to-service communication through the client side and server side Envoy proxies.

For a client to call a server with mutual TLS authentication:

- Istio re-routes the outbound traffic from a client to the client's local sidecar Envoy.
- The client side Envoy starts a mutual TLS handshake with the server side Envoy.
- The client side Envoy and the server side Envoy establish a mutual TLS connection, and Istio forwards the traffic from the client side Envoy to the server side Envoy.
- After authorisation, the server side Envoy forwards the traffic to the server service through local TCP connections.



Permissive mode

Istio mutual TLS has a permissive mode, which allows a service to accept both plaintext traffic and mutual TLS traffic at the same time.

The server's installed Istio sidecar takes mutual TLS traffic immediately without breaking existing plaintext traffic.



Secure naming

In Istio, each workload is automatically assigned with an identity.

Server identities are encoded in certificates, but service names are retrieved through the discovery service or DNS.

The secure naming information maps the server identities to the service names.

A mapping of identity A to service name B means “ A is authorised to run service B ”.

The control plane watches the apiserver, generates the secure naming mappings, and distributes them securely to the Envoy proxies.

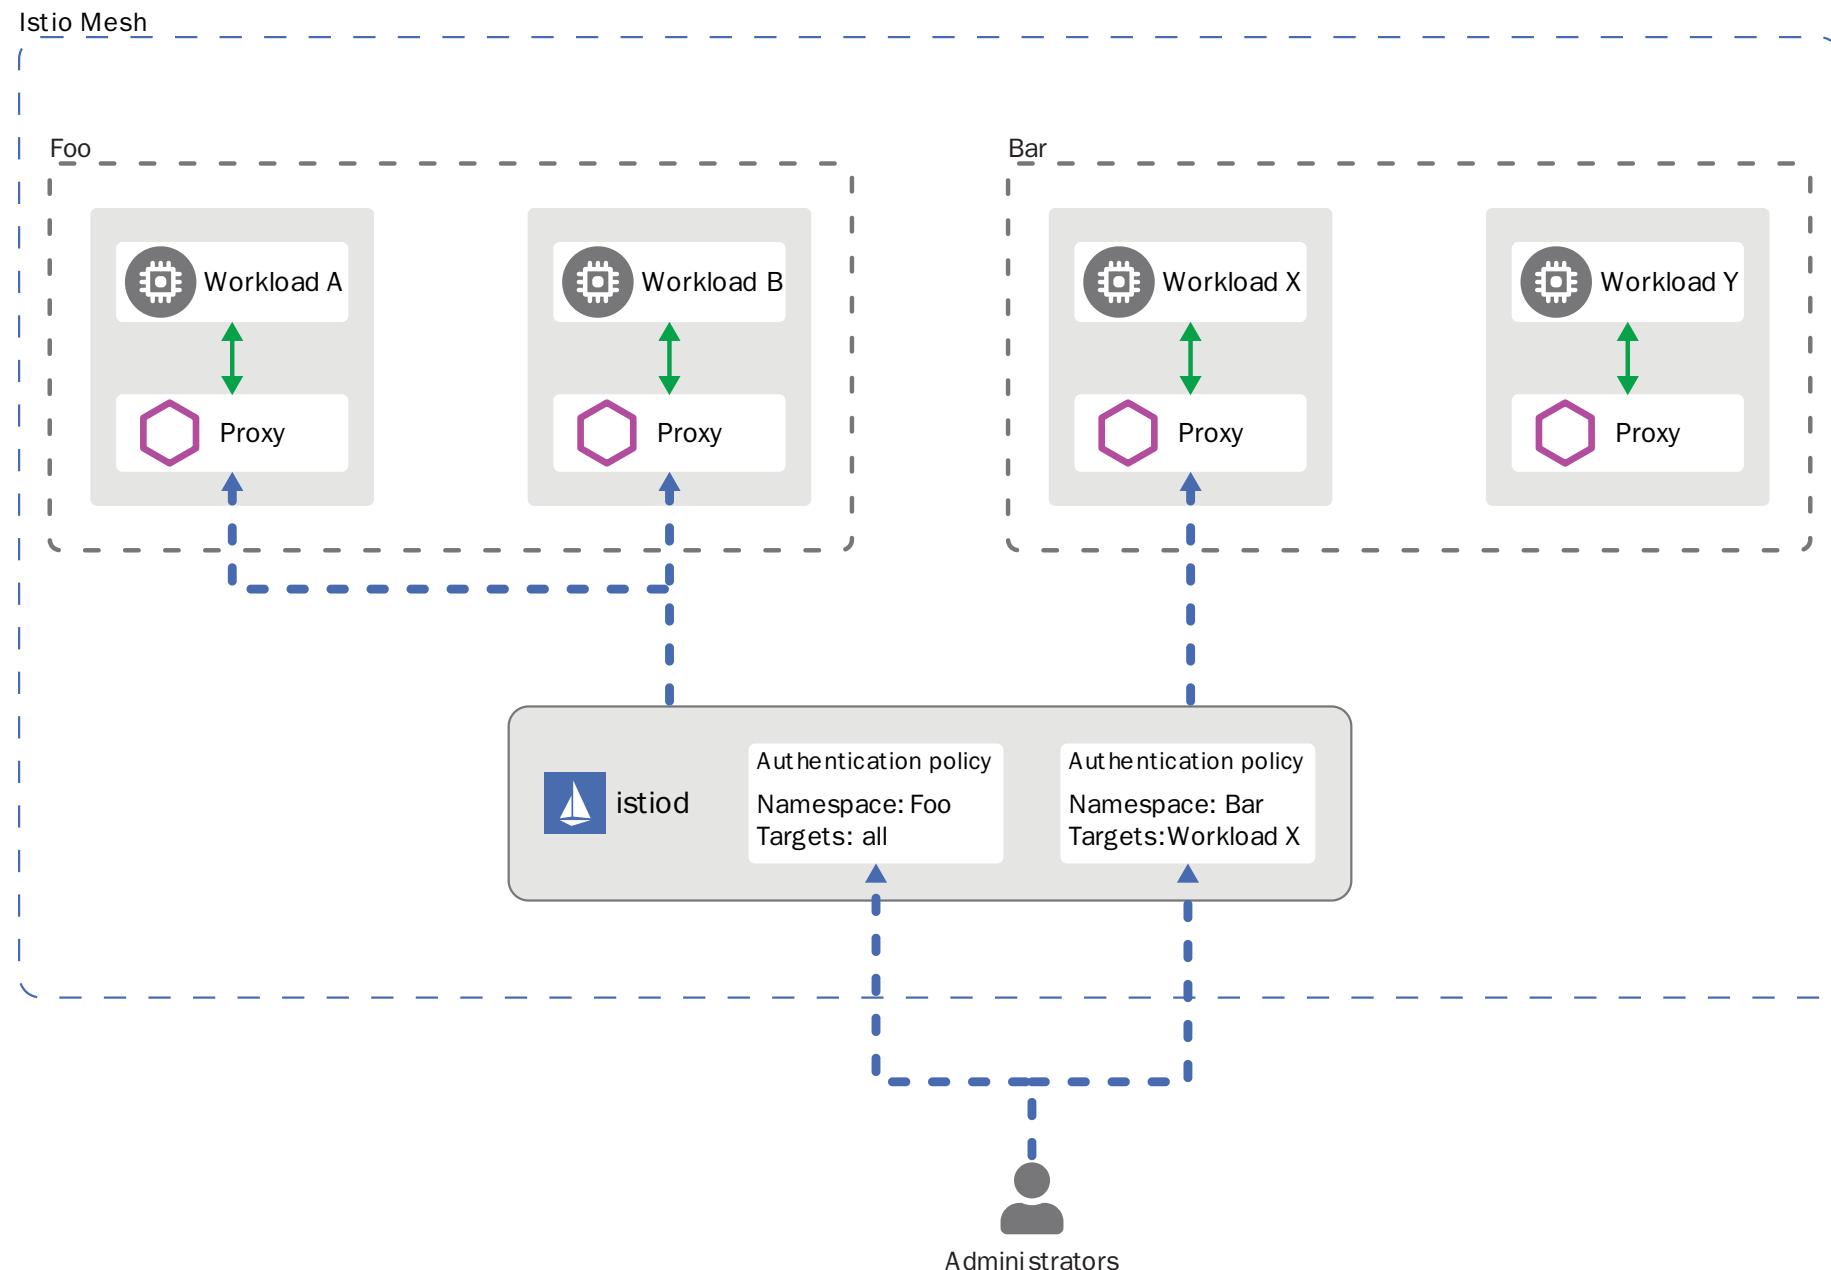


Secure naming offers the following benefits:

- Decouples the identity system from the service naming system.
- The certificate offers the identity of the server, instead of the service name the server is running for.
- It works well with the Kubernetes resource model.



Authentication architecture





Authentication

Istio provides two types of authentication:

- Peer authentication
- Request authentication



PeerAuthentication and RequestAuthentication policies use selector fields to specify the label of the workloads to which the policy applies.

The following snippet shows the selector field of a policy that applies to workloads with the app:product-page label:

```
selector:  
  matchLabels:  
    app: product-page
```

If you don't provide a value for the selector field, Istio matches the policy to all workloads in the storage scope of the policy.

Therefore, the selector fields help you specify the scope of the policies:

- Mesh-wide policy: A policy specified for the root namespace without or with an empty selector field.
- Namespace-wide policy: A policy specified for a non-root namespace without or with an empty selector field.
- Workload-specific policy: a policy defined in the regular namespace, with non-empty selector field.



Peer authentication

Used for service-to-service authentication to verify the client making the connection.

Istio offers mutual TLS as a full stack solution for transport authentication.

Authentication policies apply to requests that a service receives.

The following peer authentication policy specifies that transport authentication for the `reviews` service must use mutual TLS:

```
apiVersion: "security.istio.io/v1beta1"
kind: "PeerAuthentication"
metadata:
  name: "example-peer-policy"
  namespace: "foo"
spec:
  selector:
    matchLabels:
      app: reviews
  mtls:
    mode: STRICT
```



Peer authentication

Peer authentication policies specify the mutual TLS mode Istio enforces on target workloads. The following modes are supported:

- **PERMISSIVE**: Workloads accept both mutual TLS and plain text traffic. This mode is most useful during migrations when workloads without sidecar cannot use mutual TLS. Once workloads are migrated with sidecar injection, you should switch the mode to STRICT.
- **STRICT**: Workloads only accept mutual TLS traffic.
- **DISABLE**: Mutual TLS is disabled. From a security perspective, you shouldn't use this mode unless you provide your own security solution.



Peer authentication

```
apiVersion: "security.istio.io/v1beta1"
kind: "PeerAuthentication"
metadata:
  name: "example-workload-policy"
  namespace: "foo"
spec:
  selector:
    matchLabels:
      app: example-app
  portLevelMtls:
    80:
      mode: DISABLE
```



Request authentication

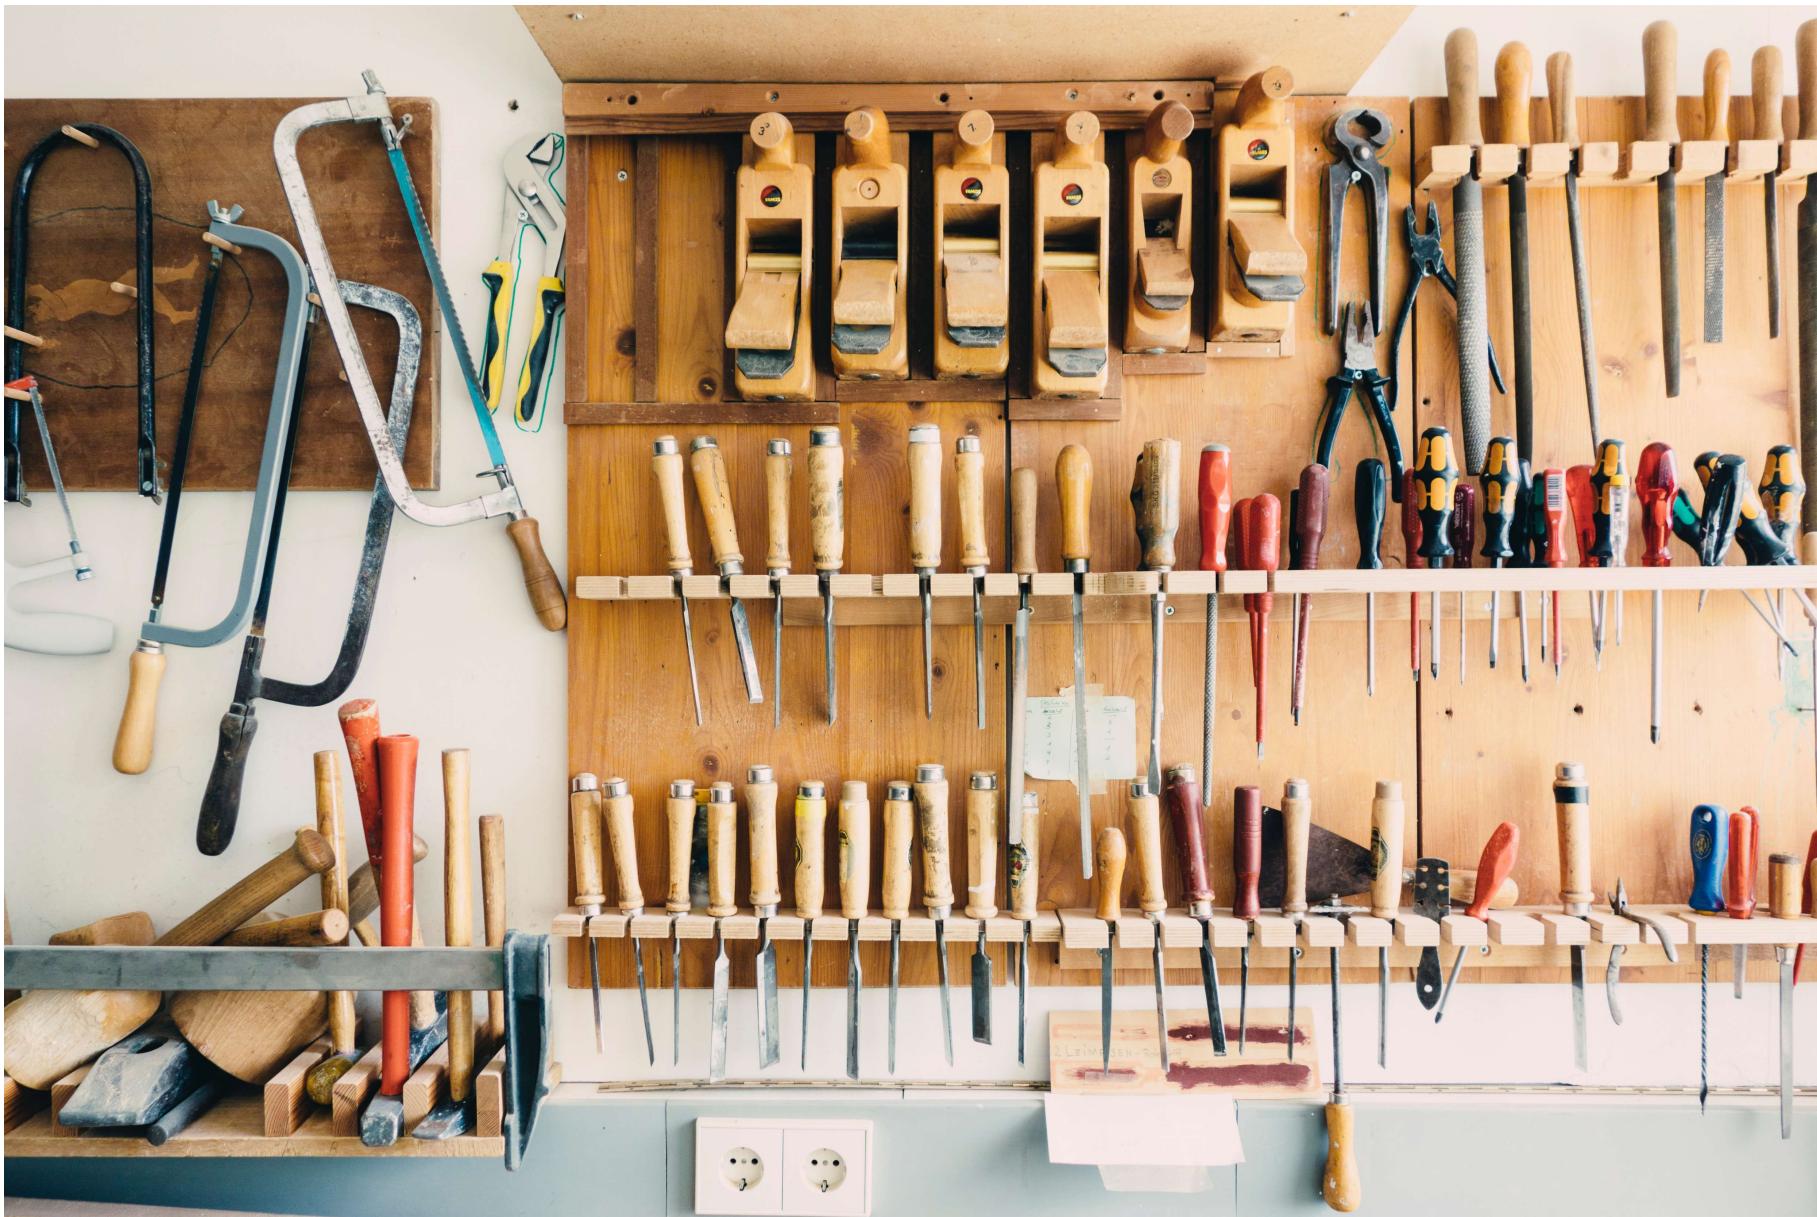
Used for end-user authentication to verify the credential attached to the request.

The following request authentication policy requires an end-user JWT for the ingress gateway:

```
apiVersion: "security.istio.io/v1beta1"
kind: "RequestAuthentication"
metadata:
  name: "jwt-example"
  namespace: istio-system
spec:
  selector:
    matchLabels:
      istio: ingressgateway
  jwtRules:
  - issuer: "testing@secure.istio.io"
    jwksUri: "https://raw.githubusercontent.com/istio/istio/release-1.5/
```

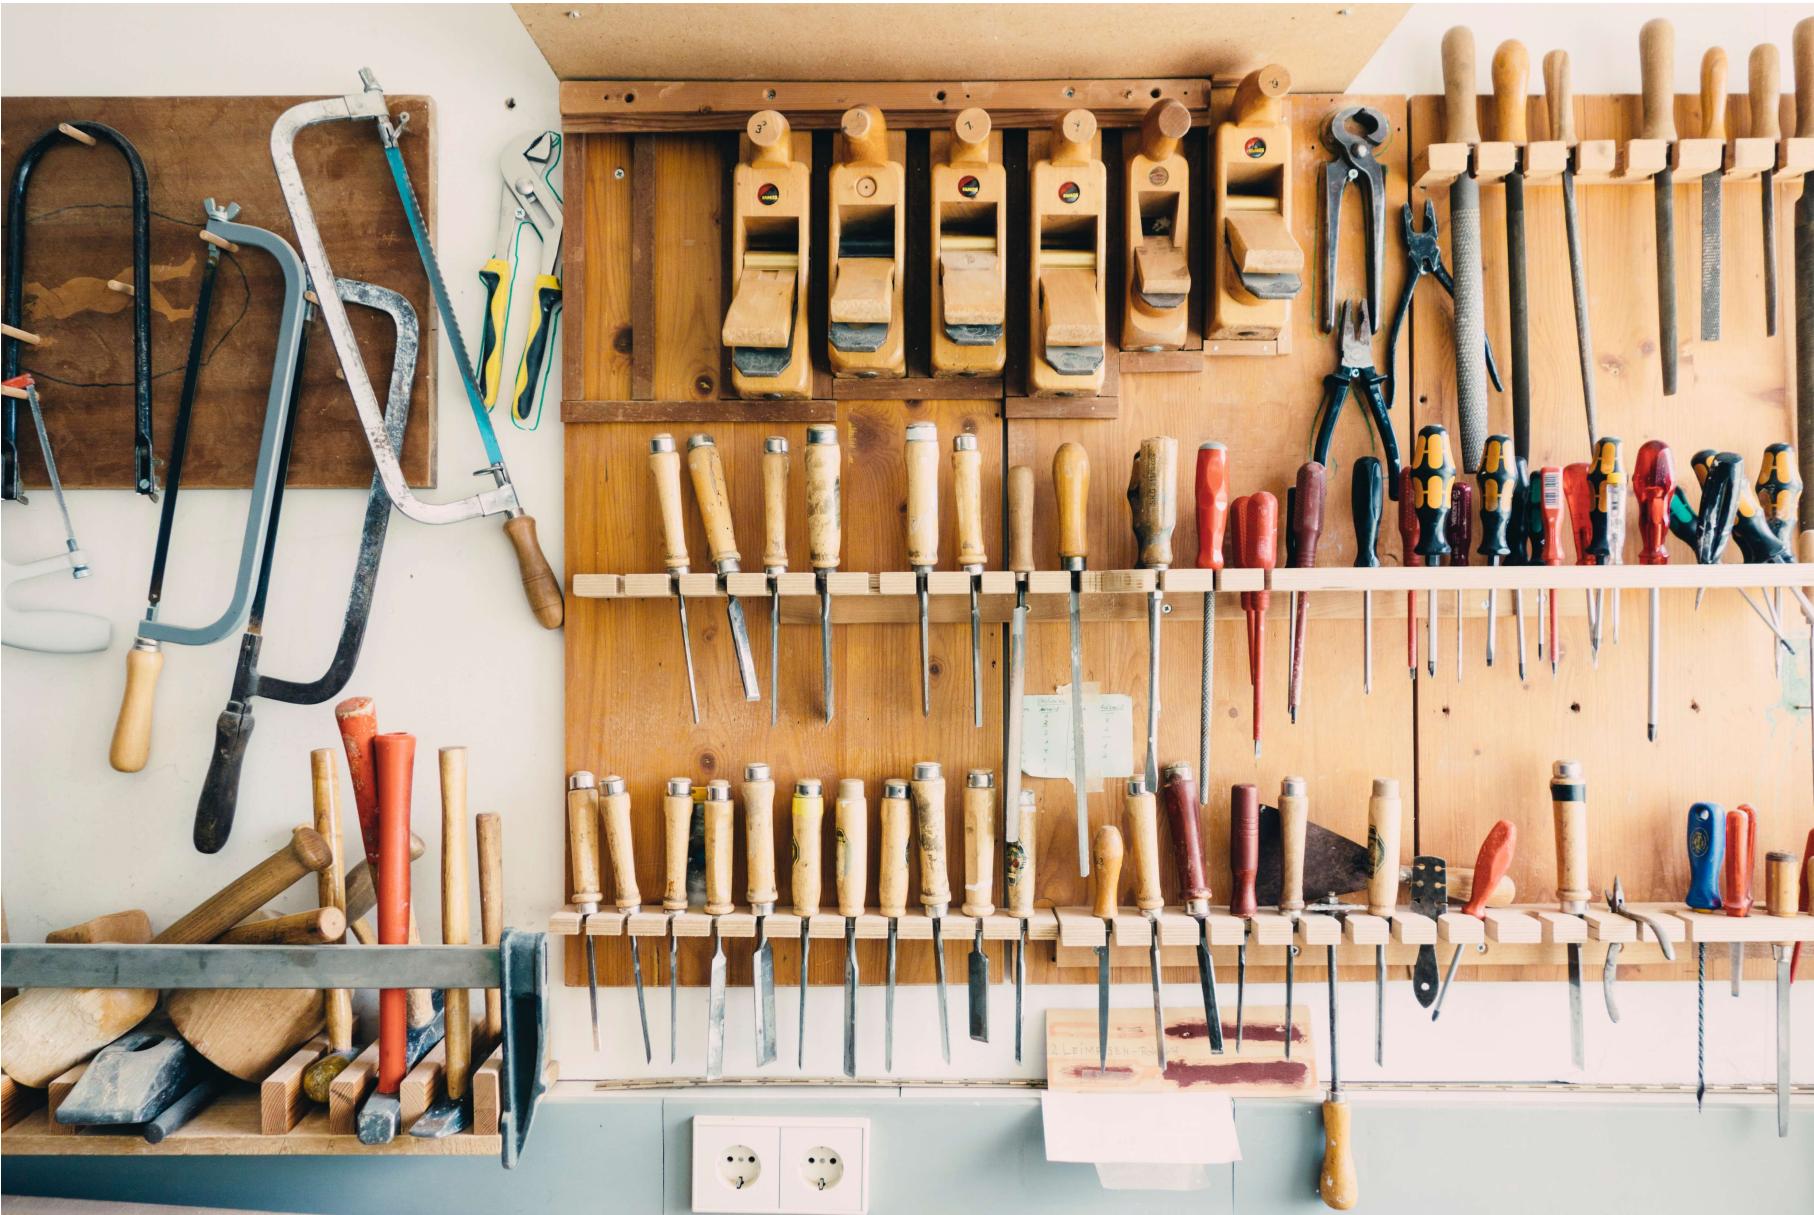


Workshop - Authentication





Workshop - OIDC Authentication





Authorisation

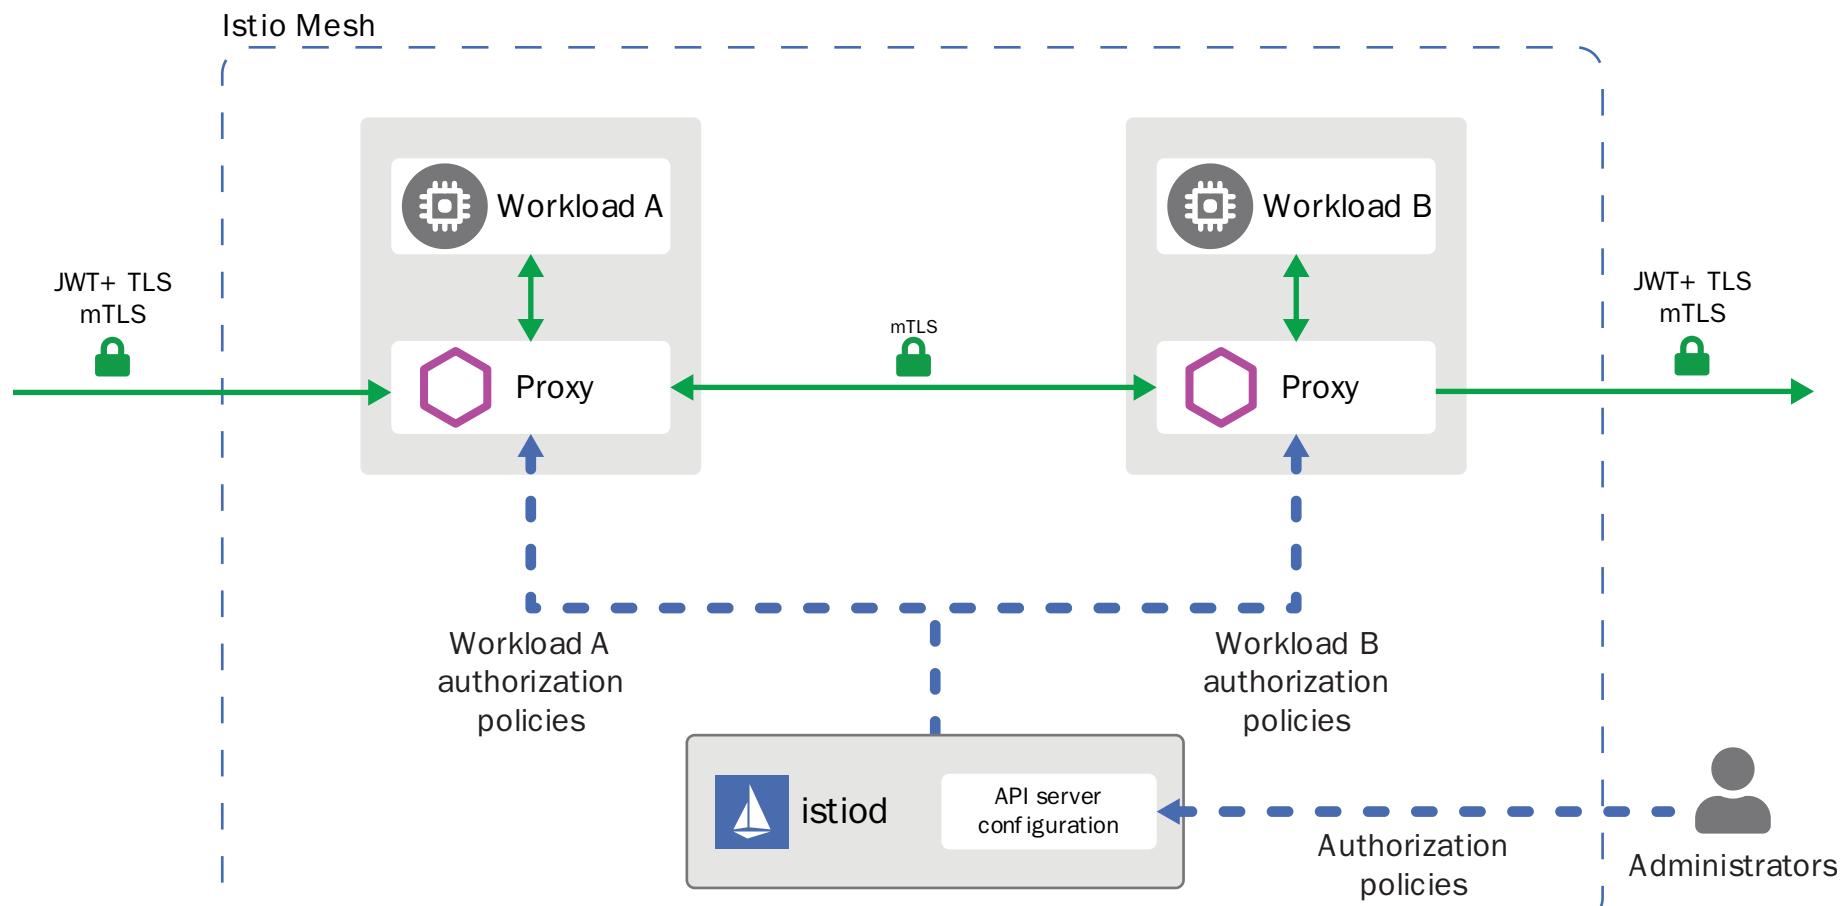
Istio's authorisation features provides mesh-level, namespace-level, and workload-level access control on workloads in an Istio Mesh.

This includes:

- Workload-to-workload and end-user-to-workload authorisation.
- A single `AuthorizationPolicy` CRD.
- Flexible semantics, operators can define custom conditions on Istio attributes.
- High performance, as Istio authorisation is enforced natively on Envoy.
- High compatibility, supports HTTP, HTTPS and HTTP2 natively, as well as any plain TCP protocols.



Authorisation architecture



<https://istio.io/docs/concepts/security/authz.svg>



There is no need to explicitly enable Istio's authorisation feature, only the AuthorizationPolicy needs applying on workloads to enforce access control.

If no AuthorizationPolicy applies to a workload, no access control will be enforced. In other words, all requests will be allowed.

If any AuthorizationPolicy applies to a workload, access to that workload is denied by default, unless explicitly allowed by a rule declared in the policy.



AuthorizationPolicy

To configure an Istio authorization policy, an `AuthorizationPolicy` resource is created.

An authorization policy includes a selector, an action, and a list of rules.

- The `selector` specifies the target that the policy
- The `action` field specifies whether to allow or deny the request
- The `rules` specify when to trigger the action
 - The `from` field in the `rules` specifies the sources of the request
 - The `to` field in the `rules` specifies the operations of the request
 - The `when` field specifies the conditions needed to apply the rule



Allow-all and deny-all

The example below shows an allow-all policy which allows full access to all workloads in the default namespace.

```
apiVersion: security.istio.io/v1beta1
kind: AuthorizationPolicy
metadata:
  name: allow-all
  namespace: default
spec:
  rules:
  - {}
```



The example below shows a deny-all policy which denies access to all workloads in the admin namespace.

```
apiVersion: security.istio.io/v1beta1
kind: AuthorizationPolicy
metadata:
  name: deny-all
  namespace: admin
spec:
  { }
```



Authenticated and unauthenticated identity

To make workloads publicly accessible, the `source` section is left empty.

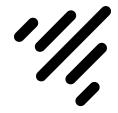
This allows sources from all (both authenticated and unauthenticated) users and workloads.

```
apiVersion: security.istio.io/v1beta1
kind: AuthorizationPolicy
metadata:
  name: httpbin
  namespace: foo
spec:
  selector:
    matchLabels:
      app: httpbin
      version: v1
  rules:
  - to:
    - operation:
        methods: [ "GET", "POST" ]
```



To allow only authenticated users, set principals to "*" instead, for example:

```
apiVersion: security.istio.io/v1beta1
kind: AuthorizationPolicy
metadata:
  name: httpbin
  namespace: foo
spec:
  selector:
    matchLabels:
      app: httpbin
      version: v1
  rules:
  - from:
    - source:
        principals: [ "*" ]
    to:
    - operation:
        methods: [ "GET", "POST" ]
```



Policy Target

Policy scope (target) is determined by metadata/namespace and an optional selector.

```
apiVersion: security.istio.io/v1beta1
kind: AuthorizationPolicy
metadata:
  name: allow-read
  namespace: default
spec:
  selector:
    matchLabels:
      app: products
  rules:
  - to:
    - operation:
        methods: ["GET", "HEAD"]
```

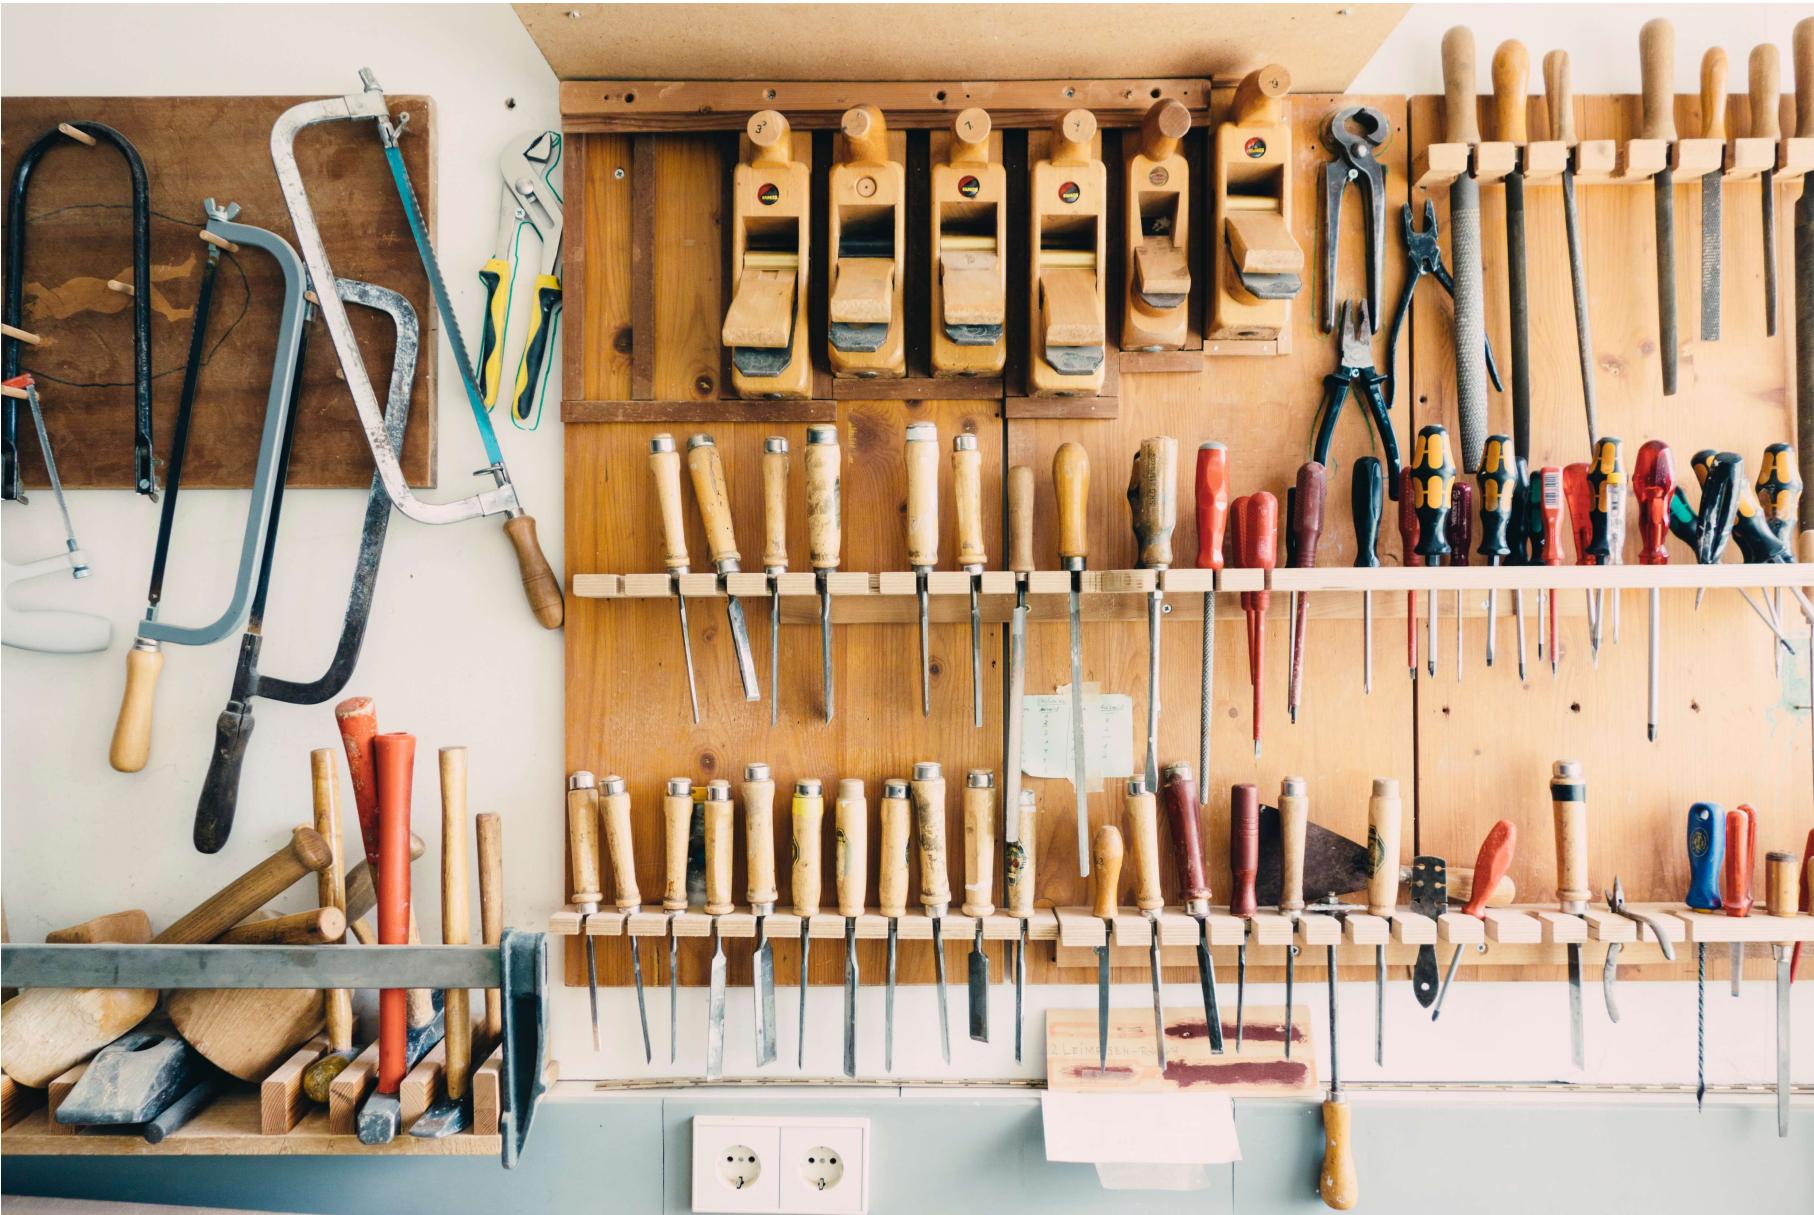


Full example

```
apiVersion: security.istio.io/v1beta1
kind: AuthorizationPolicy
metadata:
  name: httpbin
  namespace: foo
spec:
  selector:
    matchLabels:
      app: httpbin
      version: v1
  rules:
  - from:
    - source:
        principals: ["cluster.local/ns/default/sa/sleep"]
    - source:
        namespaces: ["dev"]
    to:
    - operation:
        methods: ["GET"]
        paths: ["/info*"]
  when:
  - key: request.auth.claims[groups]
    values: ["group1"]
```



Workshop - Istio authorisation





Observability

Istio generates detailed telemetry for all service communications within a mesh.

Through Istio, operators gain a thorough understanding of how monitored services are interacting, both with other services and with the Istio components themselves.



Istio generates the following types of telemetry in order to provide overall service mesh observability:

- *Metrics* – Istio generates a set of service metrics based on latency, traffic, errors, and saturation as well as detailed metrics for the mesh control plane.



Istio generates the following types of telemetry in order to provide overall service mesh observability:

- *Distributed Traces* - Istio generates distributed trace spans for each service.



Istio generates the following types of telemetry in order to provide overall service mesh observability:

- *Access Logs* – As traffic flows into a service within a mesh, Istio can generate a full record of each request, including source and destination metadata.



Metrics

Metrics provide a way of monitoring and understanding behaviour across the mesh, generating metrics for all service traffic in, out, and within an Istio service mesh.



In addition to monitoring the behaviour of services within a mesh, it is also important to monitor the behaviour of the mesh itself.

Istio components export metrics on their own internal behaviours to provide insight on the health and function of the mesh control plane.



Proxy-level metrics

Istio metrics collection begins with the sidecar proxies (Envoy).

Each proxy generates a rich set of metrics about all traffic passing through the proxy (both inbound and outbound).



As a result, understanding the connection between mesh services and Envoy resources is required for monitoring the Envoy metrics.

Istio enables operators to select which of the Envoy metrics are generated and collected at each workload instance.

This enables targeted debugging of networking behaviour, while reducing the overall cost of monitoring across the mesh.



Example proxy-level Metrics:

```
envoy_cluster_internal_upstream_rq{response_code_class="2xx",cluster_name="xds-grpc"} 1000  
envoy_cluster_upstream_rq_completed{cluster_name="xds-grpc"} 7164  
envoy_cluster_ssl_connection_error{cluster_name="xds-grpc"} 0  
envoy_cluster_lb_subsets_removed{cluster_name="xds-grpc"} 0  
envoy_cluster_internal_upstream_rq{response_code="503",cluster_name="xds-grpc"} 0
```



Service-level metrics

In addition to the proxy-level metrics, Istio provides a set of service-oriented metrics for monitoring service communications.

These metrics cover the four basic service monitoring needs: latency, traffic, errors, and saturation.



Example service-level metric:

```
istio_requests_total{  
    connection_security_policy="mutual_tls",  
    destination_app="details",  
    destination_principal="cluster.local/ns/default/sa/default",  
    destination_service="details.default.svc.cluster.local",  
    destination_service_name="details",  
    destination_service_namespace="default",  
    destination_version="v1",  
    destination_workload="details-v1",  
    destination_workload_namespace="default",  
    reporter="destination",  
    request_protocol="http",  
    response_code="200",  
    response_flags="-",  
    source_app="productpage",  
    source_principal="cluster.local/ns/default/sa/default",  
    source_version="v1",  
    source_workload="productpage-v1",  
    source_workload_namespace="default"  
} 214
```



istioctl dashboard prometheus

Prometheus Alerts Graph Status ▾ Help

Enable query history

istio_requests_total

Load time: 14ms
Resolution: 14s
Total time series: 15

Execute - insert metric at cursor - ▾

Graph Console

Element

```
istio_requests_total(connection_mtls="false",destination_app="details",destination_namespace="default",destination_principal="unknown",destination_service="details.default.svc.cluster.local",destination_service_name="details",destination_service_namespace="default",destination_v1,destination_workload_namespace="default",instance="172.17.0.18:42422",job="istio-mesh",reporter="client",request_protocol="http",response_code="200",source_app="productpage",source_namespace="default",source_principal="unknown",source_version="v1",source_workload_name="details-v1")  
istio_requests_total(connection_mtls="false",destination_app="details",destination_namespace="default",destination_principal="unknown",destination_service="details.default.svc.cluster.local",destination_service_name="details",destination_service_namespace="default",destination_v1,destination_workload_namespace="default",instance="172.17.0.18:42422",job="istio-mesh",reporter="server",request_protocol="http",response_code="200",source_app="productpage",source_namespace="default",source_principal="unknown",source_version="v1",source_workload_name="details-v1")  
istio_requests_total(connection_mtls="false",destination_app="policy",destination_namespace="istio-system",destination_principal="unknown",destination_service="istio-policy.istio-system.svc.cluster.local",destination_service_name="istio-policy",destination_service_namespace="istio-system",destination_version="unknown",destination_workload="istio-policy",destination_workload_namespace="istio-system",instance="172.17.0.18:42422",job="istio-mesh",reporter="server",request_protocol="grpc",response_code="200",source_app="details",source_namespace="default",source_principal="unknown",source_version="v1",source_workload="details-v1",source_workload_namespace="default")  
istio_requests_total(connection_mtls="false",destination_app="policy",destination_namespace="istio-system",destination_principal="unknown",destination_service="istio-policy.istio-system.svc.cluster.local",destination_service_name="istio-policy",destination_service_namespace="istio-system",destination_version="unknown",destination_workload="istio-policy",destination_workload_namespace="istio-system",instance="172.17.0.18:42422",job="istio-mesh",reporter="server",request_protocol="grpc",response_code="200",source_app="productpage",source_namespace="default",source_principal="unknown",source_version="v1",source_workload="productpage-v1",source_workload_namespace="default")  
istio_requests_total(connection_mtls="false",destination_app="policy",destination_namespace="istio-system",destination_principal="unknown",destination_service="istio-policy.istio-system.svc.cluster.local",destination_service_name="istio-policy",destination_service_namespace="istio-system",destination_version="unknown",destination_workload="istio-policy",destination_workload_namespace="istio-system",instance="172.17.0.18:42422",job="istio-mesh",reporter="server",request_protocol="grpc",response_code="200",source_app="reviews",source_namespace="default",source_principal="unknown",source_version="v1",source_workload="reviews-v1",source_workload_namespace="default")  
istio_requests_total(connection_mtls="false",destination_app="policy",destination_namespace="istio-system",destination_principal="unknown",destination_service="istio-policy.istio-system.svc.cluster.local",destination_service_name="istio-policy",destination_service_namespace="istio-system",destination_version="unknown",destination_workload="istio-policy",destination_workload_namespace="istio-system",instance="172.17.0.18:42422",job="istio-mesh",reporter="server",request_protocol="grpc",response_code="200",source_app="unknown",source_namespace="istio-ingressgateway",source_principal="unknown",source_version="unknown",source_workload="istio-ingressgateway",source_workload_namespace="istio-system")  
istio_requests_total(connection_mtls="false",destination_app="productpage",destination_namespace="default",destination_principal="unknown",destination_service="productpage.default.svc.cluster.local",destination_service_name="productpage",destination_service_namespace="default",destination_v1,destination_workload_namespace="default",instance="172.17.0.18:42422",job="istio-mesh",reporter="client",request_protocol="http",response_code="200",source_app="unknown",source_namespace="istio-system",source_principal="unknown",source_version="unknown",source_workload_name="productpage-v1")
```



Control plane metrics

Each Istio component (Pilot, Galley, Mixer) also provides a collection of self-monitoring metrics.

These metrics allow monitoring of the behaviour of Istio itself (as distinct from that of the services within the mesh).



Grafana

```
kubectl apply -f ${ISTIO_INSTALL_DIR}/samples/addons/grafana.yaml  
istioctl dashboard grafana
```



Istio Mesh Dashboard

Last 5 minutes Refresh every 5s

Global Request Volume: 1.4 ops

Global Success Rate (non-5xx responses): 100%

4xxs: N/A

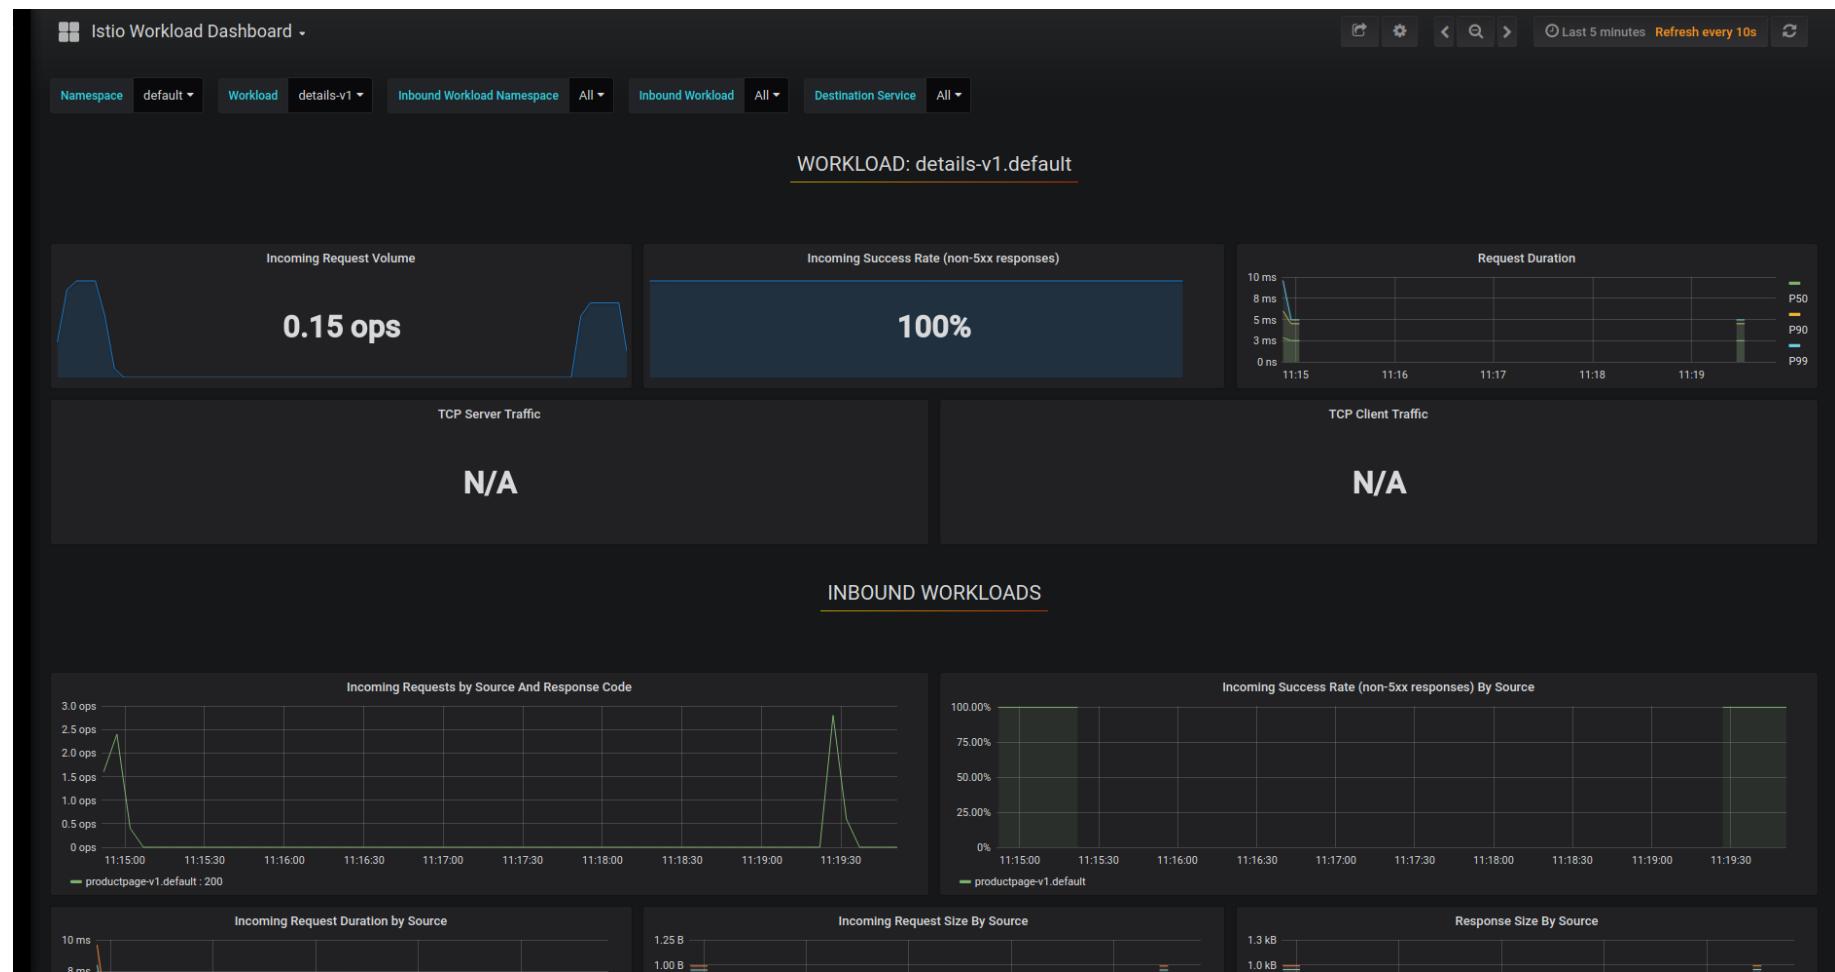
5xxs: N/A

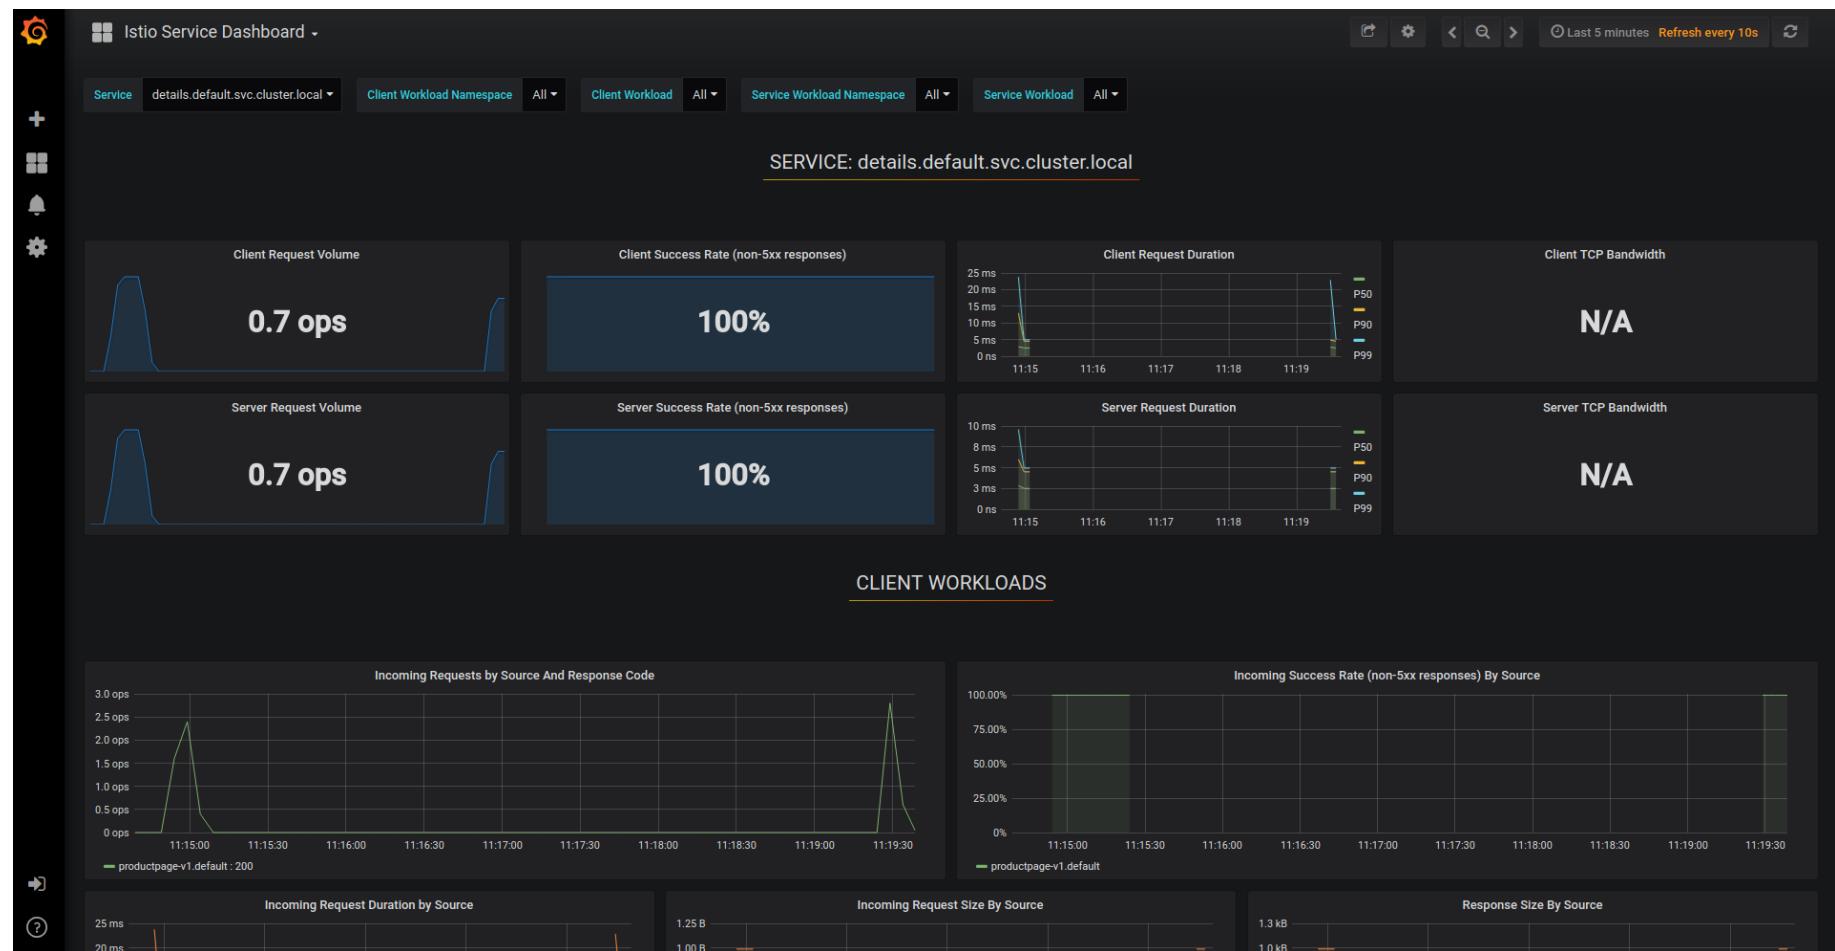
Istio is an [open platform](#) that provides a uniform way to connect, [manage](#), and [secure](#) microservices. Need help? Join the [Istio community](#).

HTTP/GRPC Workloads

Service	Workload	Requests	P50 Latency	P90 Latency	P99 Latency	Success Rate
reviews.default.svc.cluster.local	reviews-v3.default	0.02 ops	18 ms	24 ms	25 ms	100.00%
reviews.default.svc.cluster.local	reviews-v2.default	0.05 ops	31 ms	46 ms	50 ms	100.00%
reviews.default.svc.cluster.local	reviews-v1.default	0.07 ops	8 ms	19 ms	24 ms	100.00%
ratings.default.svc.cluster.local	ratings-v1.default	0.07 ops	3 ms	5 ms	5 ms	100.00%
productpage.default.svc.cluster.local	productpage-v1.default	0.15 ops	25 ms	60 ms	96 ms	100.00%
istio-telemetry.istio-system.svc.cluster.local	istio-telemetry.istio-system	0.85 ops	3 ms	7 ms	18 ms	100.00%
istio-policy.istio-system.svc.cluster.local	istio-policy.istio-system	0.29 ops	3 ms	5 ms	5 ms	100.00%
details.default.svc.cluster.local	details-v1.default	0.15 ops	3 ms	5 ms	5 ms	100.00%

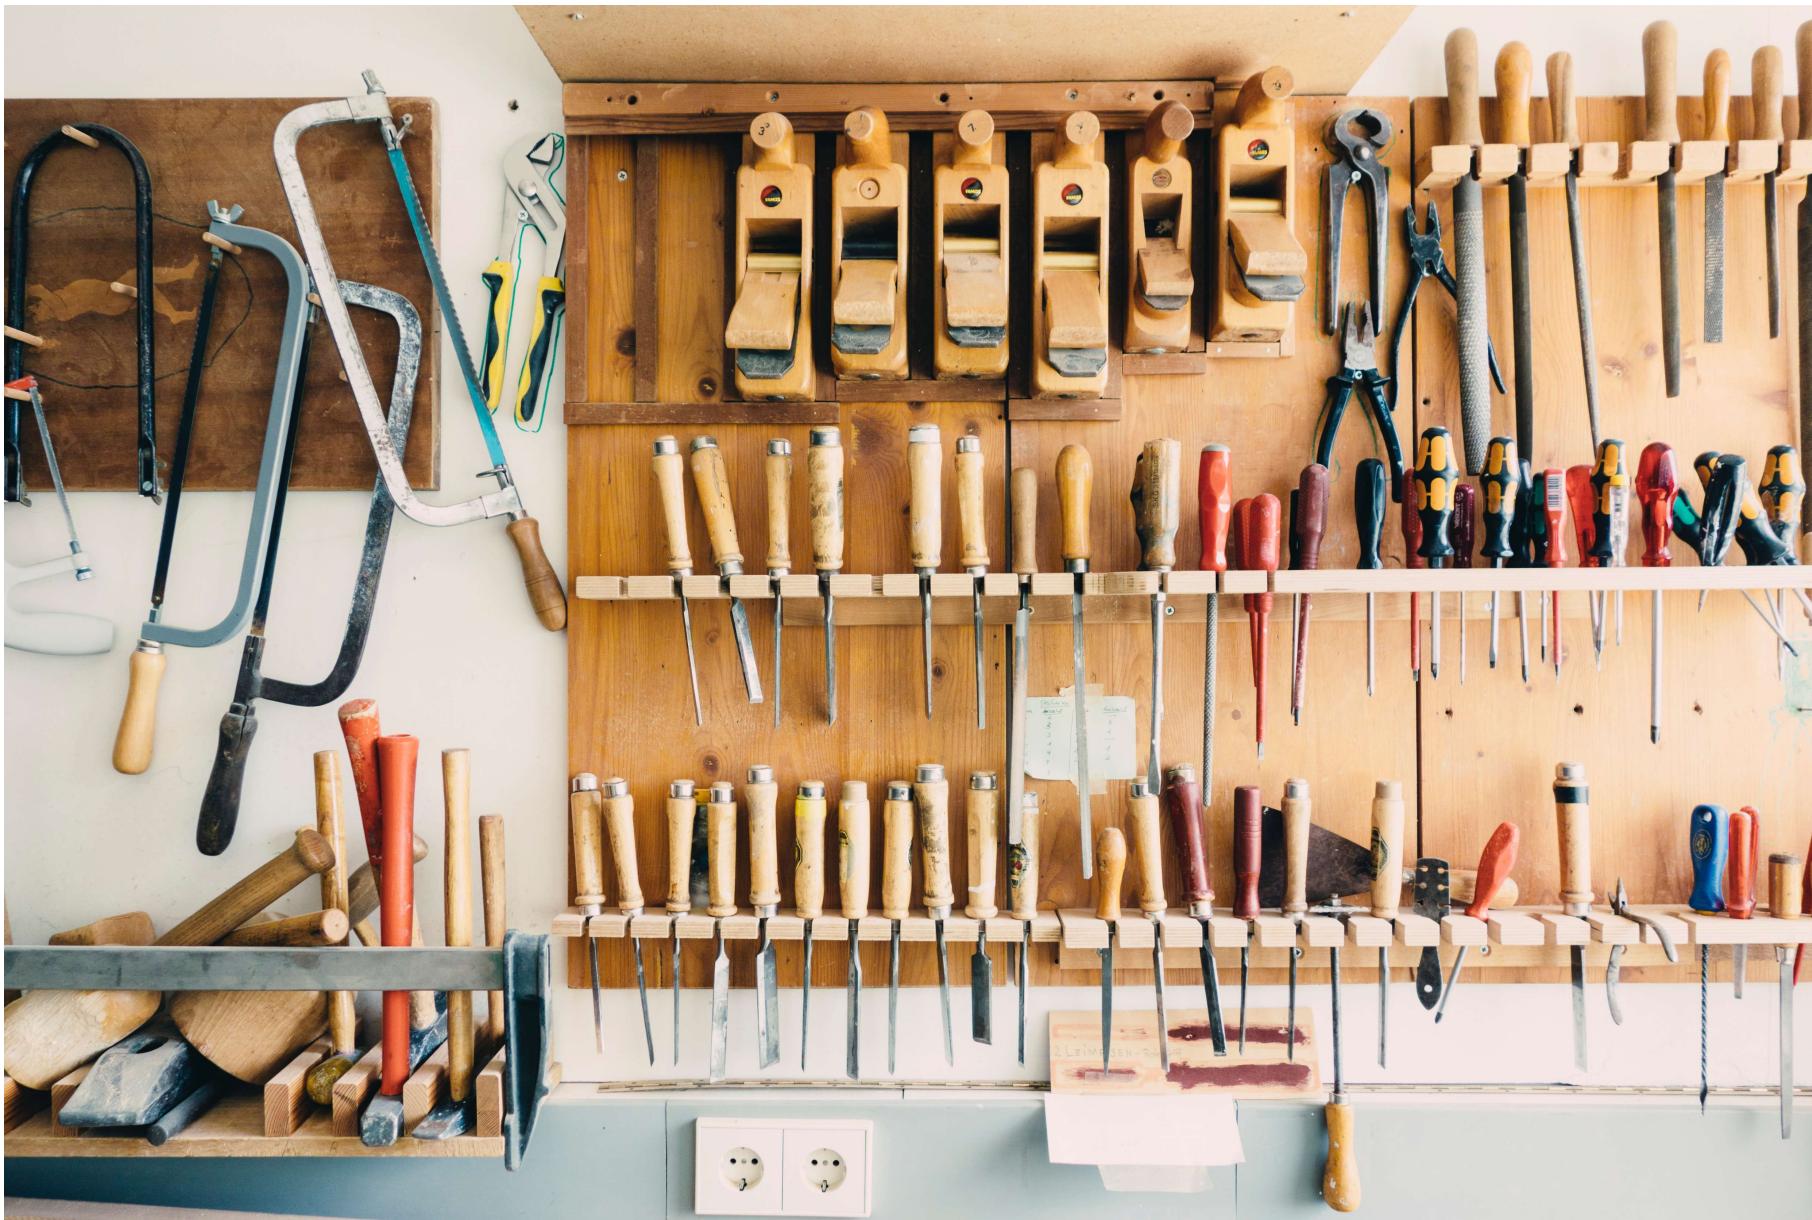
Back Help







Workshop - Monitoring Blocked External Service Traffic





Observability

Logging



Access Logs

Access logs provide a way to monitor and understand behaviour from the perspective of an individual workload instance.

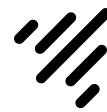
Istio can generate access logs for service traffic in a configurable set of formats, providing operators with full control of the how, what, when and where of logging.



Istio exposes a full set of source and destination metadata to the access logging mechanisms, allowing detailed audit of network transactions.

Access logs may be generated locally or exported to custom backends, including Fluentd.

Logs can be generated in TEXT or JSON format



Example Istio access log (formatted in JSON):

```
{  
    "protocol": "HTTP/1.1",  
    "duration": 16,  
    "upstream_local_address": "10.60.2.8:43826",  
    "response_flags": "-",  
    "response_code": 418,  
    "path": "/status/418",  
    "upstream_transport_failure_reason": null,  
    "x_forwarded_for": null,  
    "upstream_host": "10.60.1.7:80",  
    "upstream_cluster": "outbound|8000||httpbin.default.svc.cluster.local"  
    "connection_termination_details": null,  
    "start_time": "2021-07-29T08:55:33.467Z",  
    "request_id": "9c9ecb6e-c240-4d17-b8c3-7bbac329faa2",  
    "user_agent": "curl/7.78.0-DEV",  
    "upstream_service_time": "15",  
    "response_code_details": "via_upstream",  
    "requested_server_name": null,  
    "method": "GET",  
    "bytes_sent": 135,  
    "authority": "httpbin:8000",  
    "downstream_remote_address": "10.60.2.8:47826",  
    "bytes_received": 0,  
    "downstream_local_address": "10.63.243.74:8000",  
    "route_name": "default"  
}
```



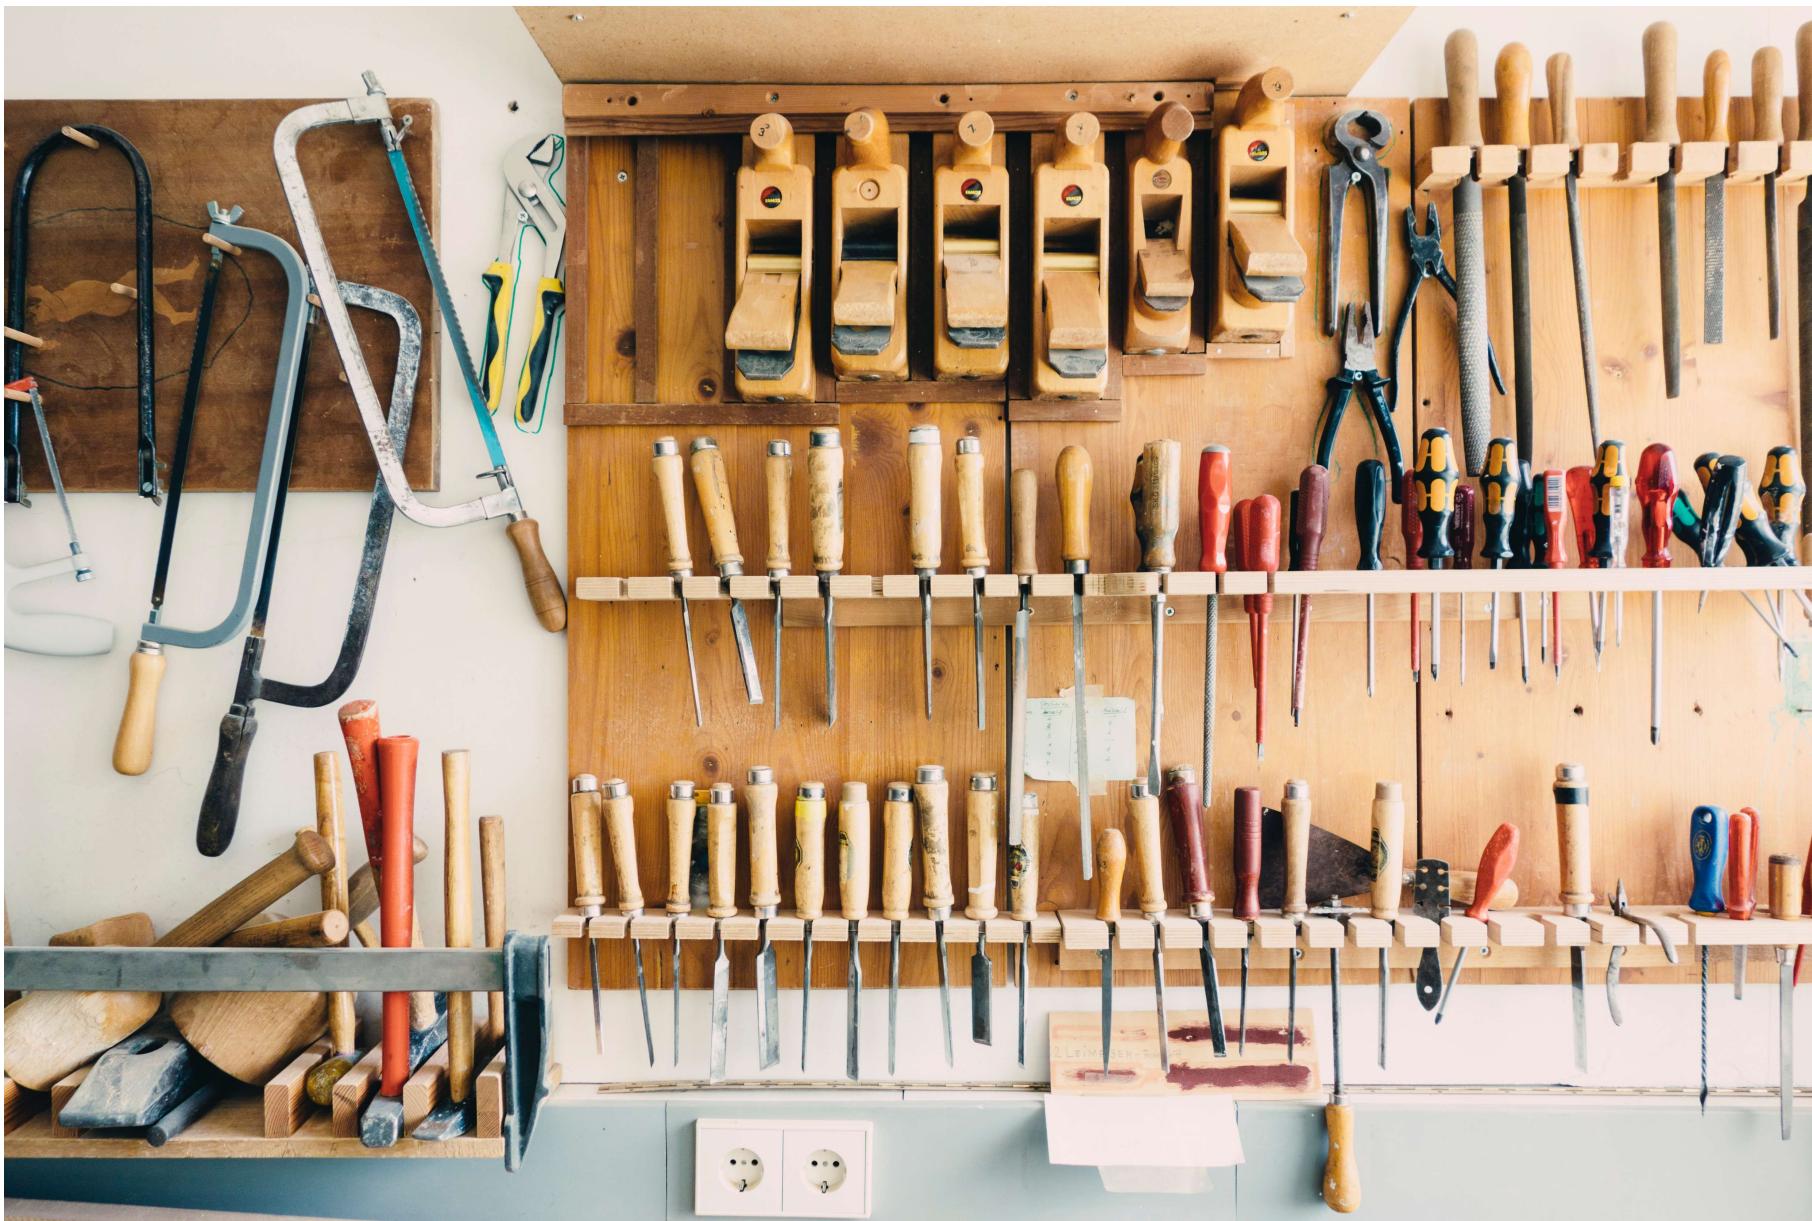
Logging configuration

Istio provide these 3 settings to configure the access log on the Envoy proxies:

- `meshConfig.accessLogFile`
- `meshConfig.accessLogEncoding`
- `meshConfig.accessLogFormat`



Workshops – Visualising metrics with Grafana





Distributed Tracing



Distributed tracing provides a way to monitor and understand behaviour by monitoring individual requests as they flow through a mesh.

Traces empower mesh operators to understand service dependencies and the sources of latency within their service mesh.



Istio supports distributed tracing through the Envoy proxies.

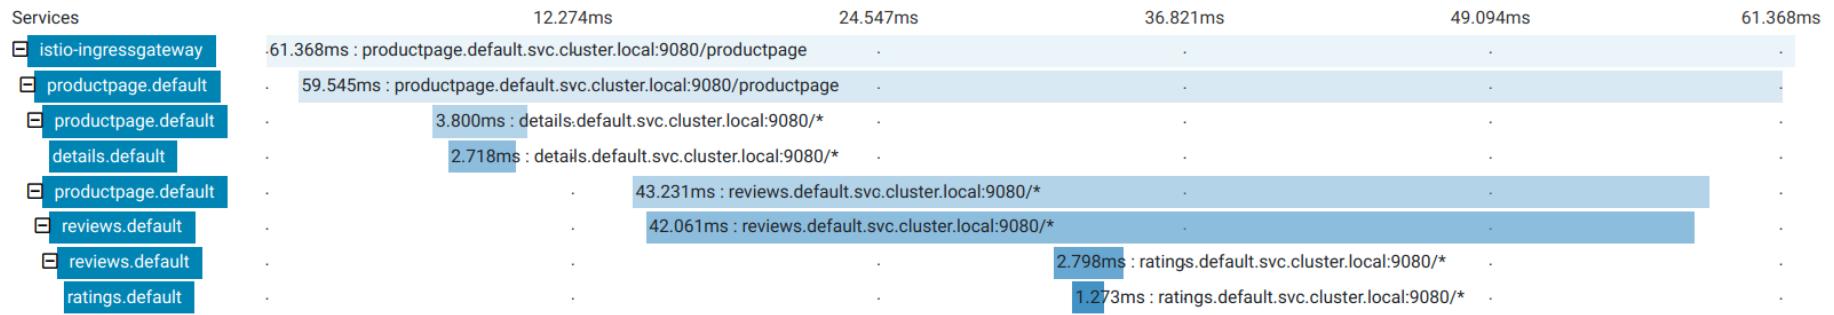
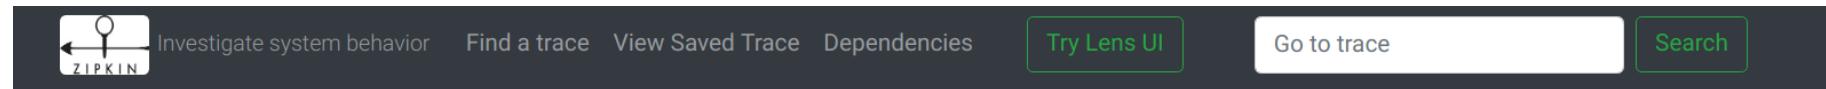
The proxies automatically generate trace spans on behalf of the applications they proxy, requiring only that the applications forward the appropriate request context.



Istio supports a number of tracing backends, including Zipkin, Jaeger, LightStep, and Datadog.



istioctl dashboard zipkin|jaeger





Trace context propagation

Istio proxies can automatically send spans, however applications need to propagate the appropriate HTTP headers so that spans can be correlated correctly into a single trace.



To do this, an application needs to collect and propagate the following headers from the incoming request to any outgoing requests:

- x-request-id
- x-b3-traceid
- x-b3-spanid
- x-b3-parentspanid
- x-b3-sampled
- x-b3-flags
- x-ot-span-context



```
def getForwardHeaders(request):
    headers = {}
    # x-b3-*** headers can be populated using the opentracing span
    span = get_current_span()
    carrier = {}
    tracer.inject(
        span_context=span.context,
        format=Format.HTTP_HEADERS,
        carrier=carrier)
    headers.update(carrier)
    incoming_headers = ['x-request-id']
    for ihdr in incoming_headers:
        val = request.headers.get(ihdr)
        if val is not None:
            headers[ihdr] = val
    return headers
```

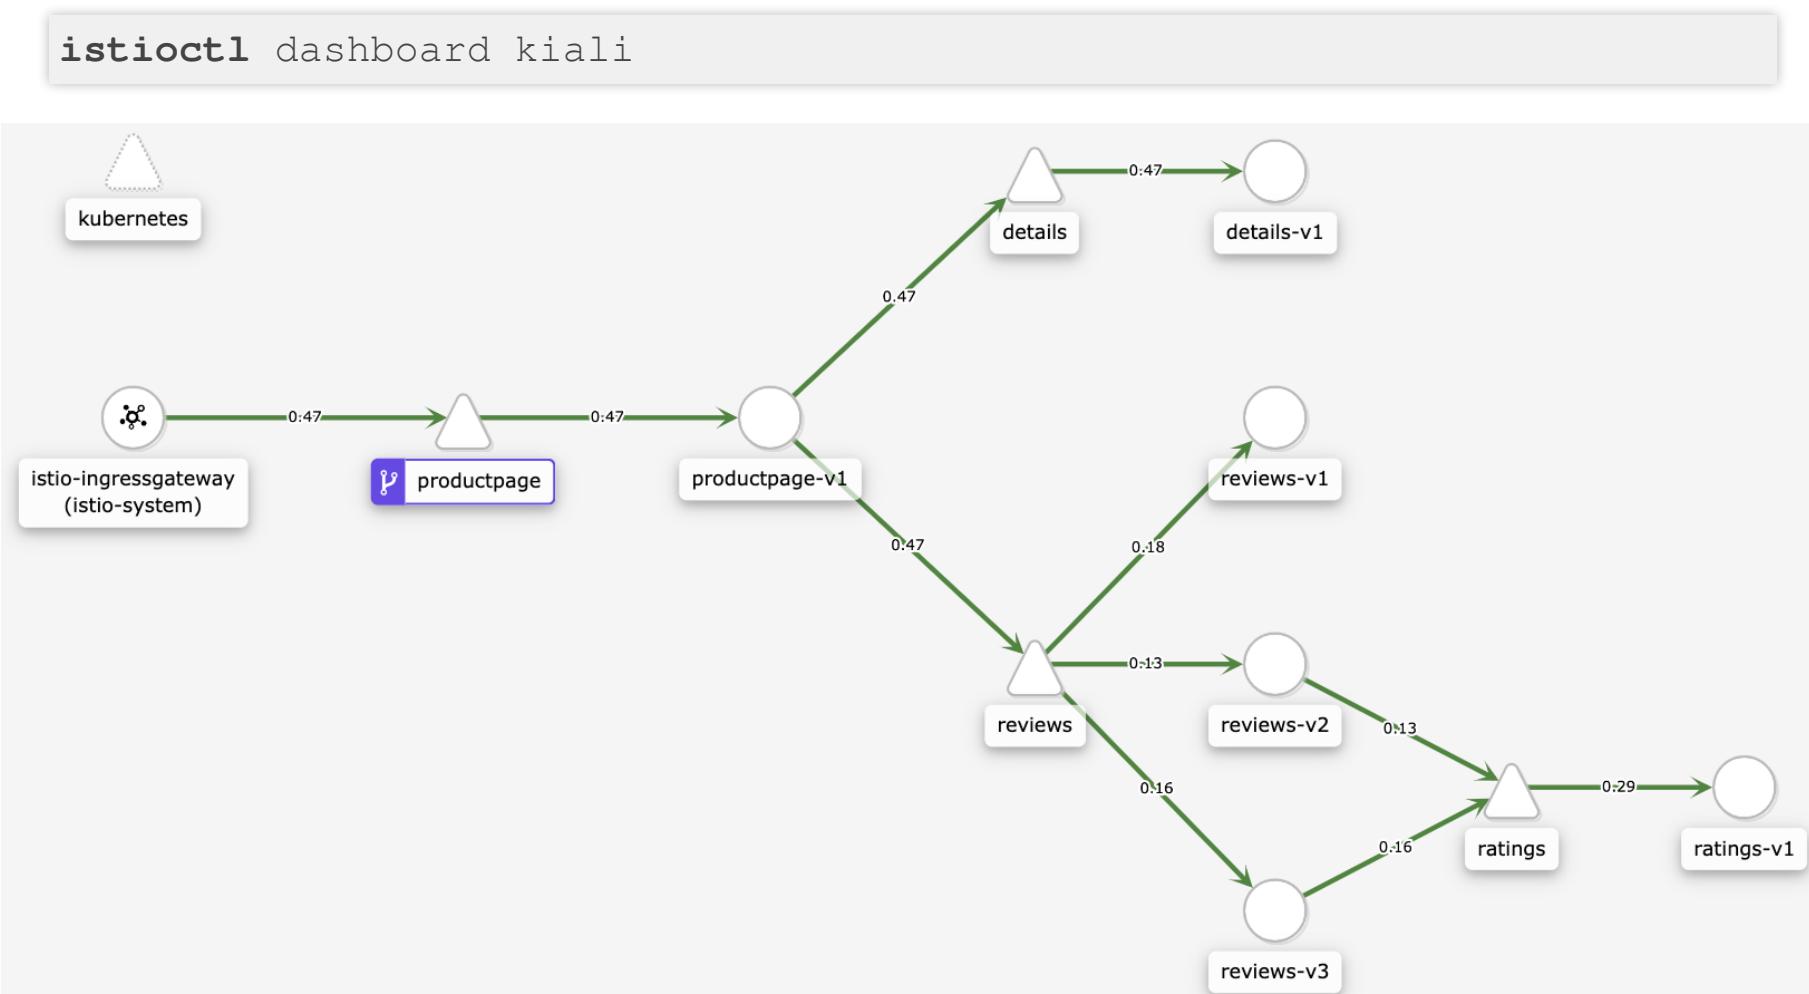


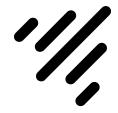
Workshops – Tracing requests with Jaeger



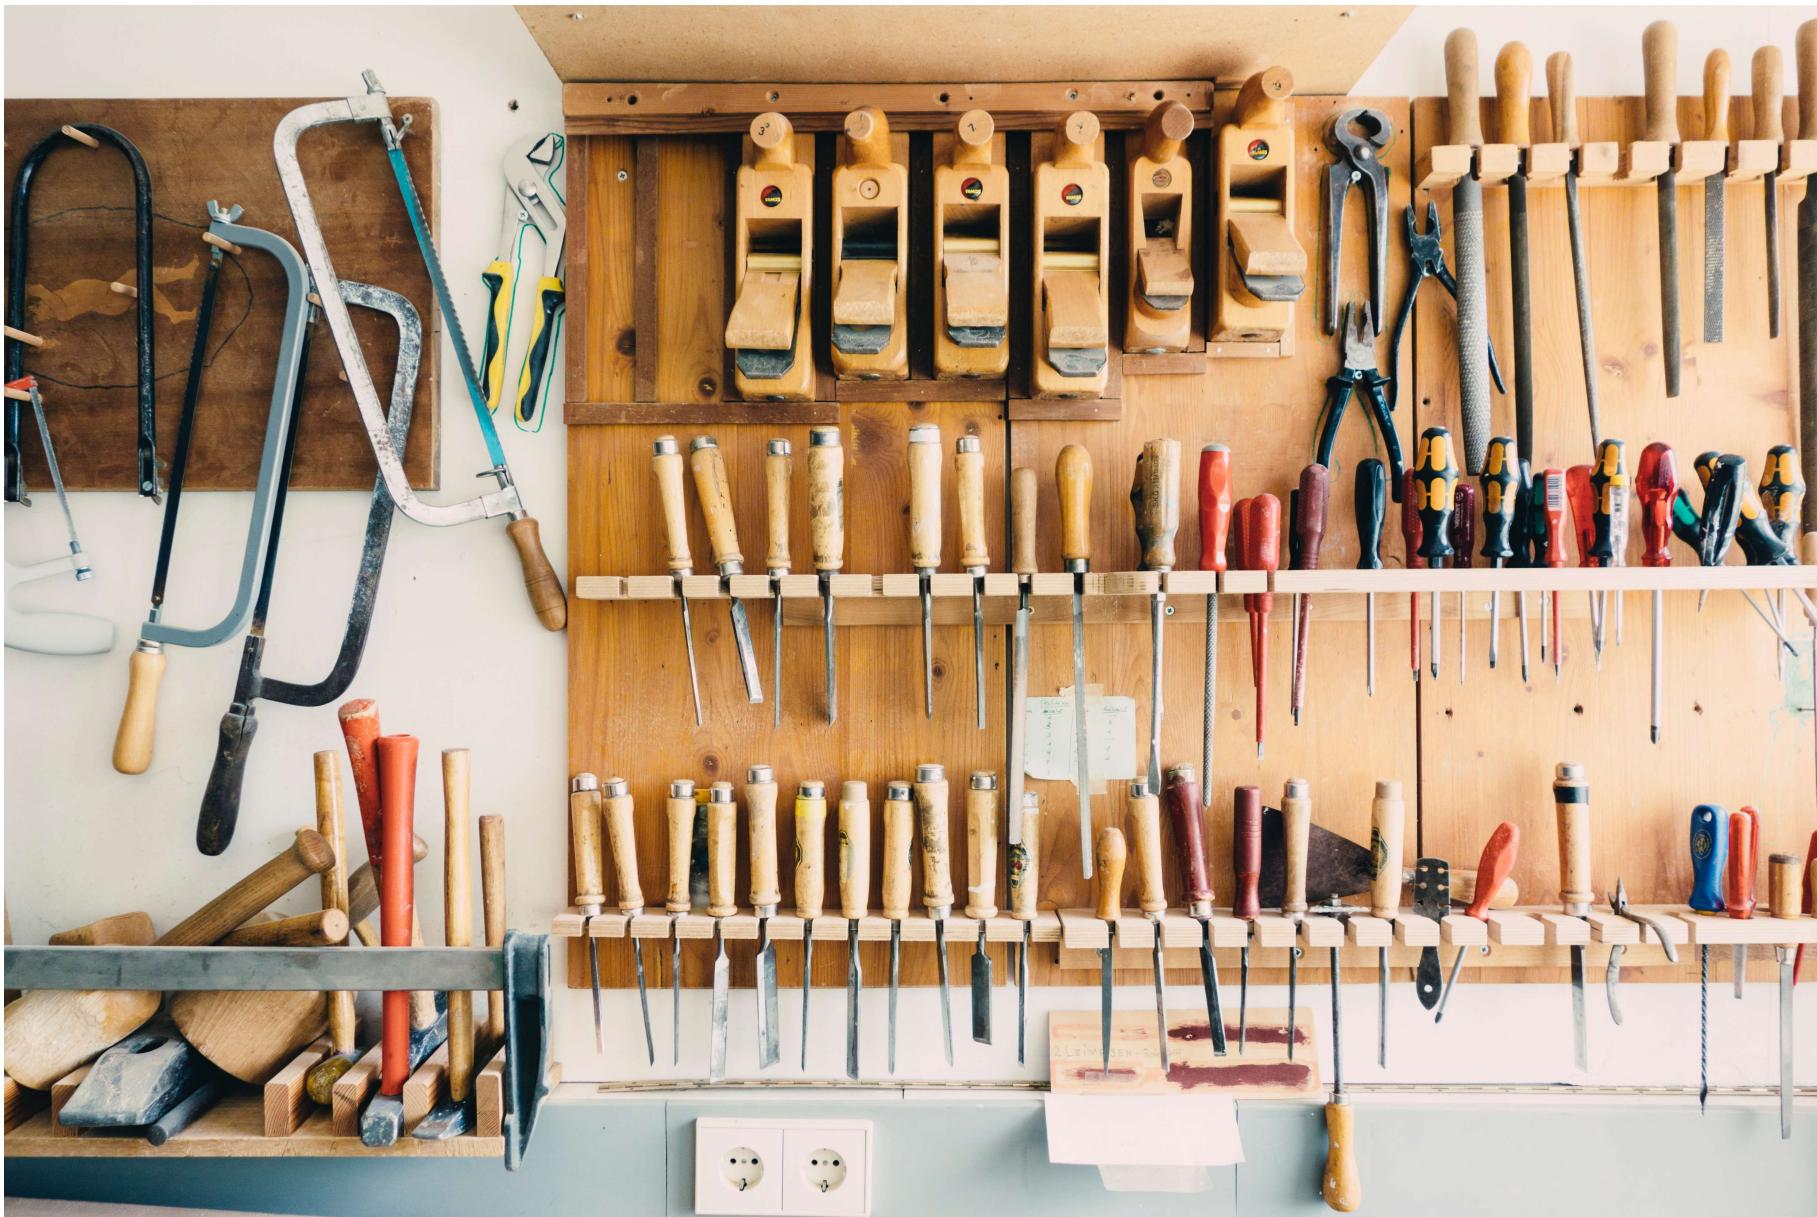


Kiali





Workshops – Observability with Kiali





Finally...





Training

- Kubernetes in Practice (Day One and Day Two)
- Kubernetes for Application Developers (Day One and Day Two)
- Istio in Practice (Day One and Day Two)
- Operational Wargaming
- Kubernetes from Scratch
- Extending Kubernetes with Operators
- Serverless on Kubernetes with Knative



Consulting

On-site / Remote / Project-based

Production Readiness Review

Cloud Native Discovery Workshop



Jetstack Subscription

Jetstack Subscription is designed to help you verify, support and continually optimise your Kubernetes production environment.

-  Break Fix Support →
-  Customer Reliability Engineering →
-  Kubernetes Reference Architecture →
-  Automated Cluster Compliance Checking →
-  Online Training and Wargaming →

[Register Interest](#)



Jetstack is trusted by





The End

