

Istio in Practice - Day 2

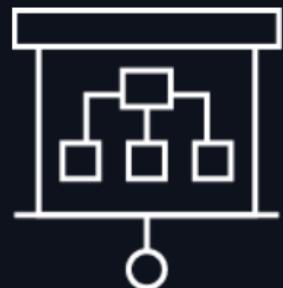




Introductions

Each person say:

- Hello & where they are from
- How much k8s/Istio experience they have
- The main thing they want to learn
- Favourite something computer related
- Favourite something not computer related



Training

Learn and engage directly alongside our team, with courses for all stages of your Kubernetes journey



Workshops

Enhancing Kubernetes knowledge with our workshops. Whichever stage of your cloud native journey you are on we can help you do it better



Consulting

Consulting and engineering to make the most of Kubernetes and move you to production quickly



Subscription

Reference architecture, online training and SLA support 24x7 for your production Kubernetes deployment



Agenda

- Envoy Deep Dive
- Mixer Deep Dive
- Extensibility
- Fault Injection
- Circuit Breaking
- Traffic Mirroring
- Deployment Models



What is a Service Mesh?

An architectural pattern that provides common network services as a feature of the infrastructure



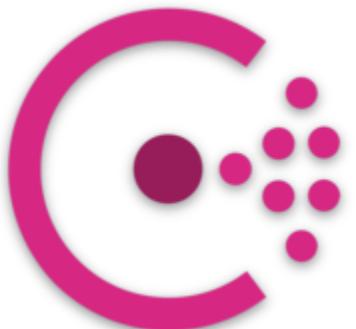
Why Do We Need a Service Mesh?

The scale and complexity of microservice architectures and distributed systems benefits greatly from a dedicated layer to orchestrate cross cutting concerns

- Facilitates service-to-service communication between microservices
- Abstracts the network and allows operators to think in terms of services rather than specific IPs and ports
- Functionality includes discovery, load balancing, failure recovery, policy enforcement, monitoring and telemetry

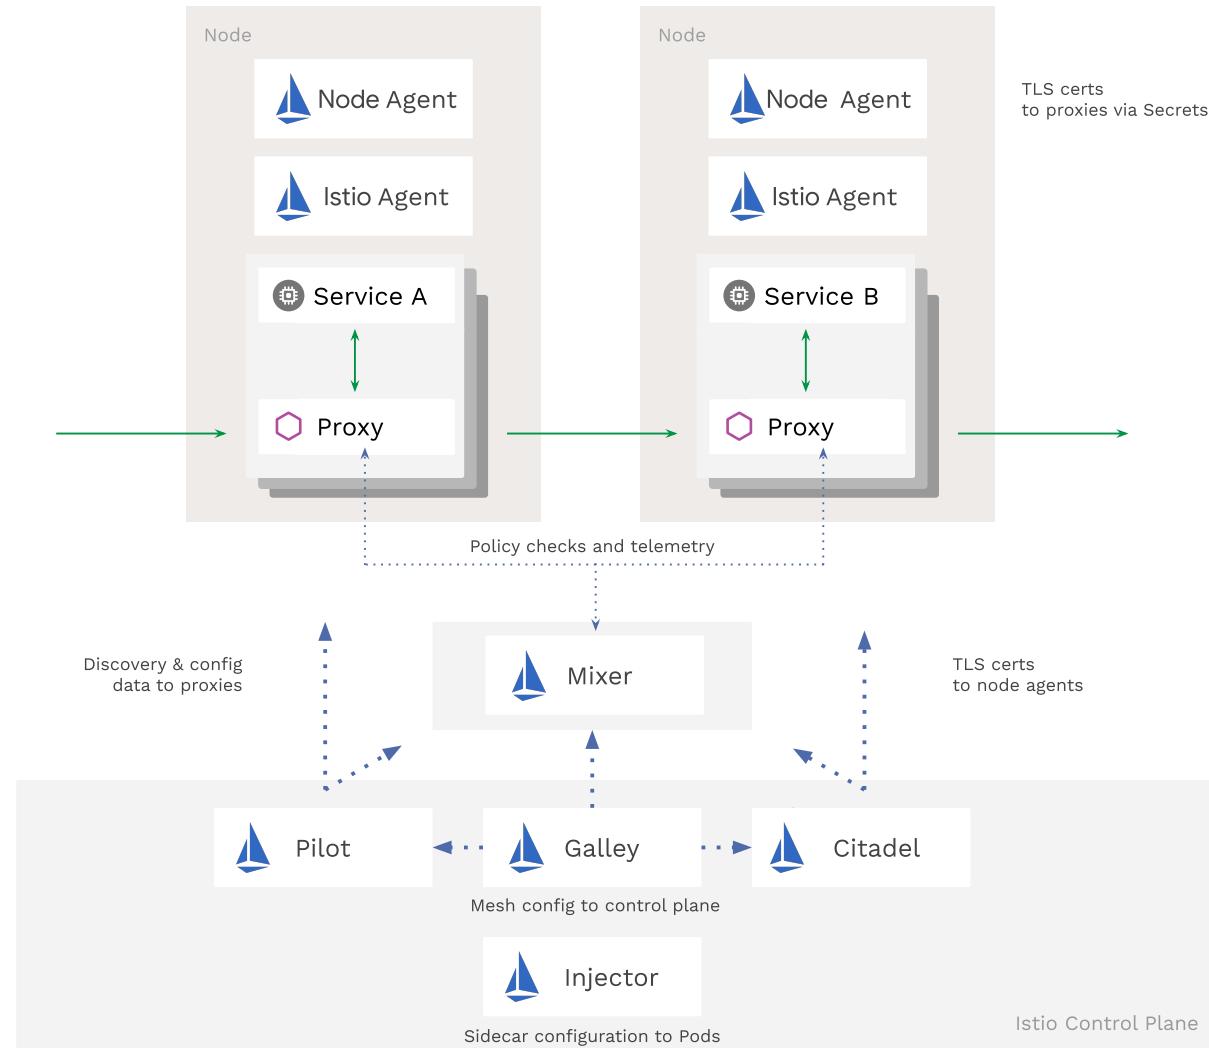


Service Mesh Implementations



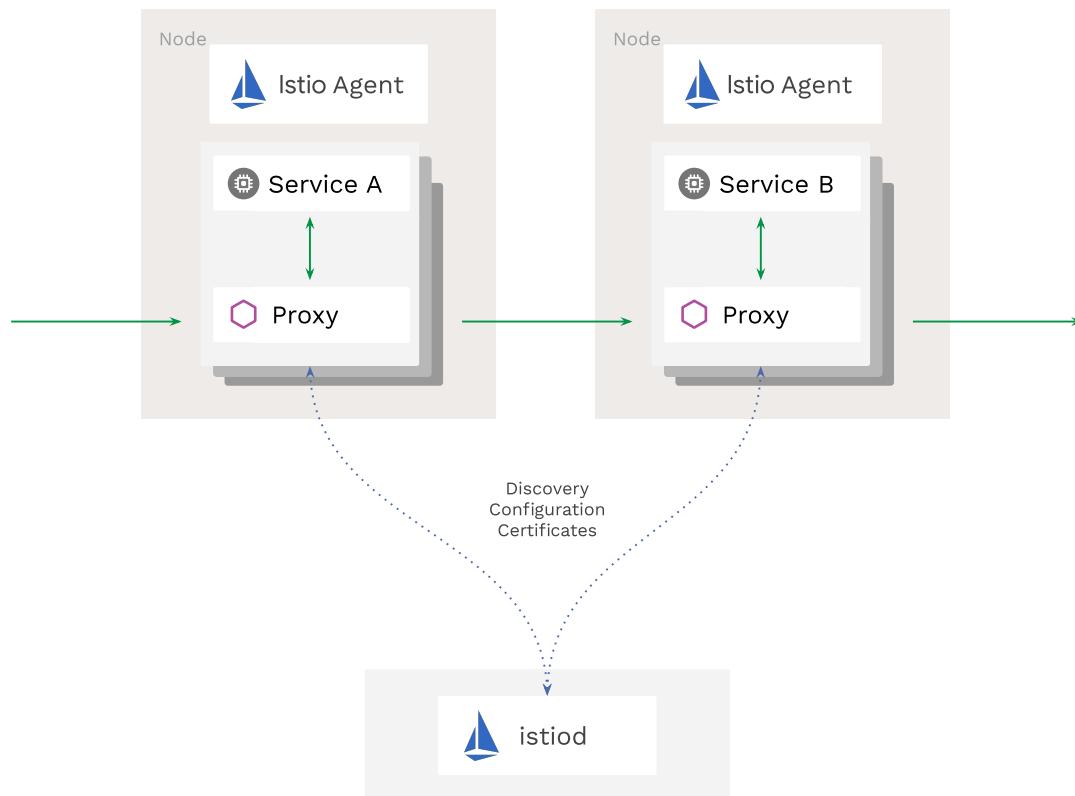


Istio Architecture v1.4.X



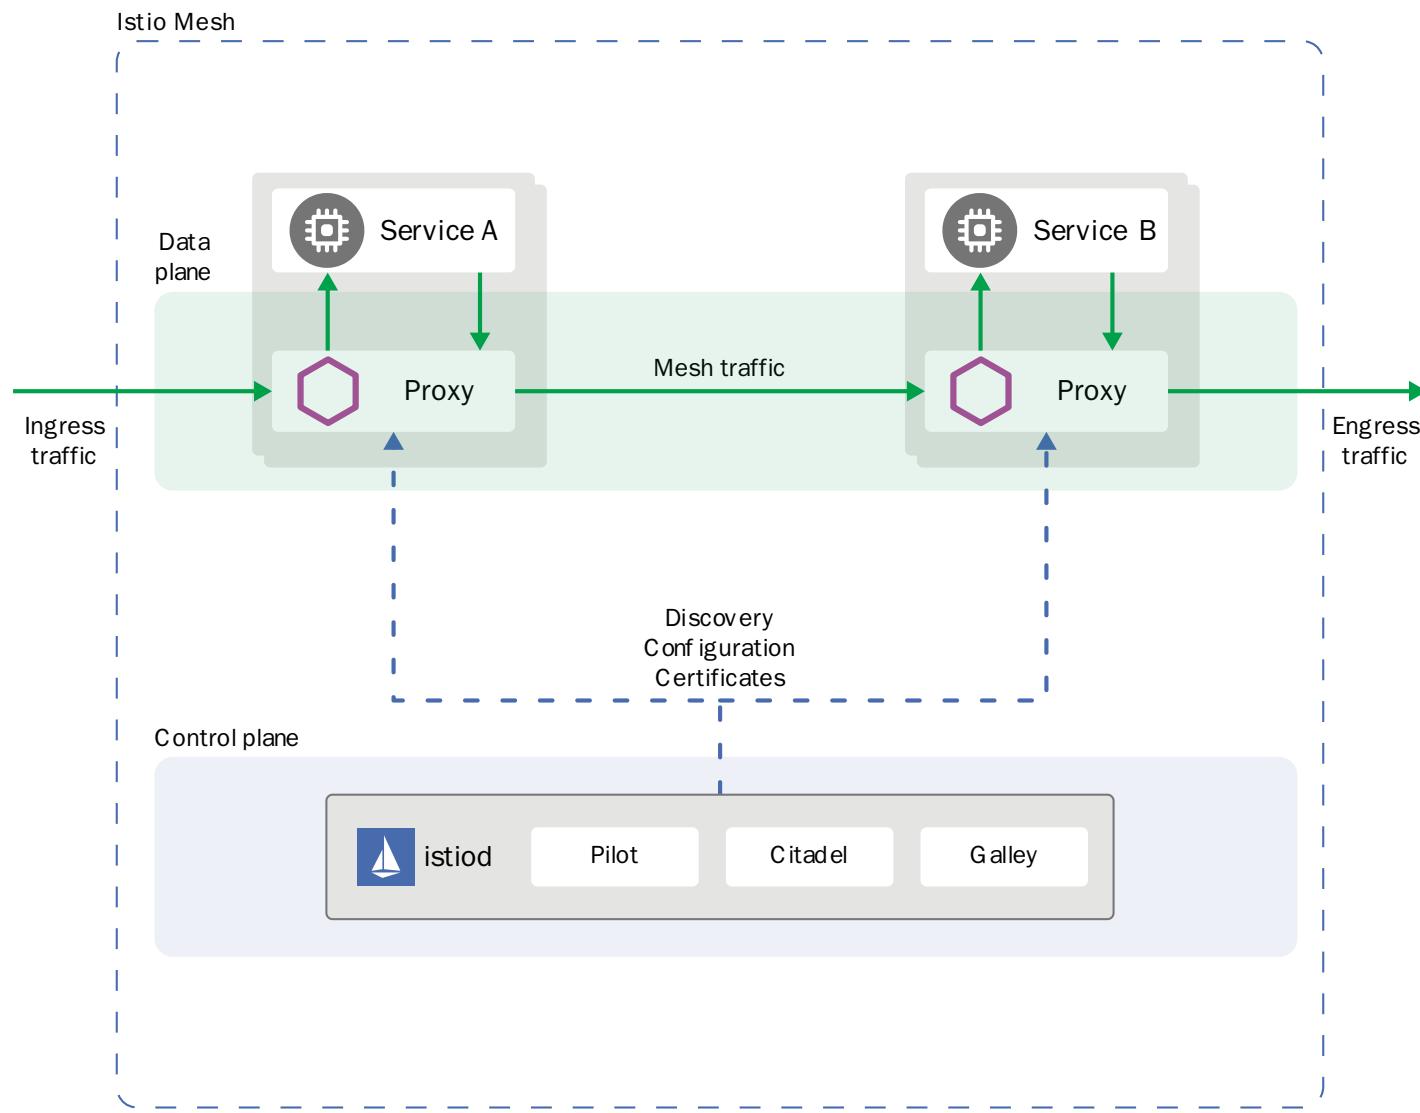


Istio Architecture v1.5.X+





Istio Architecture v1.5.X+



<https://istio.io/docs/ops/deployment/architecture/arch.svg>



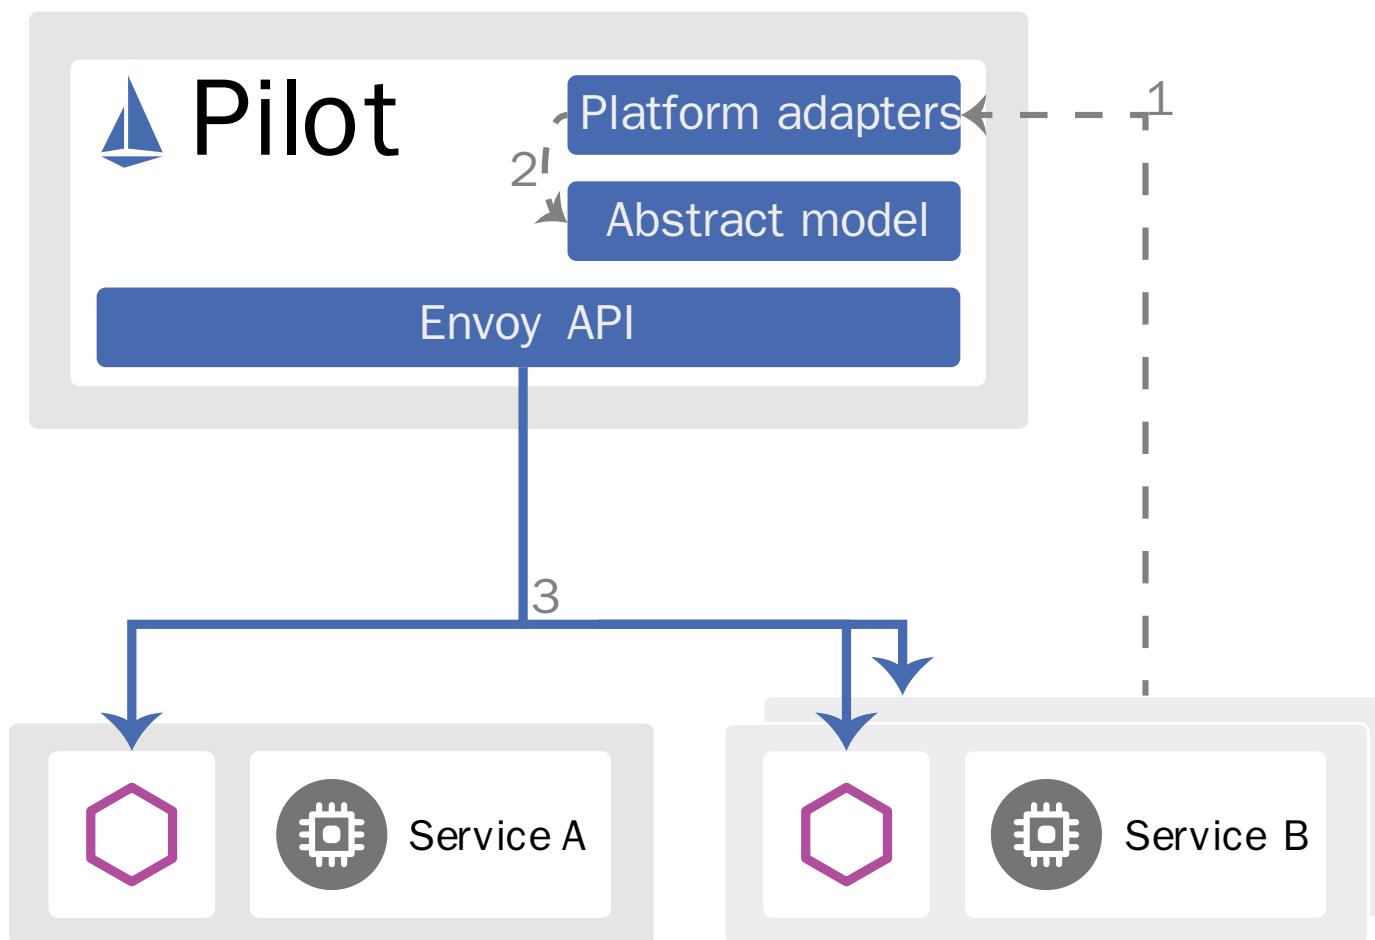
Control Plane / istiod

- Pilot
- Citadel
- Galley
- Sidecar Injection



Pilot

Serves xDS APIs to configure service discovery, traffic management capabilities and security for the Envoy sidecars

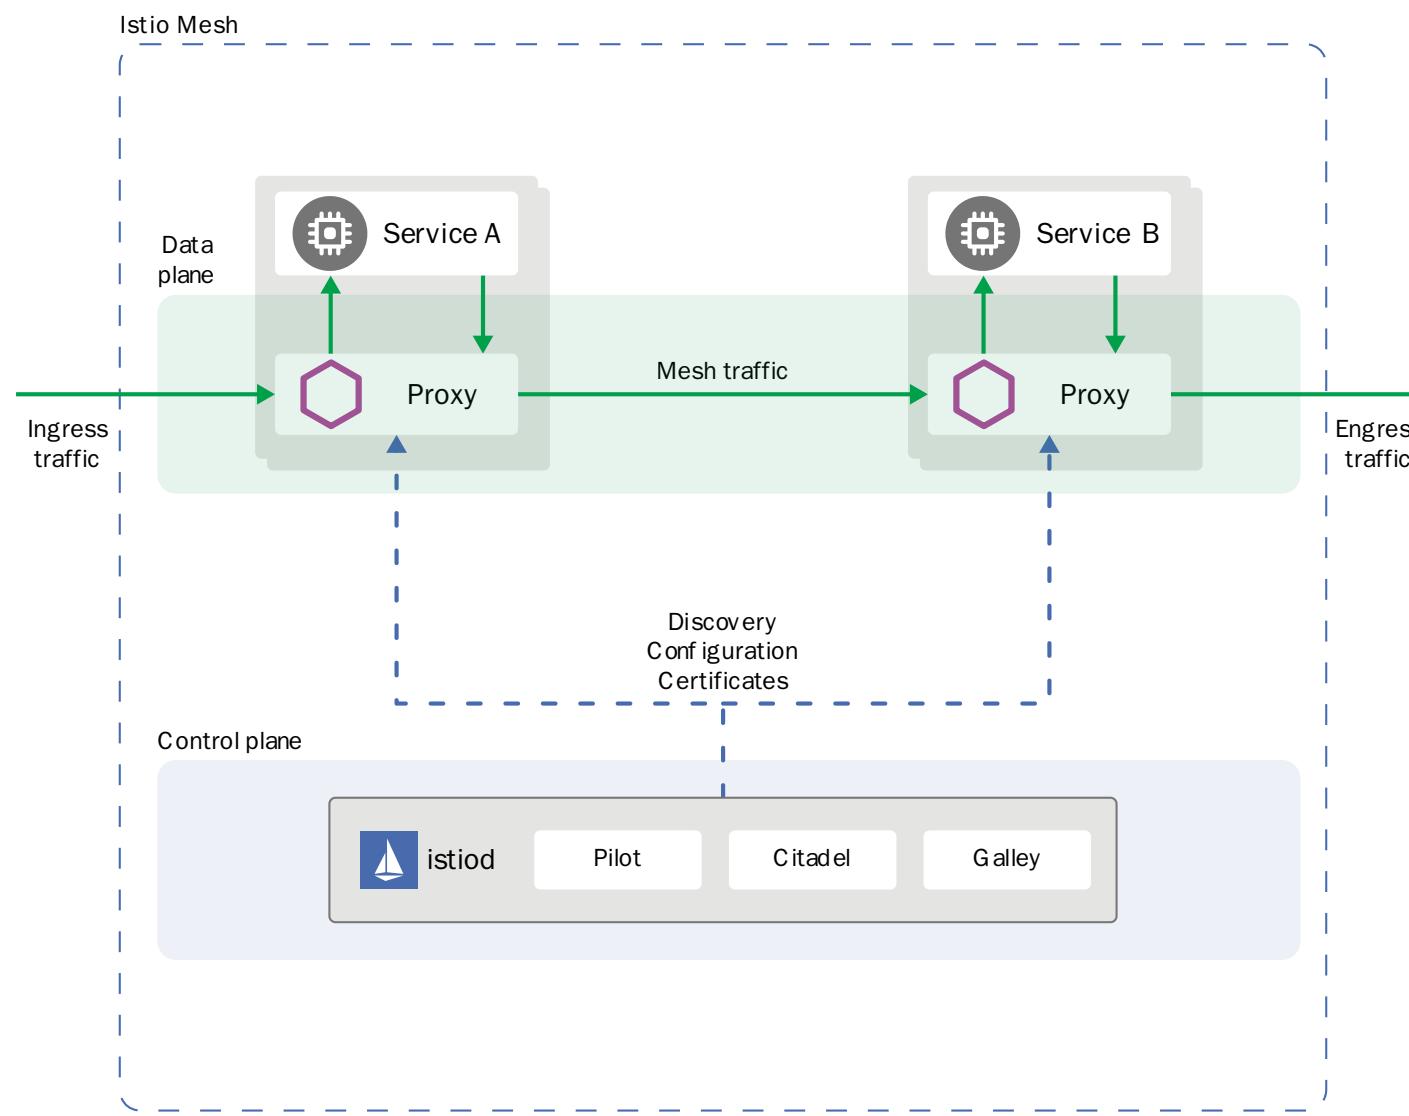


<https://istio.io/latest/docs/ops/deployment/architecture/discovery.svg>



Citadel

Enables strong service-to-service and end-user authentication with built-in identity and credential management

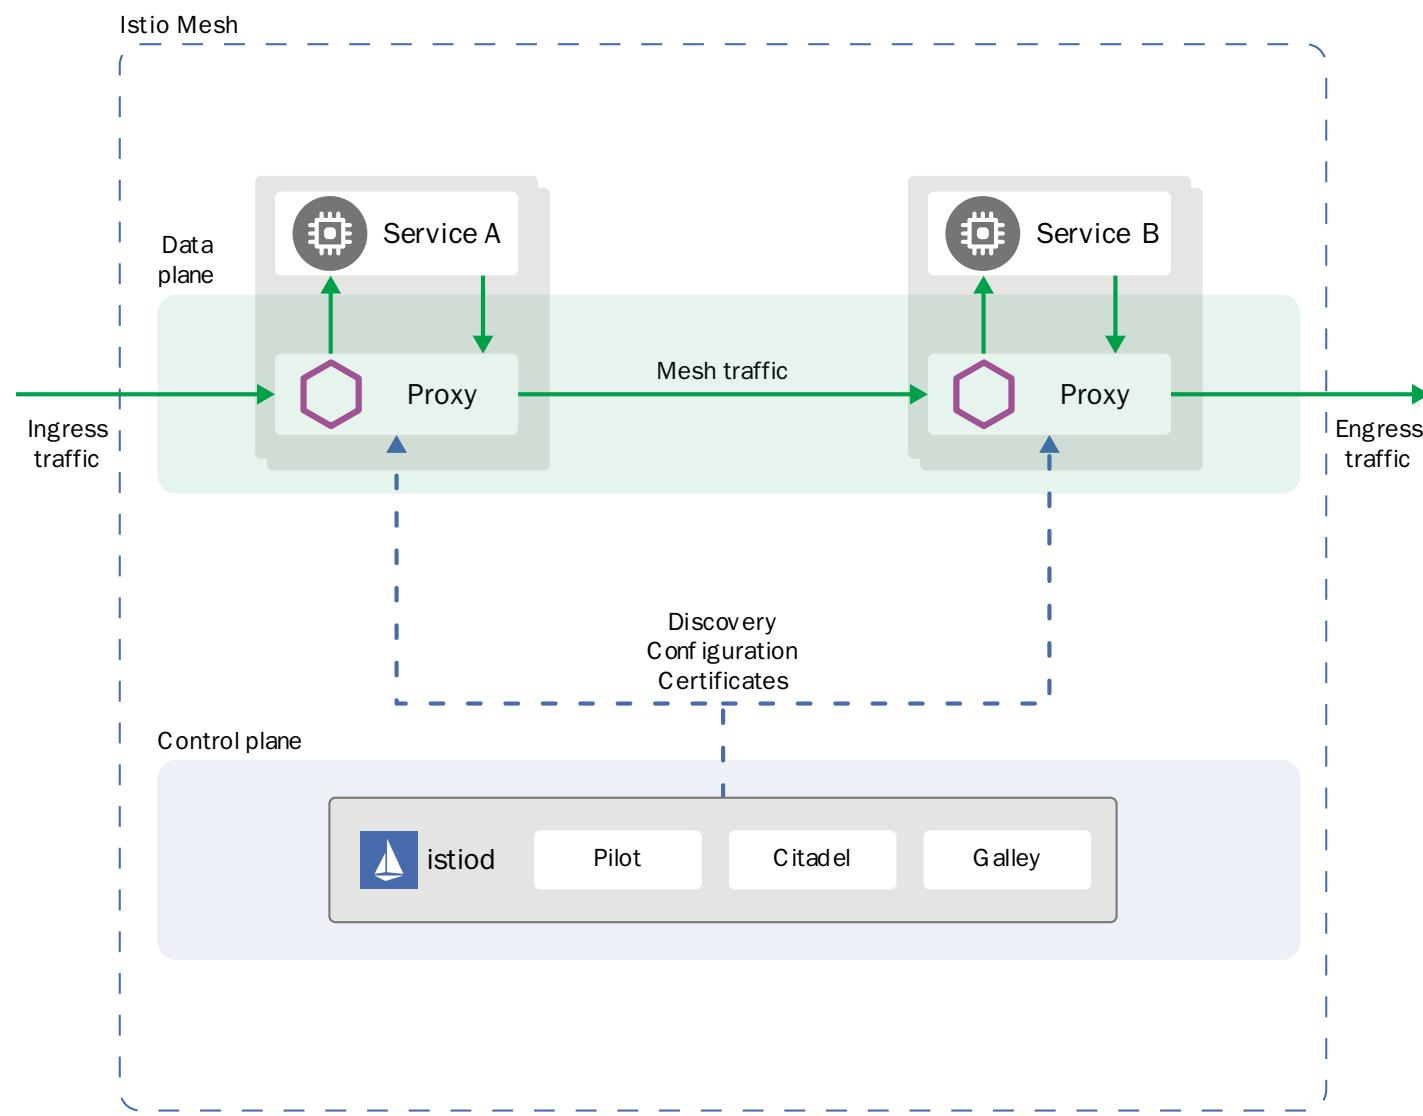


<https://istio.io/latest/docs/ops/deployment/architecture/arch.svg>



Galley

Converts platform specific CRDs into component configuration which can be consumed (e.g. by Pilot) over the Mesh Configuration Protocol



<https://istio.io/latest/docs/ops/deployment/architecture/arch.svg>



Sidecar Injection

Manual

```
istioctl kube-inject -f deployment.yaml | \  
    kubectl apply -f -
```

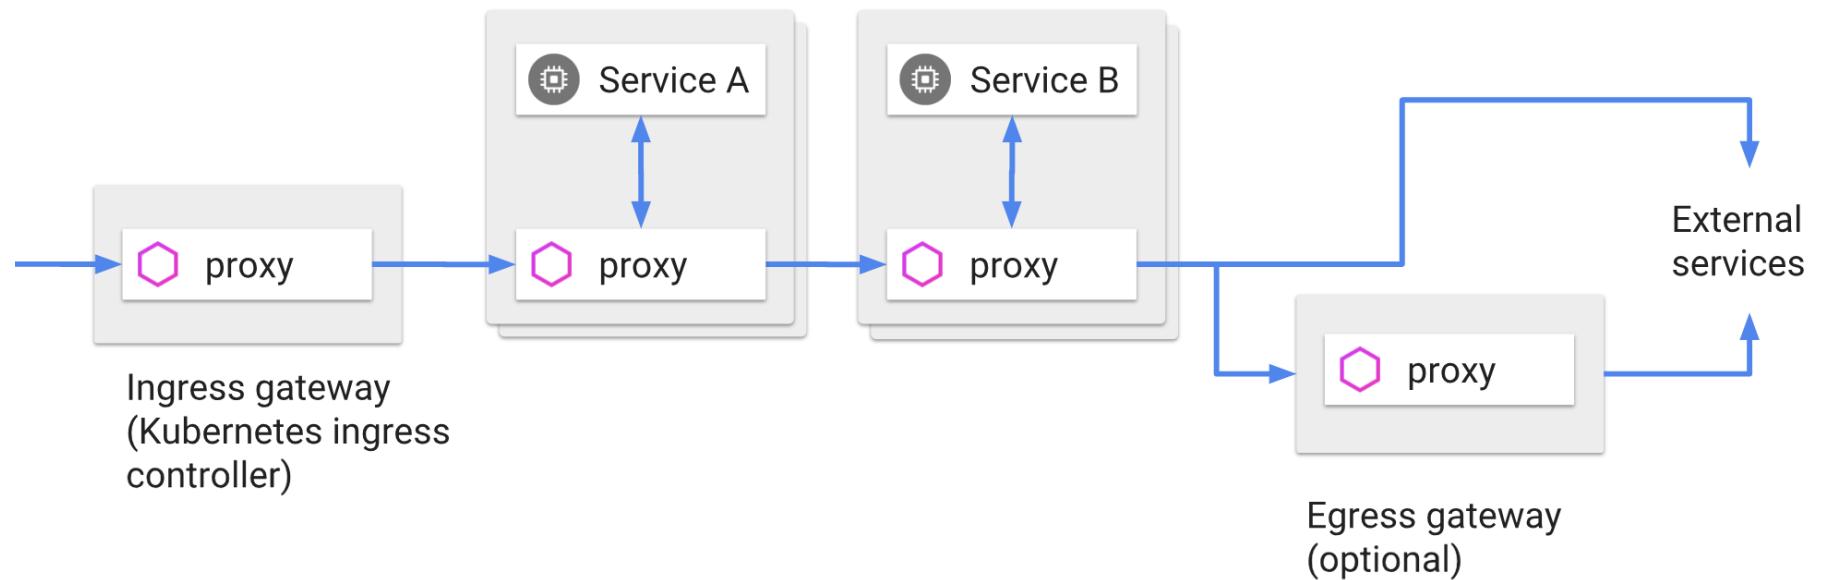
Automatic

```
apiVersion: admissionregistration.k8s.io/v1  
kind: MutatingWebhookConfiguration  
metadata:  
  name: istio-sidecar-injector  
webhooks:  
- rules:  
  - apiGroups:  
    - ""  
    apiVersions:  
    - v1  
    operations:  
    - CREATE  
    resources:  
    - pods  
    scope: '*'  
  ...
```



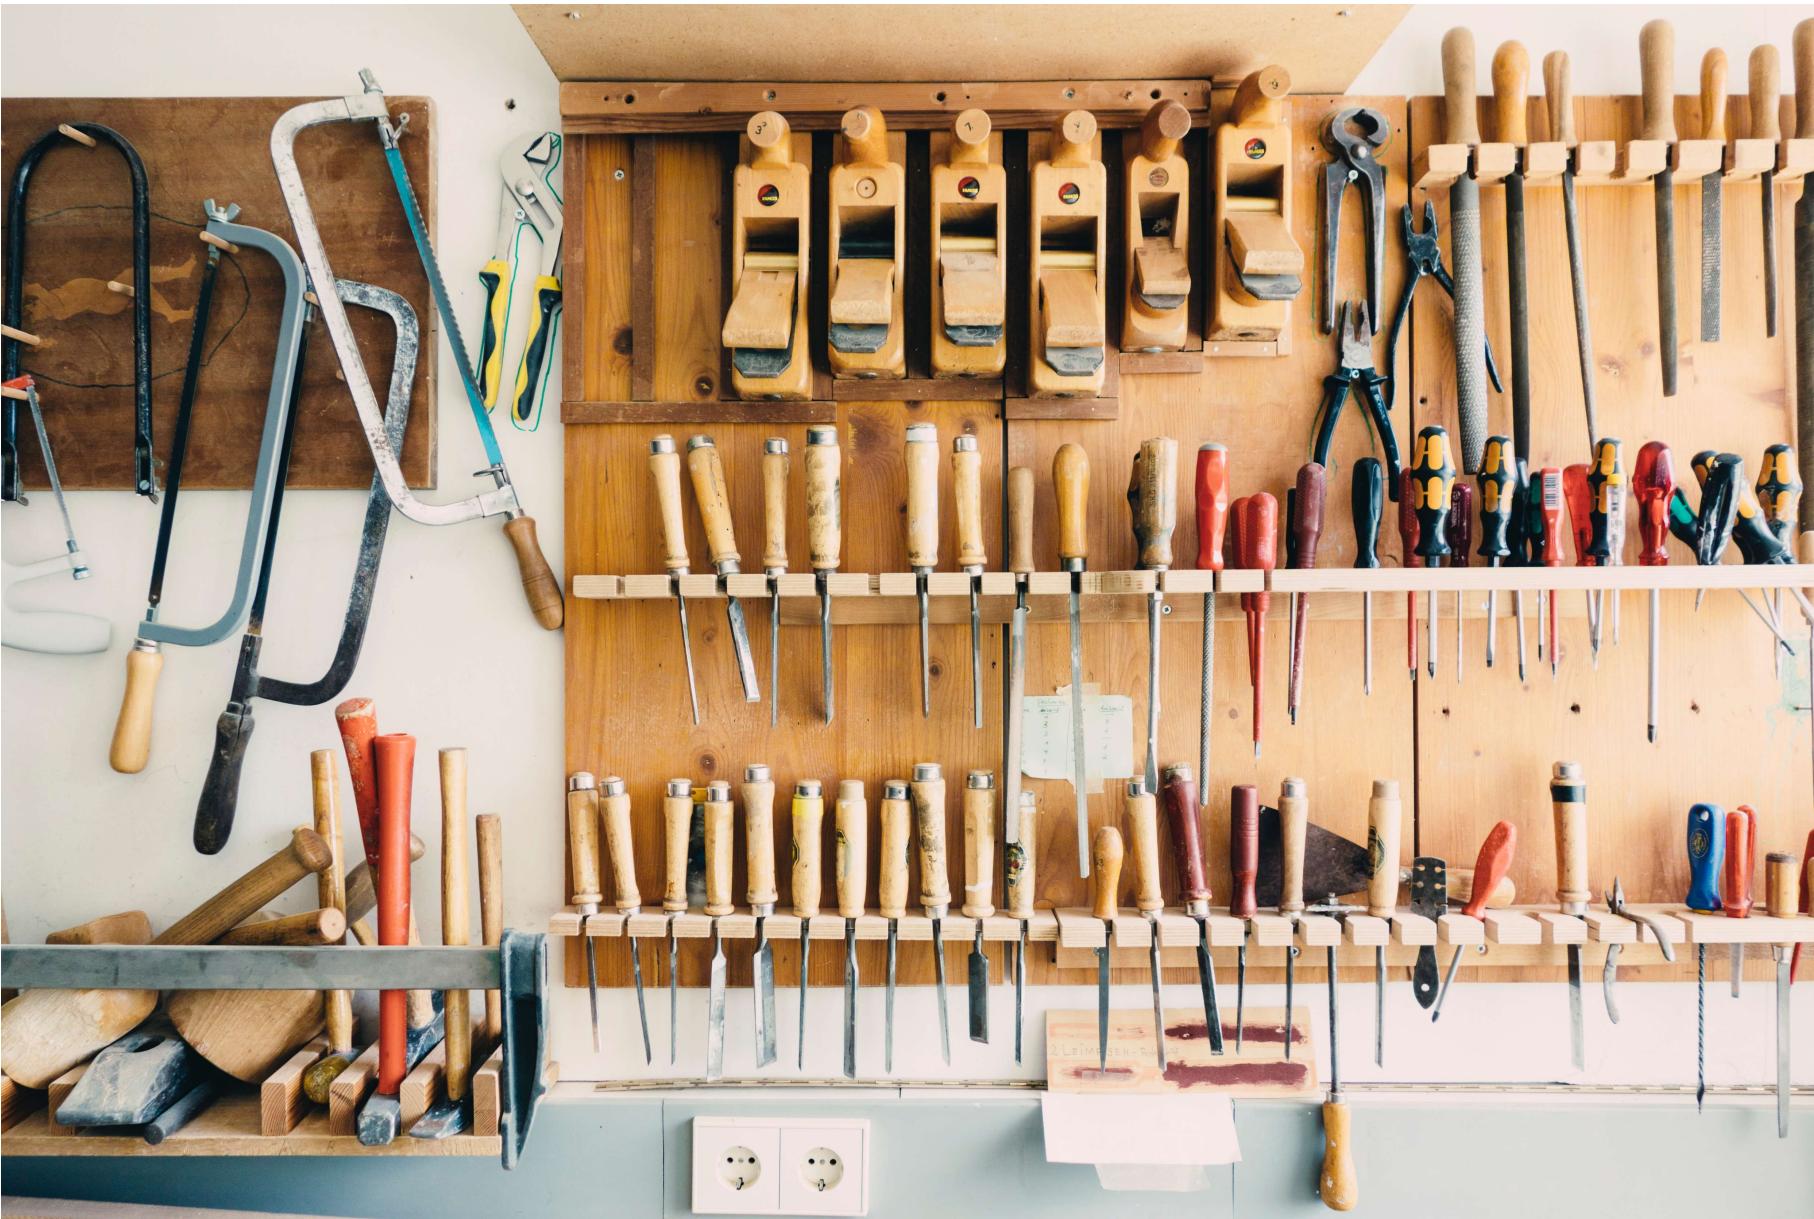
Data Plane

Envoy





Workshop: Deploy GKE Cluster and Install Istio





IstioOperator API

```
apiVersion: install.istio.io/v1alpha1
kind: IstioOperator
spec:
  profile: default
```

```
istioctl install -f istio-operator.yaml
```

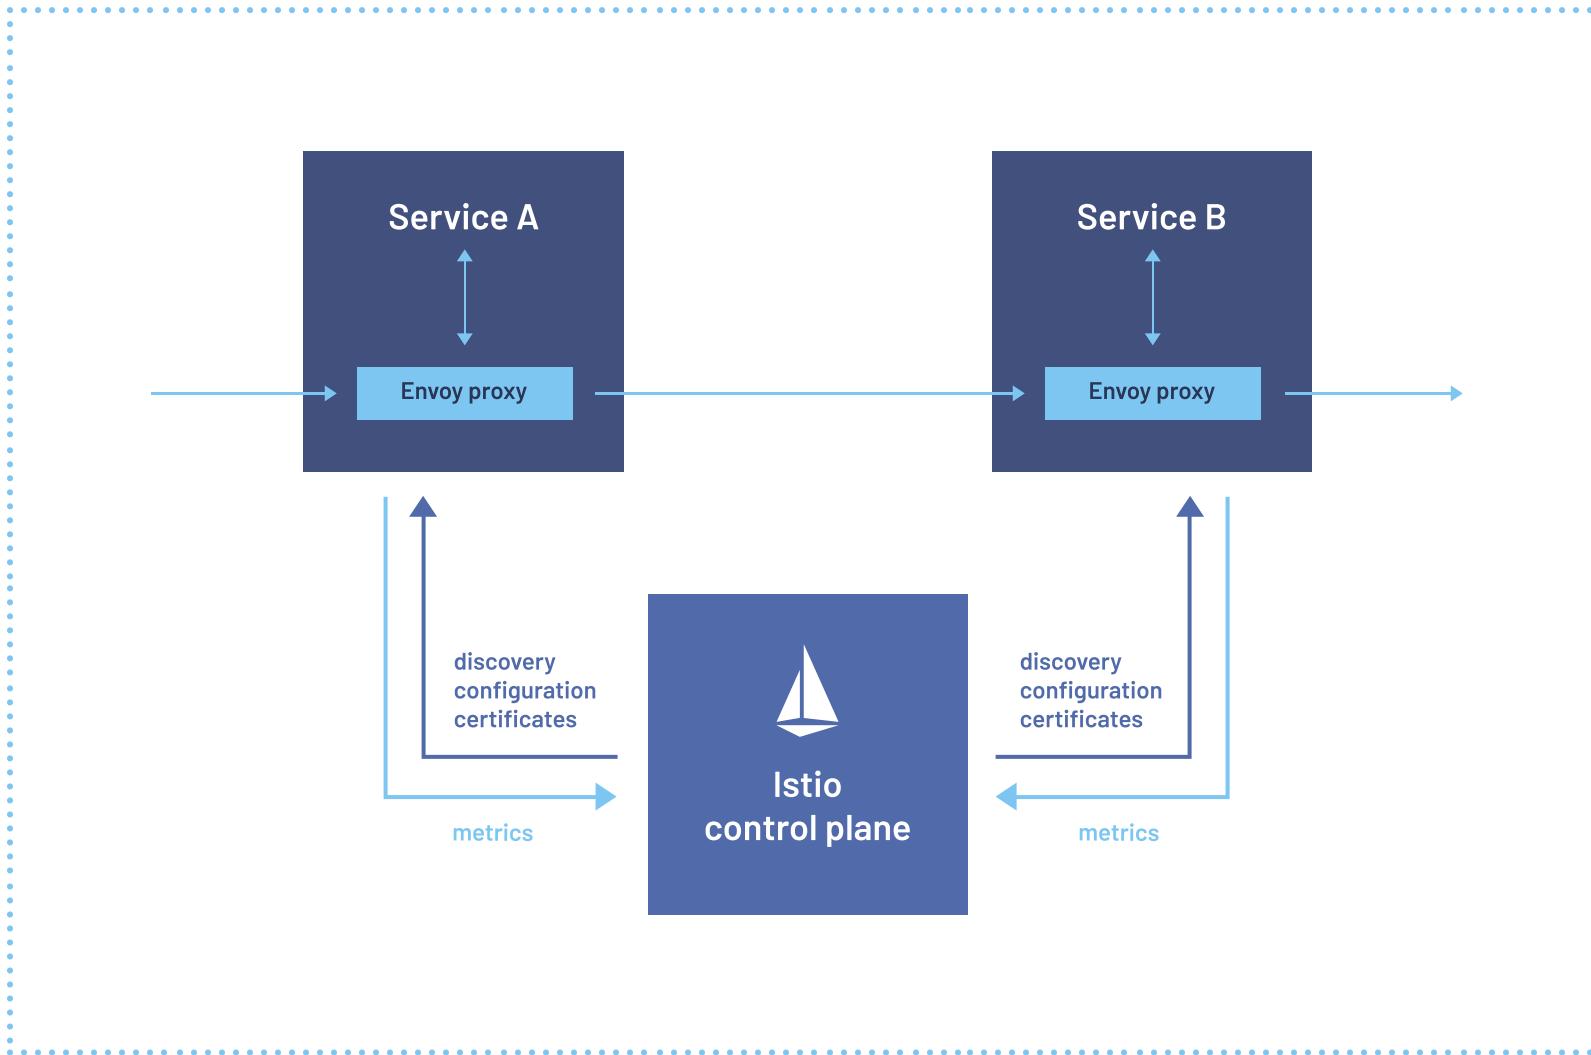


Istio Sidecar Deep Dive

The proxy that powers the Istio data plane



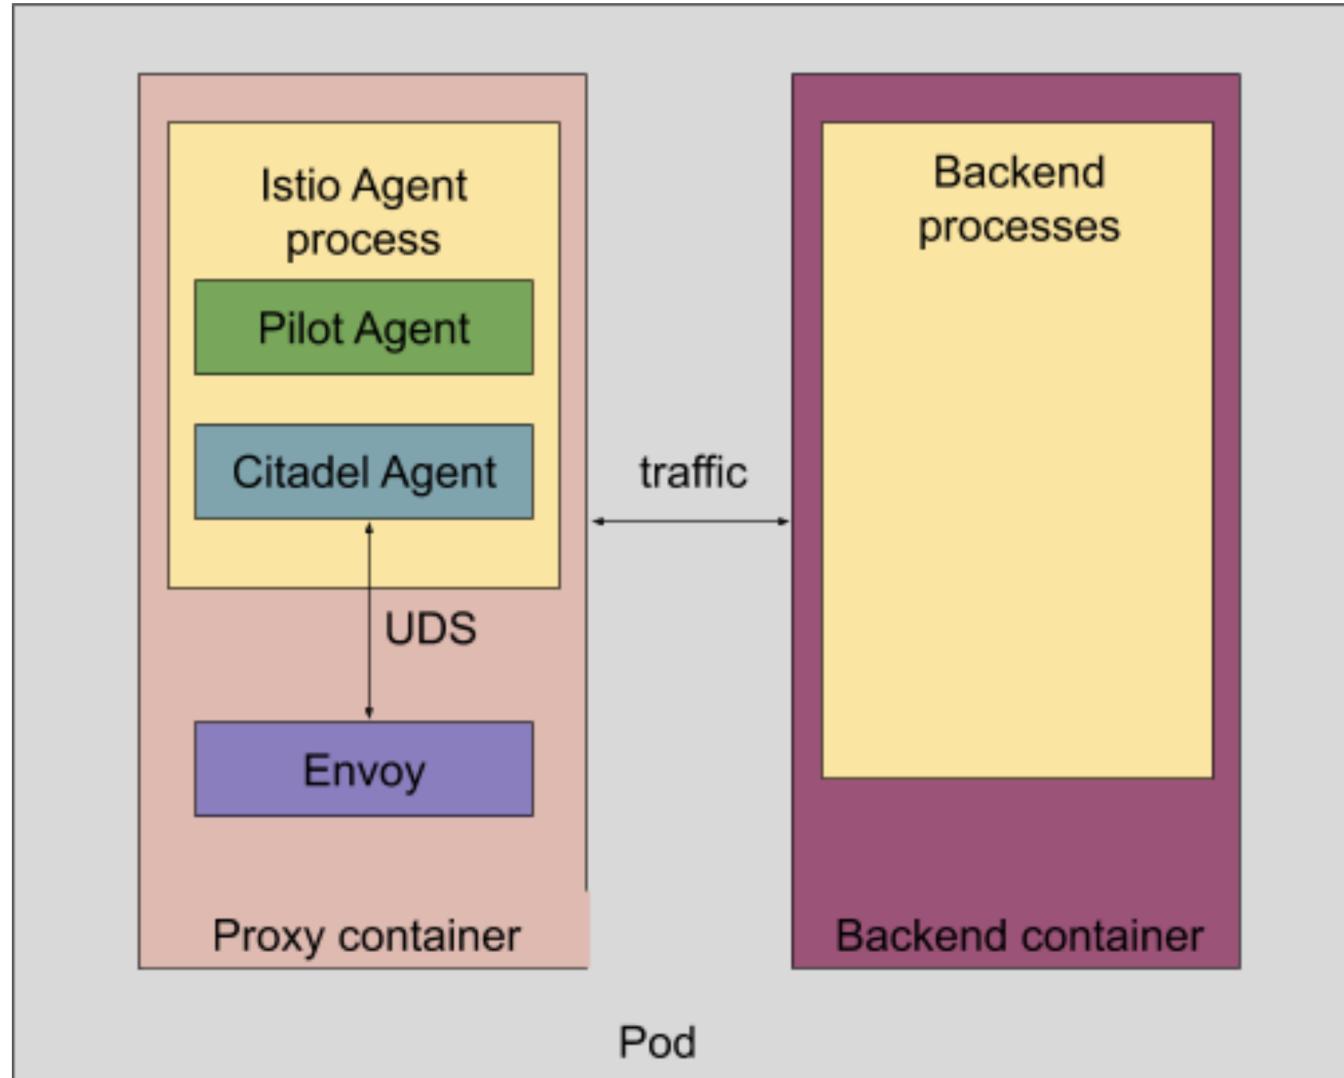
Envoy Proxy



<https://istio.io/latest/img/service-mesh.svg>

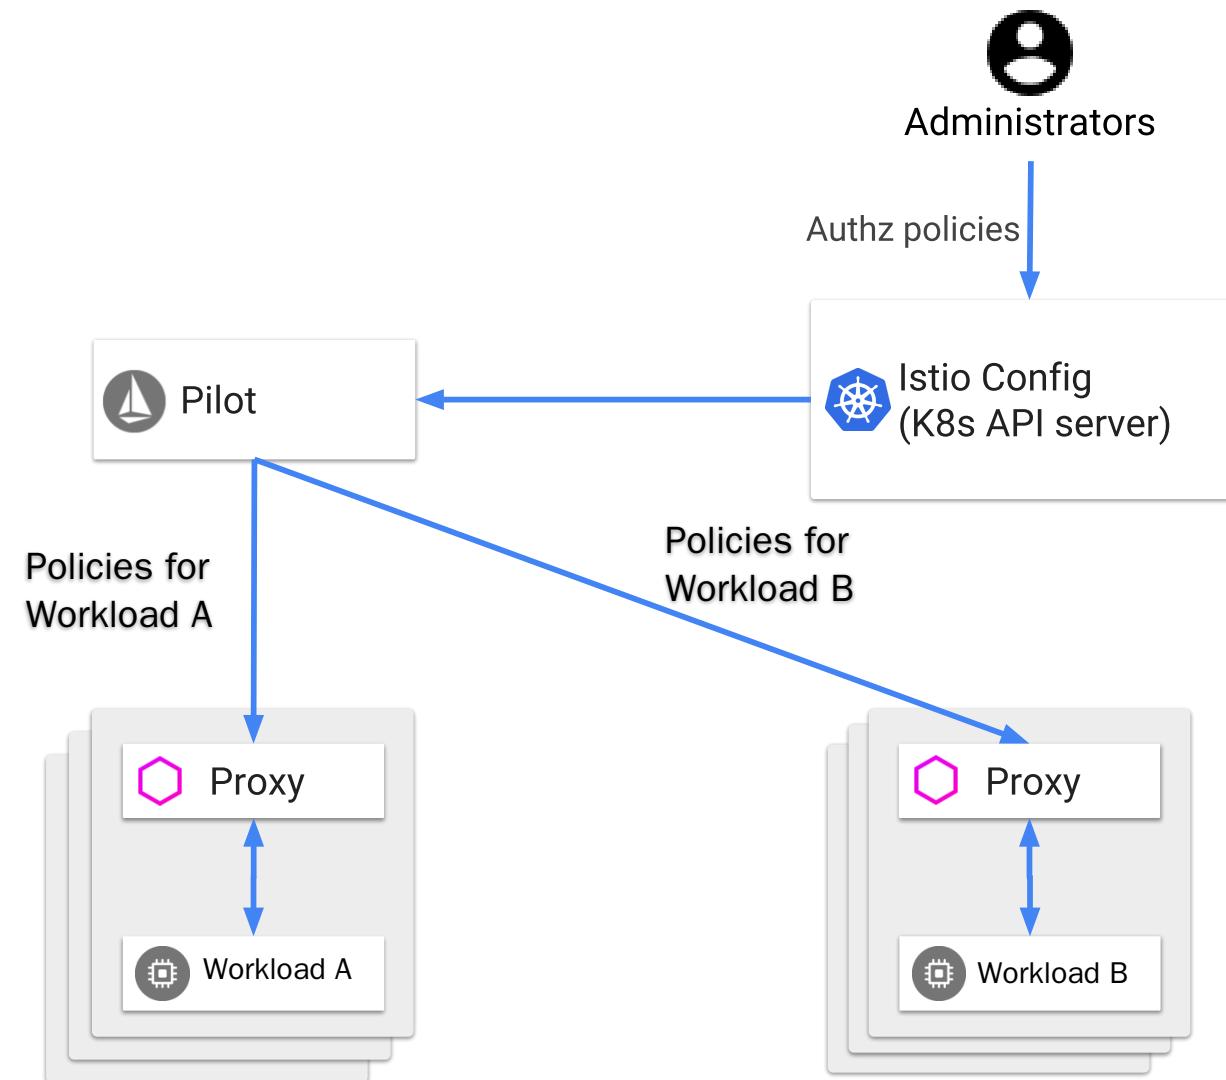


Istio Agent



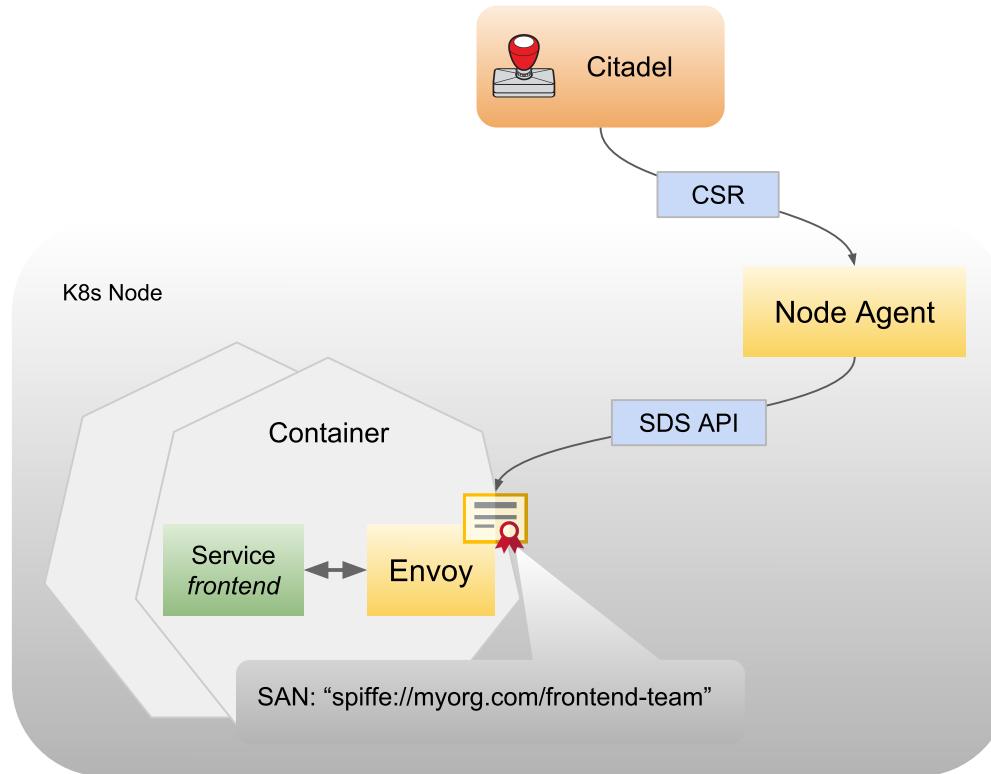


Pilot





SDS (Secret Discovery Service) 1.4.x





SDS (Secret Discovery Service) 1.5.x

```
$ ls -l /etc/istio/proxy/SDS
srw-rw-rw- 1 istio-proxy istio-proxy 0 Apr 13 19:58 /etc/istio/proxy/SDS
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  ...
spec:
  containers:
    - name: istio-proxy
      image: docker.io/istio/proxyv2:1.5.1
      args:
        - proxy
        - sidecar
        - --configPath
        - /etc/istio/proxy
        ...
      volumeMounts:
        - name: istio-envoy
          mountPath: /etc/istio/proxy
        ...
  volumes:
    - name: istio-envoy
      emptyDir:
        medium: Memory
      ...
```



PodSecurityPolicies

- Required by `istio-init` so that it's permitted to modify `IPTables` rules (unless using Istio CNI)
- ~~Required when using SDS to read the UNIX domain socket created by the node agent~~ Unnecessary in 1.5.x



Envoy Sidecar Injection

```
kubectl get mutatingwebhookconfiguration istio-sidecar-injector -o yaml
```

```
...
name: sidecar-injector.istio.io
namespaceSelector:
  matchLabels:
    istio-injection: enabled
rules:
- apiGroups:
  - ""
  apiVersions:
  - v1
  operations:
  - CREATE
  resources:
  - pods
  scope: ! * !
...
...
```



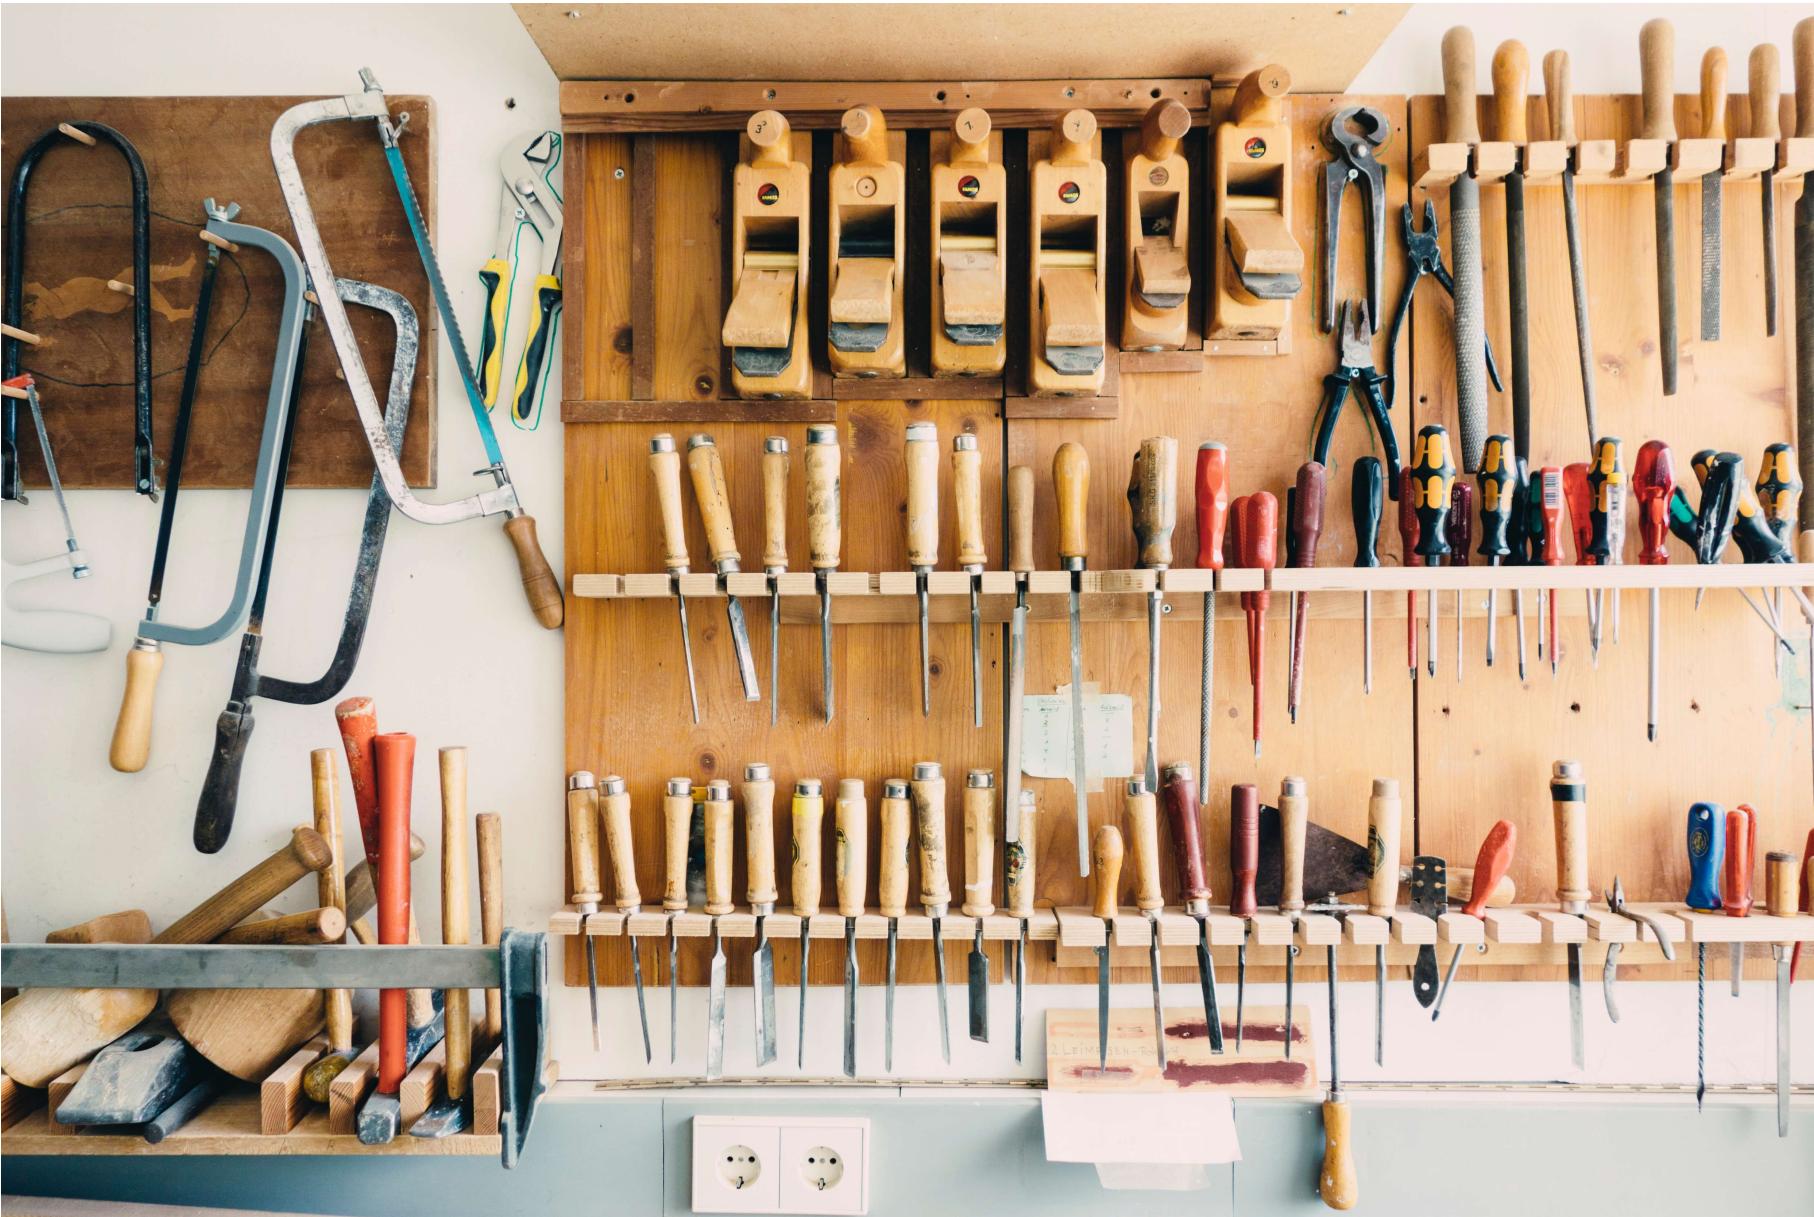
Envoy Sidecar Injection

```
kubectl label namespace default istio-injection=enabled --overwrite
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  ...
spec:
  initContainers:
  - name: istio-init
    image: docker.io/istio/proxyv2:1.5.1
    command:
      - istio-iptables
      ...
  containers:
  - name: nginx
    image: nginx
    ...
  - name: istio-proxy
    image: docker.io/istio/proxyv2:1.5.1
    args:
      - proxy
      ...
  ...
```



Workshop - Envoy





Solution

```
kubectl label namespace kube-system name=kube-system  
kubectl label namespace istio-system name=istio-system
```

```
apiVersion: networking.k8s.io/v1  
kind: NetworkPolicy  
metadata:  
  name: productpage  
  namespace: netpol  
spec:  
  podSelector:  
    matchLabels:  
      app: productpage  
  policyTypes:  
  - Ingress  
  - Egress  
  ingress:  
  - from:  
    - namespaceSelector:  
        matchLabels:  
          name: istio-system  
    podSelector:  
      matchLabels:  
        app: istio-ingressgateway  
  ports:  
  - protocol: TCP  
    port: 9080  
  - from:  
    - namespaceSelector:  
        matchLabels:
```

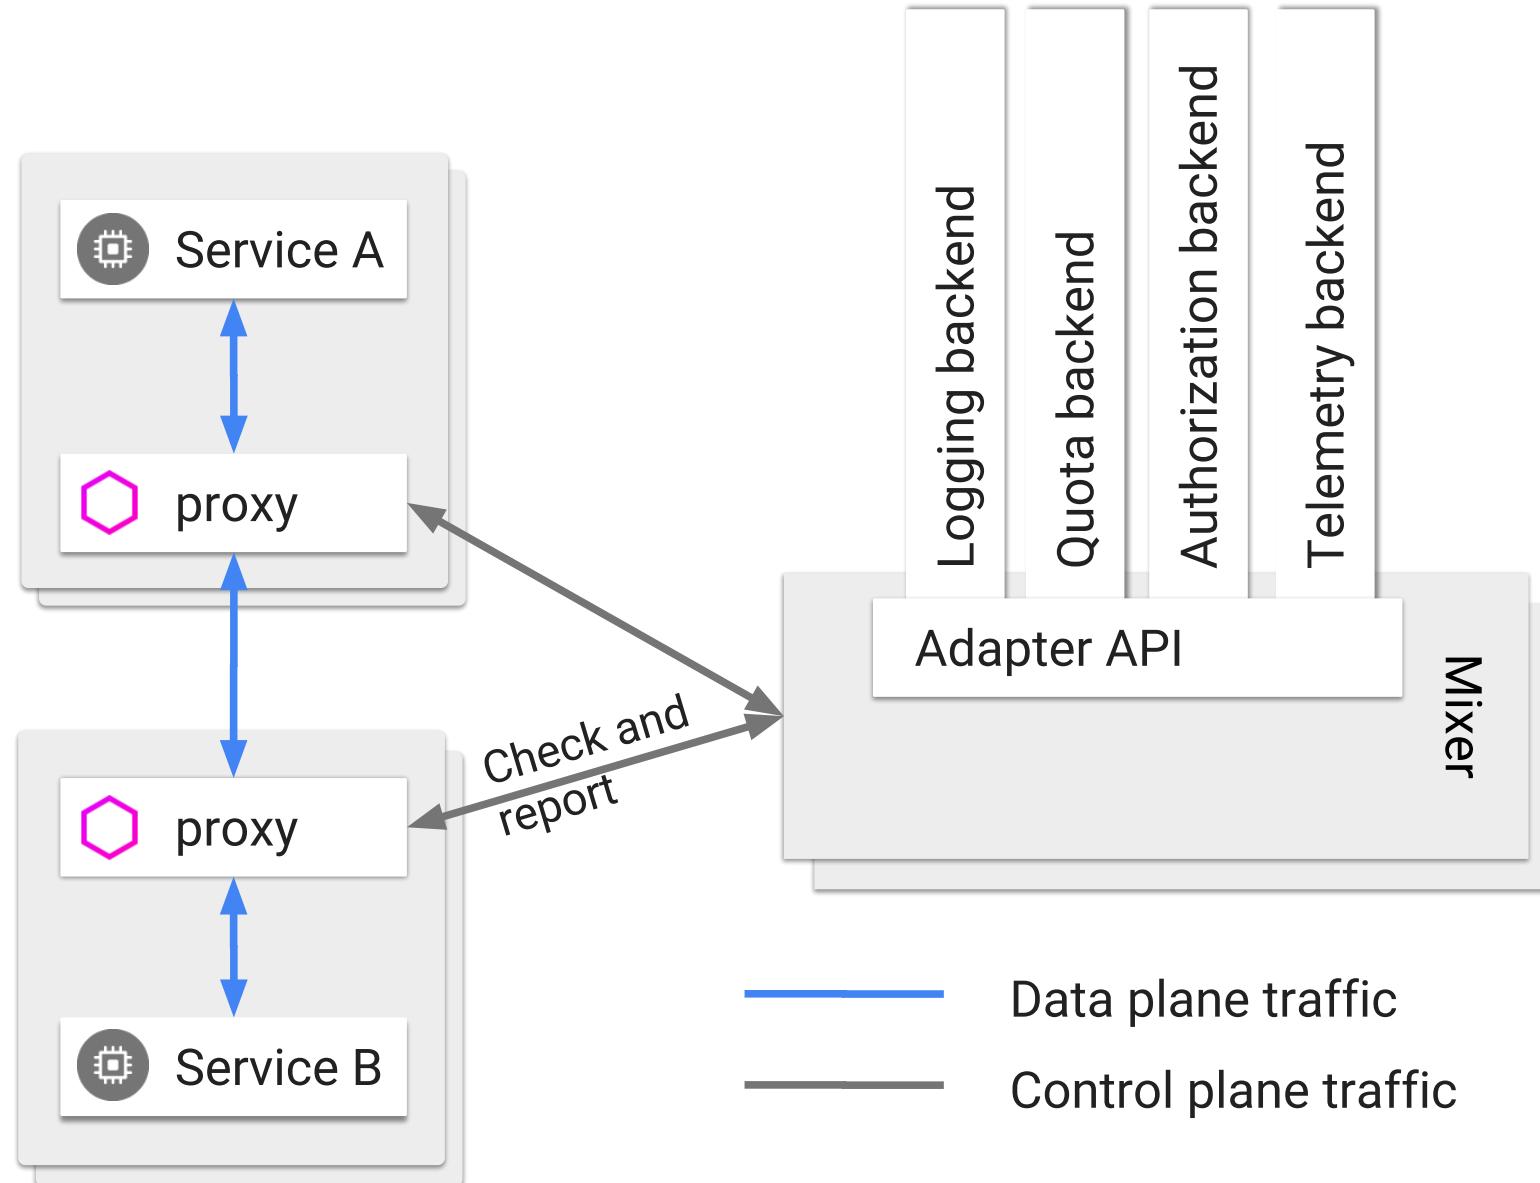


Mixer Deep Dive (deprecated)

Policy and telemetry

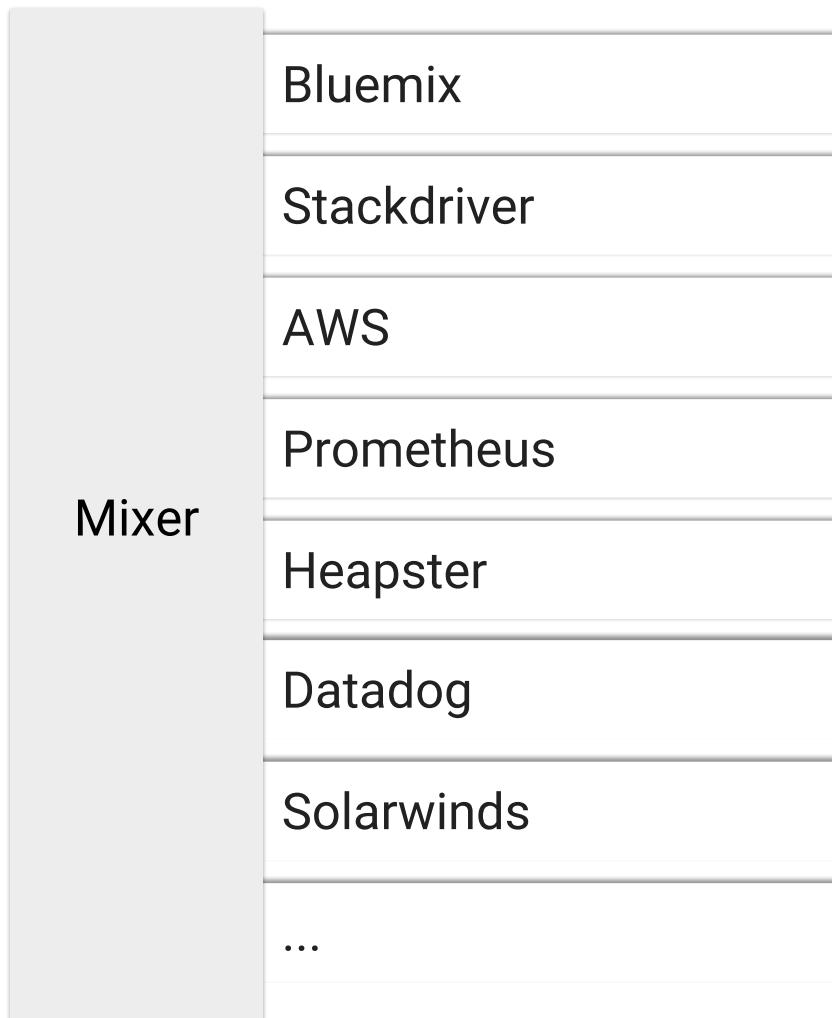


Mixer topology



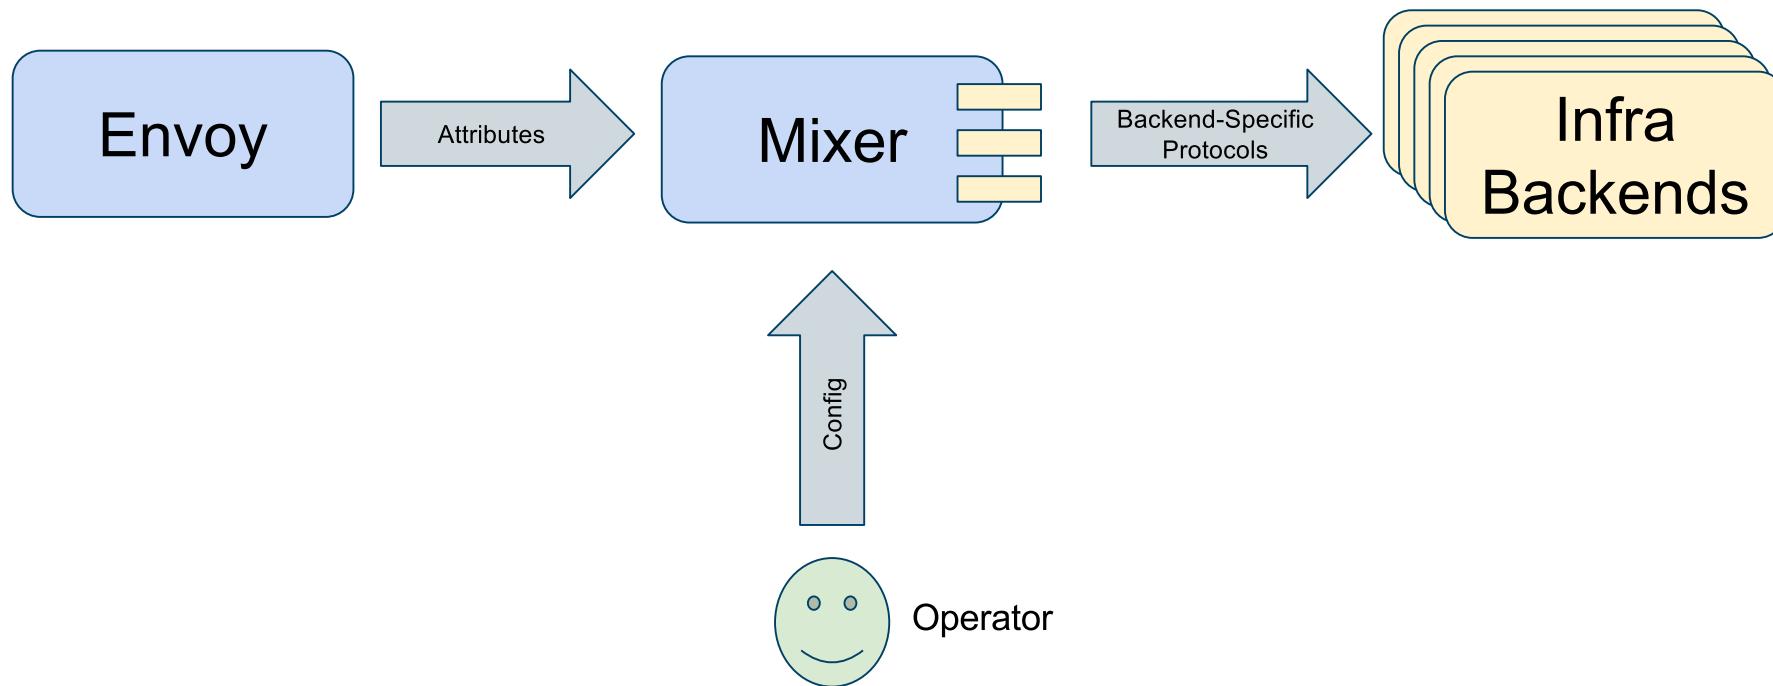


Mixer and its Adapters



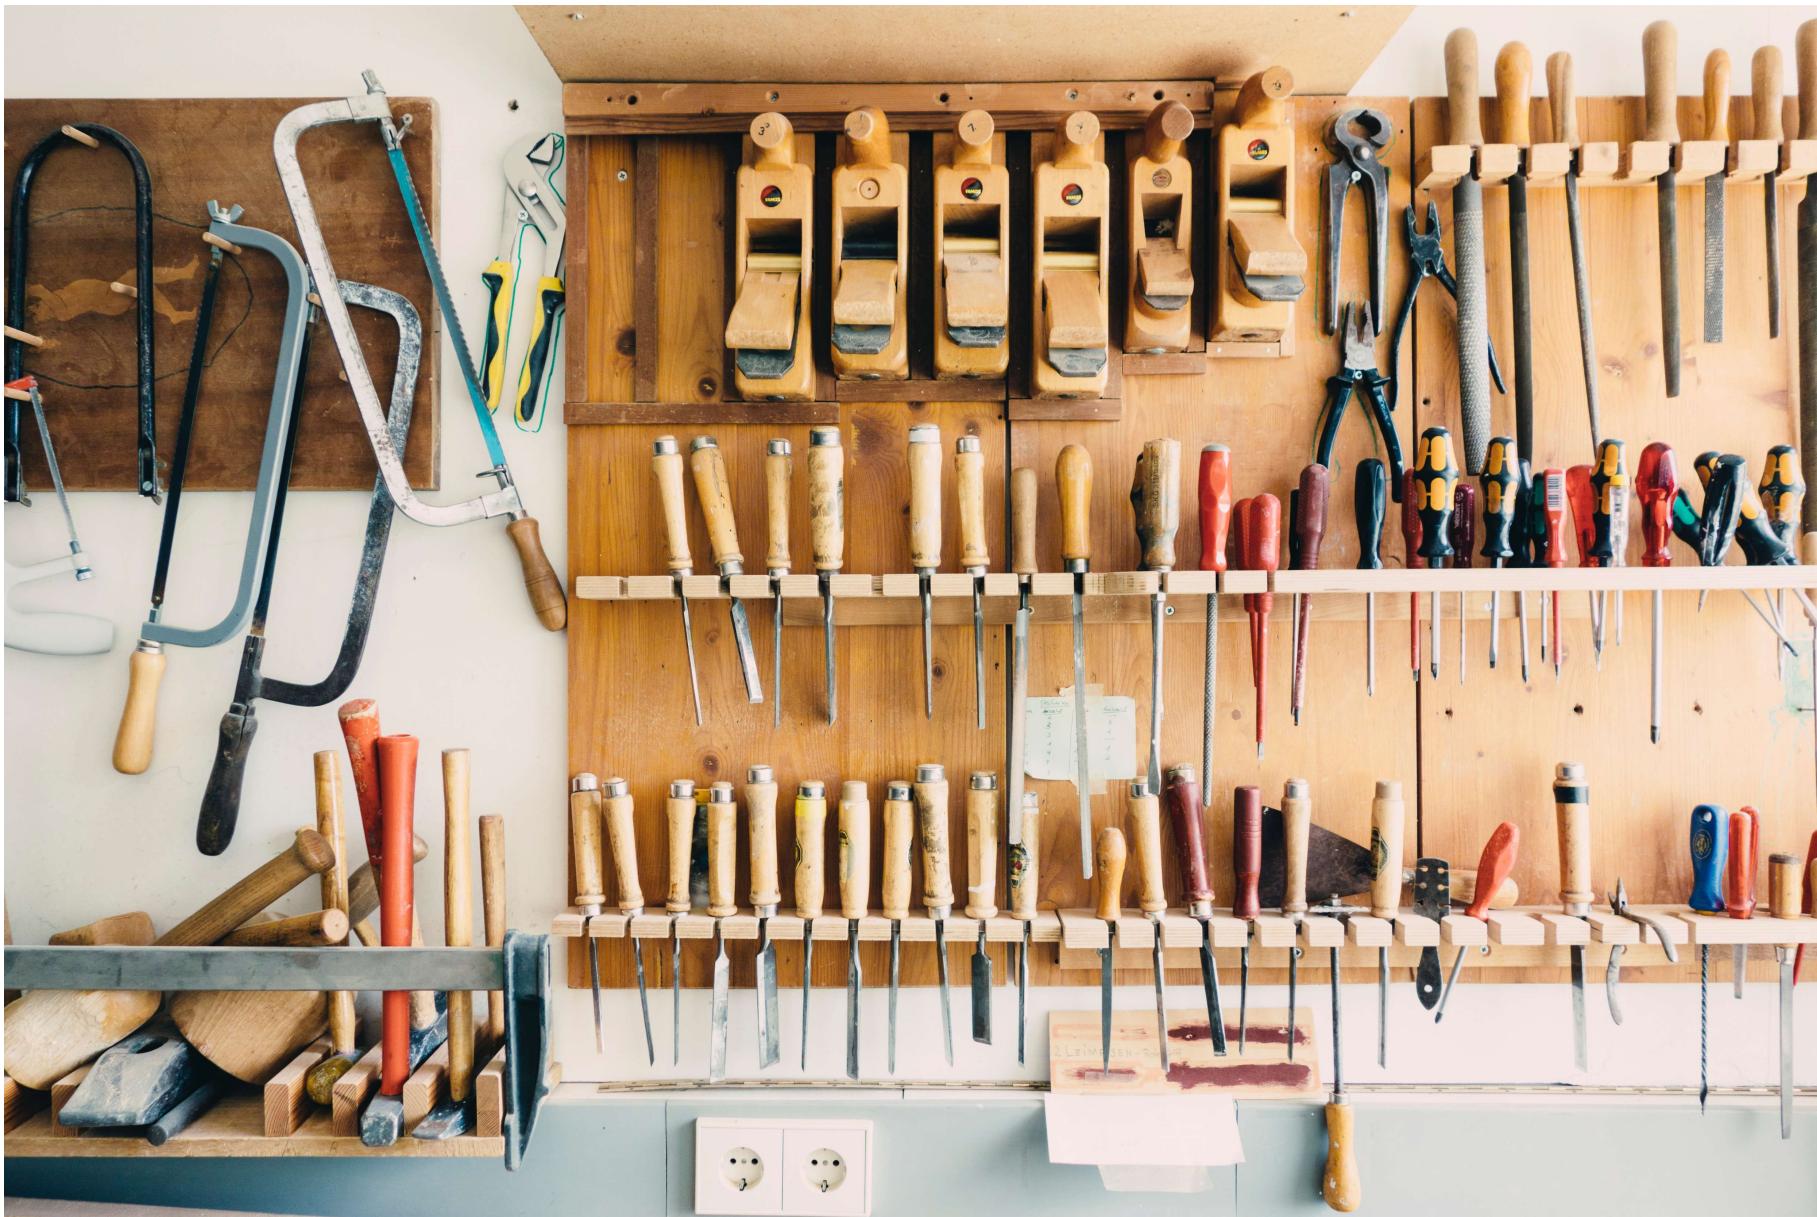


Attribute Machine





Workshop - Mixer





Solution

```
# Configuration for metric instance
apiVersion: config.istio.io/v1alpha2
kind: instance
metadata:
  name: request-count
  namespace: istio-system
spec:
  compiledTemplate: metric
  params:
    value: "1"
    dimensions:
      source_version: source.labels["version"] | "unknown"
      destination_service: destination.service.host | "unknown"
      destination_version: destination.labels["version"] | "unknown"
      response_code: response.code | 200
      monitored_resource_type: '"UNSPECIFIED"'
```



```
# Configuration for a Prometheus handler
apiVersion: config.istio.io/v1alpha2
kind: handler
metadata:
  name: prometheus-request-count
  namespace: istio-system
spec:
  compiledAdapter: prometheus
  params:
    metrics:
      - name: request_count # Prometheus metric name
        instance_name: request-count.instance.istio-system # Mixer instance name
        kind: COUNTER
        label_names:
          - source_version
          - destination_service
          - destination_version
          - response_code
```



```
# Rule to send metric instances to a Prometheus handler
apiVersion: config.istio.io/v1alpha2
kind: rule
metadata:
  name: request-count
  namespace: istio-system
spec:
  actions:
    - handler: prometheus-request-count
      instances: [ request-count ]
```



Extensibility



WebAssembly

WebAssembly is a sandboxing technology that defines a binary instruction format for a stack-based virtual machine.

Host architecture specific runtimes can be used to run WebAssembly modules portably.

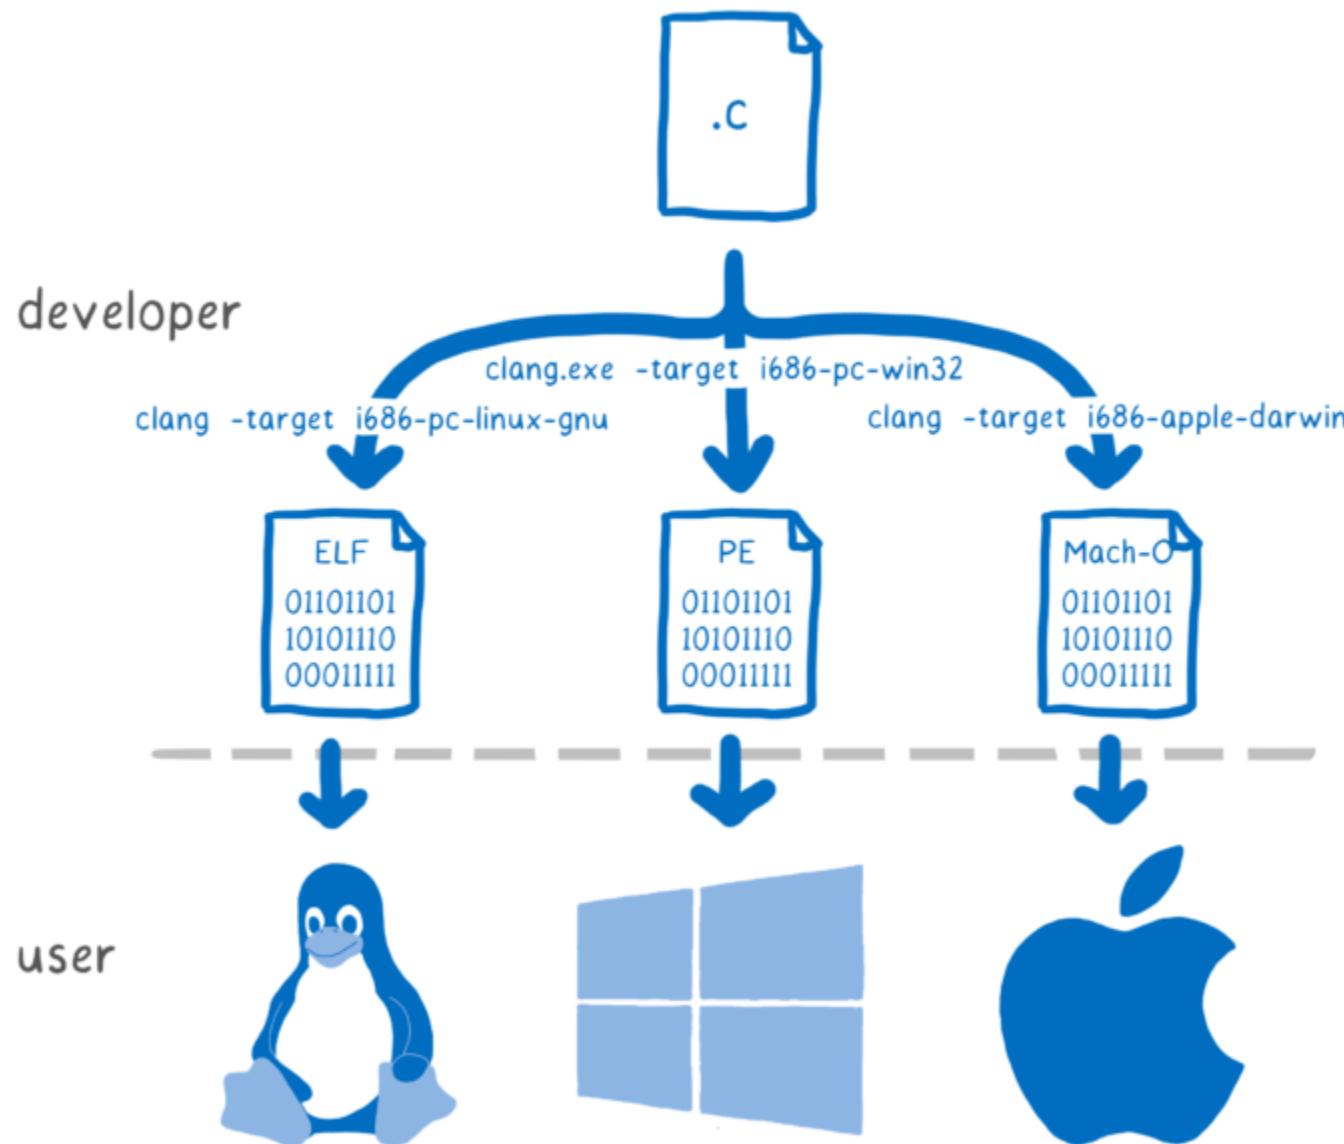
WebAssembly was originally designed to run high performance applications securely in the browser.

WebAssembly Goals:

- Fast, efficient, portable
- Readable and debuggable assembly language format
- Secure

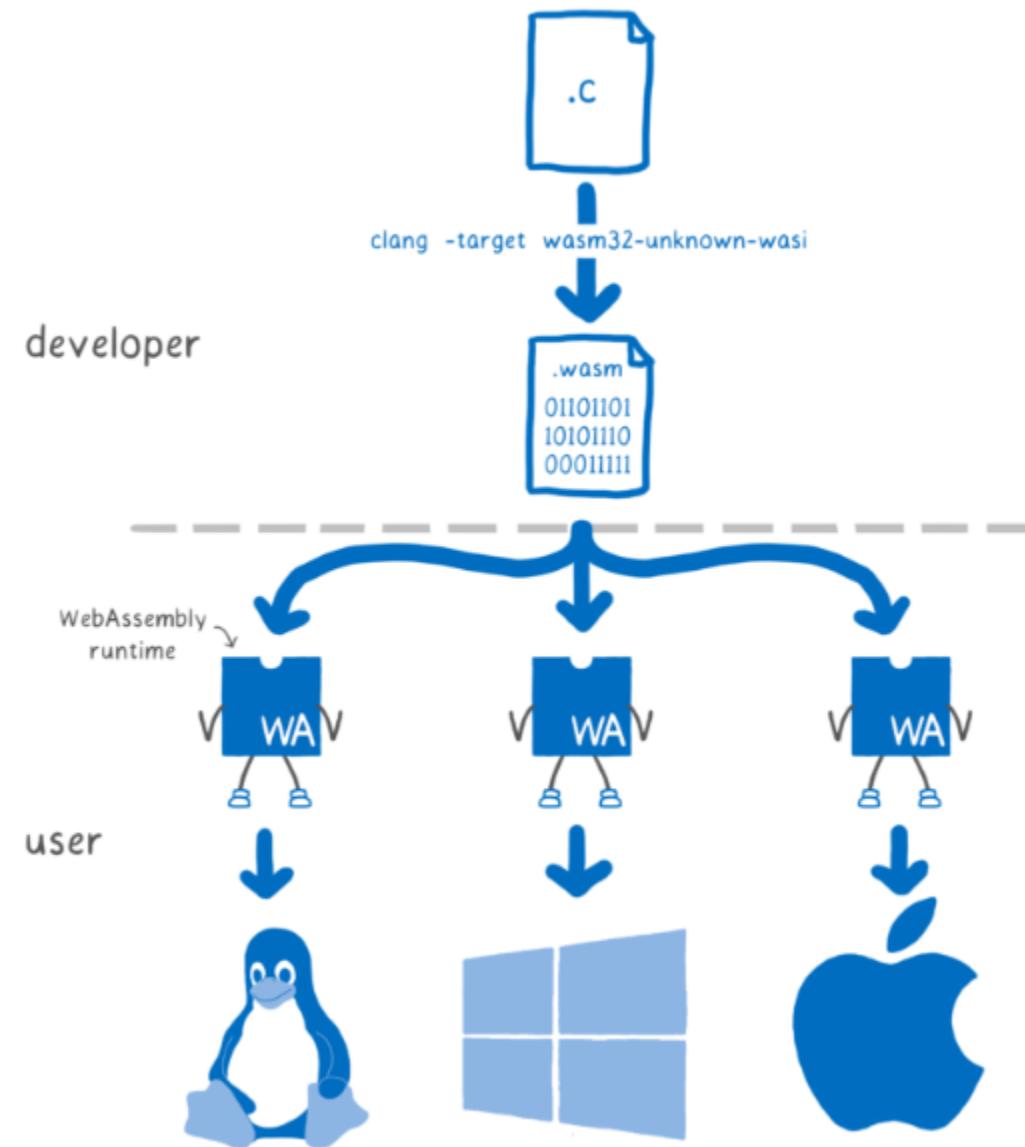


WebAssembly





WebAssembly



WebAssembly System Interface (WASI)

Defines an ABI for WebAssembly modules to interact with their host environment:

- Files and filesystems
- Sockets
- Clocks
- Random numbers

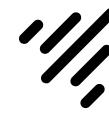


Extensibility

Each Envoy listener (usually programmed through xDS or the Listener Discover Service (LDS) specifically) defines a set of filters that sit in the data path, collectively forming a chain.

EnvoyFilter is an Istio resource that provides a mechanism for modifying these filters. Use cases include:

- Interpreting custom protocols
- Generating telemetry
- Performing RBAC



Extensibility

```
apiVersion: networking.istio.io/v1alpha3
kind: EnvoyFilter
metadata:
  name: reviews-lua
  namespace: bookinfo
spec:
  workloadSelector:
    labels:
      app: reviews
  configPatches:
    # The first patch adds the lua filter to the listener/http connection
    - applyTo: HTTP_FILTER
      match:
        context: SIDECAR_INBOUND
      listener:
        portNumber: 8080
        filterChain:
          filter:
            name: "envoy.http_connection_manager"
            subFilter:
              name: "envoy.router"
      patch:
        operation: INSERT_BEFORE
        value: # lua filter specification
          name: envoy.lua
          typed_config:
            "@type": "type.googleapis.com/envoy.config.filter.http.lua.v2."
            inlineCode: |
              ...
            
```



Extensibility

Envoy supports a number of in-built filters, including `envoy.filters.http.wasm` which supports running WASM modules. This filter can be used to dynamically extend Envoy.

Proxy-Wasm is a proxy specific ABI allowing these modules to interact with their host proxy. This replaces Mixer as the primary extension mechanism in Istio.

Istio 1.6 will provide a uniform configuration API for Proxy-Wasm plugins.

Proxy-Wasm will eventually become part of WASI.



Extensibility

WebAssembly sandbox goals:

- **Efficiency** - An extension adds low latency, CPU, and memory overhead.
- **Function** - An extension can enforce policy, collect telemetry, and perform payload mutations.
- **Isolation** - A programming error or crash in one plugin does not affect other plugins.
- **Configuration** - The plugins are configured using an API that is consistent with other Istio APIs. An extension can be configured dynamically.
- **Operator** - An extension can be canaried and deployed as log-only, fail-open or fail-close.
- **Extension developer** - The plugin can be written in several programming languages.



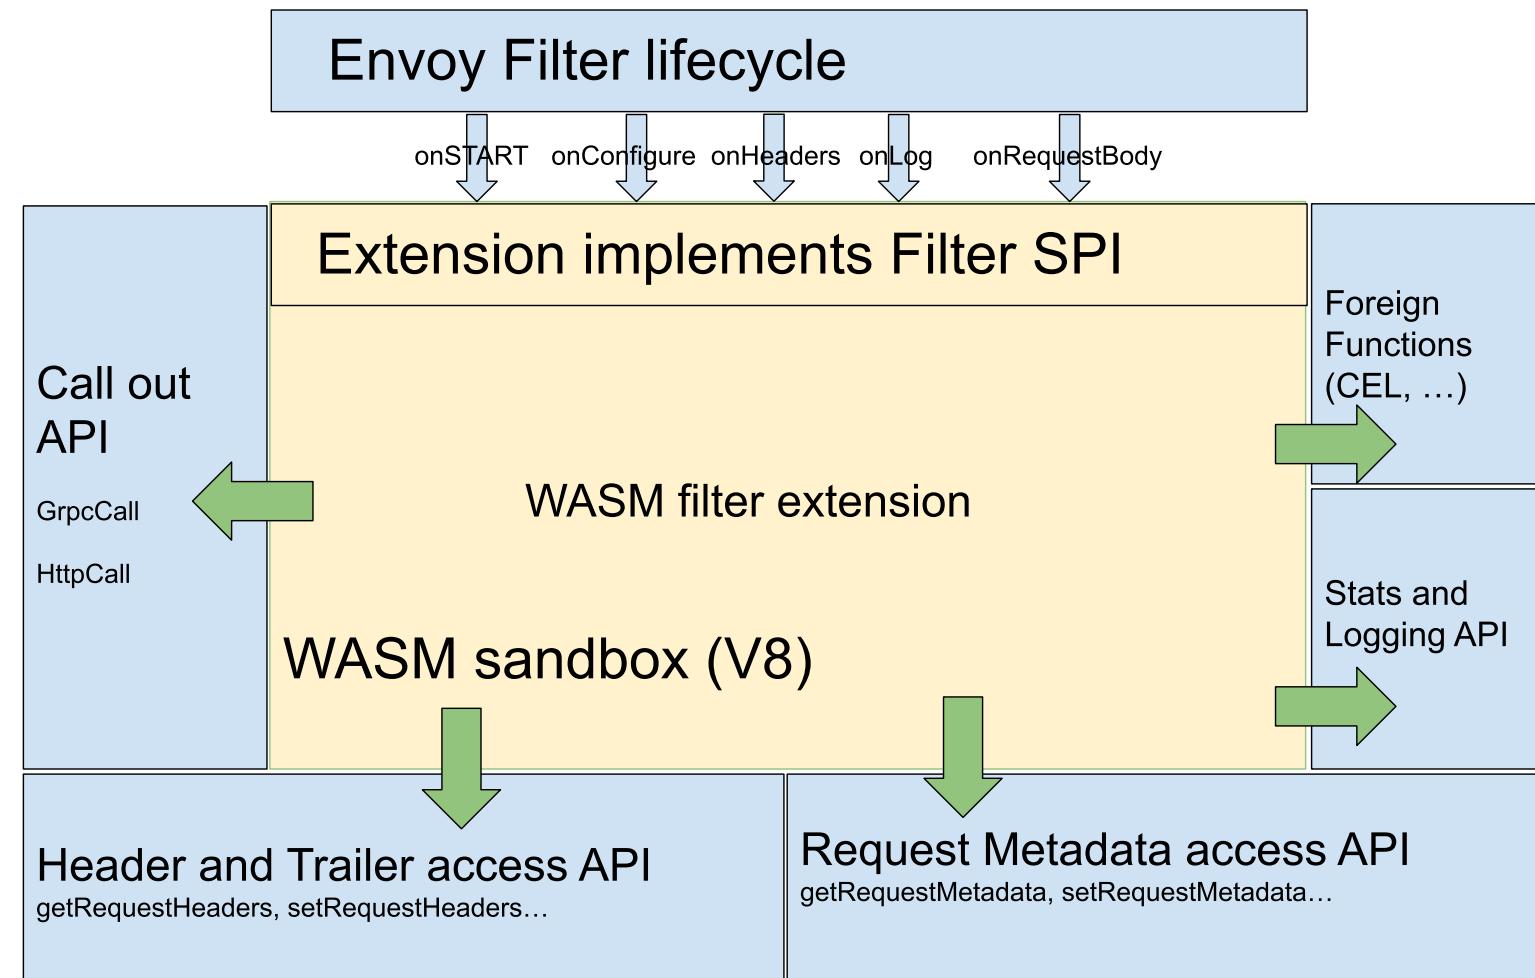
High-level Architecture

Istio extensions (Proxy-Wasm plugins) have several components:

- **Filter Service Provider Interface (SPI)** for building Proxy-Wasm plugins for filters.
- **Sandbox V8 Wasm Runtime** embedded in Envoy.
- **Host APIs** for headers, trailers and metadata.
- **Call out APIs** for gRPC and HTTP calls.
- **Stats and Logging APIs** for metrics and monitoring.



High-level Architecture





Fault Injection



Fault Injection

Fault injection is a form of chaos experimentation that introduces errors into a system to ensure that it can withstand and recover from error conditions.

This process can be particularly useful to ensure that your failure recovery policies aren't incompatible or too restrictive, potentially resulting in critical service unavailability.



Fault Injection

You can inject two types of faults, both configured using a `VirtualService`

Delays

Delays are timing failures. They mimic increased network latency or an overloaded upstream service.

Aborts

Aborts are crash failures. They mimic failures in upstream services. Aborts usually manifest in the form of HTTP error codes or TCP connection failures.



Fault Injection

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
    - ratings
  http:
    - fault:
        delay:
          percentage:
            value: 0.1
          fixedDelay: 5s
      route:
        - destination:
            host: ratings
            subset: v1
```

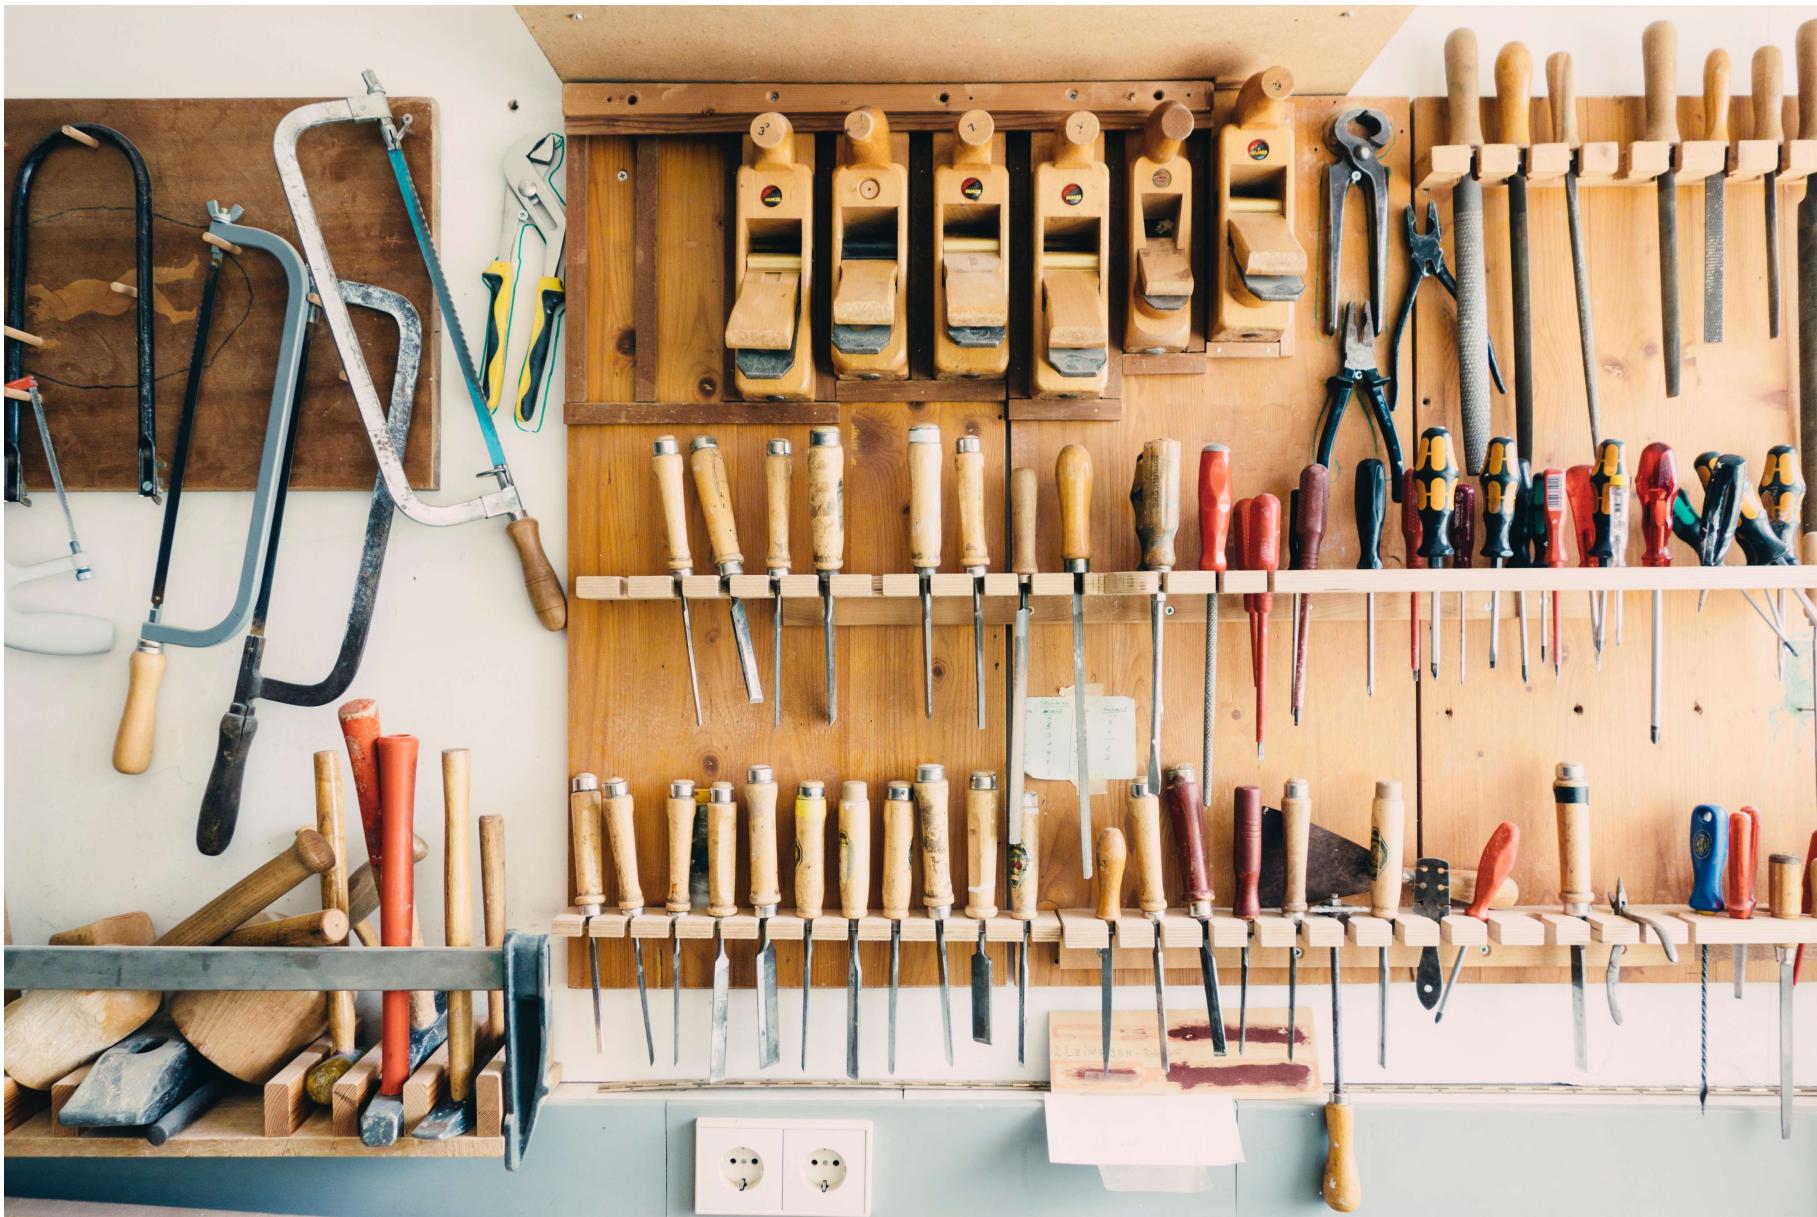


Fault Injection

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
    - ratings
  http:
    - fault:
        abort:
          httpStatus: 500
          percentage:
            value: 10
        route:
          - destination:
              host: ratings
              subset: v1
```



Workshop - Fault Injection





Circuit Breaking



Circuit Breaking

Circuit breakers are another useful mechanism Istio provides for creating resilient microservice-based applications.

A circuit breaker pattern ensures fast failure rather than having clients try to connect to an overloaded or failing host.



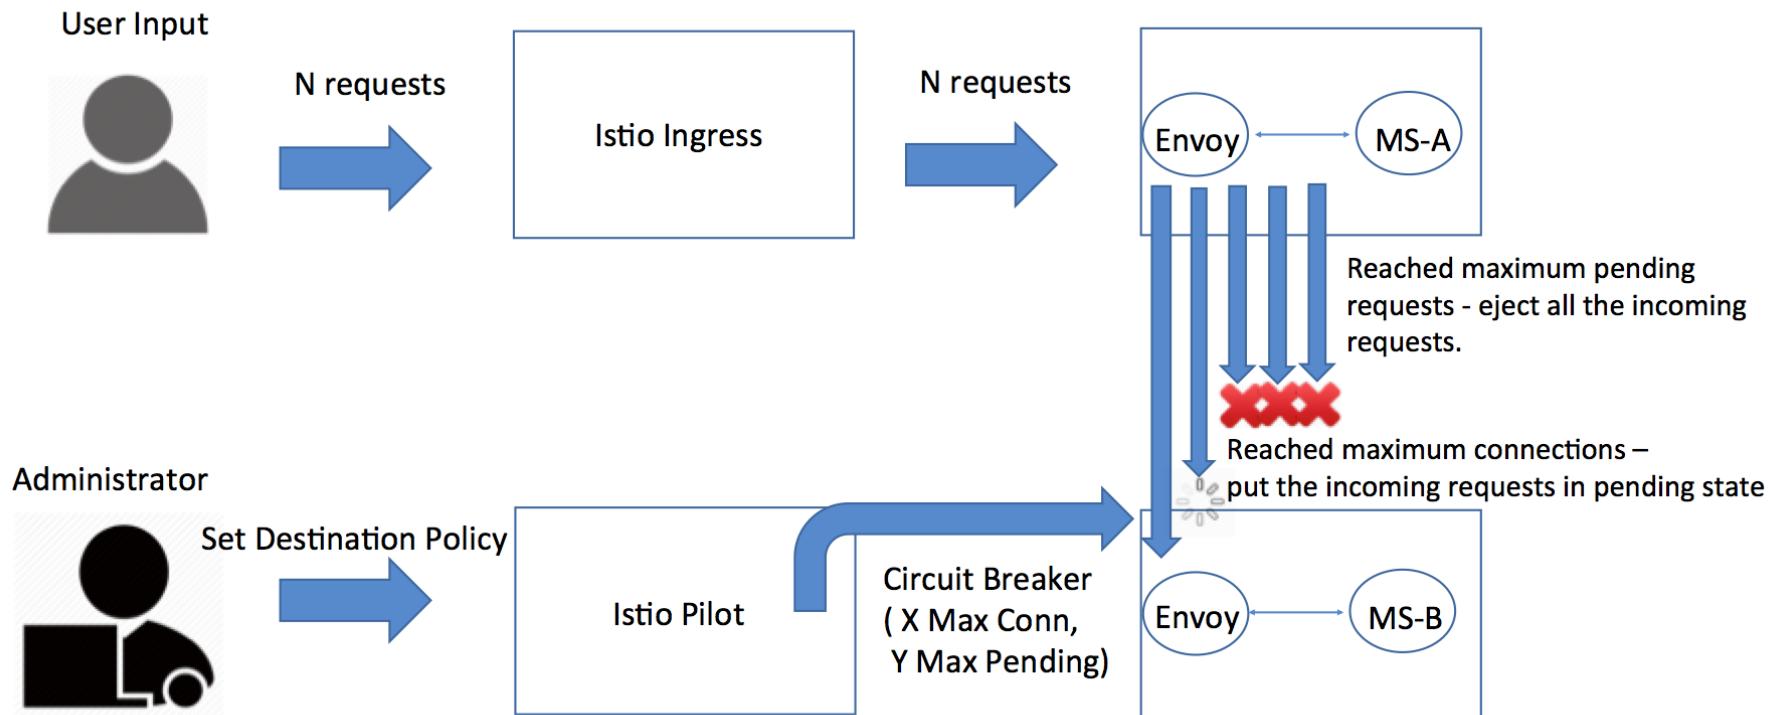
Circuit Breaking

The following example limits the number of concurrent connections for the reviews service workloads of the v1 subset to 100:

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: reviews
spec:
  host: reviews
  subsets:
  - name: v1
    labels:
      version: v1
    trafficPolicy:
      connectionPool:
        tcp:
          maxConnections: 100
```



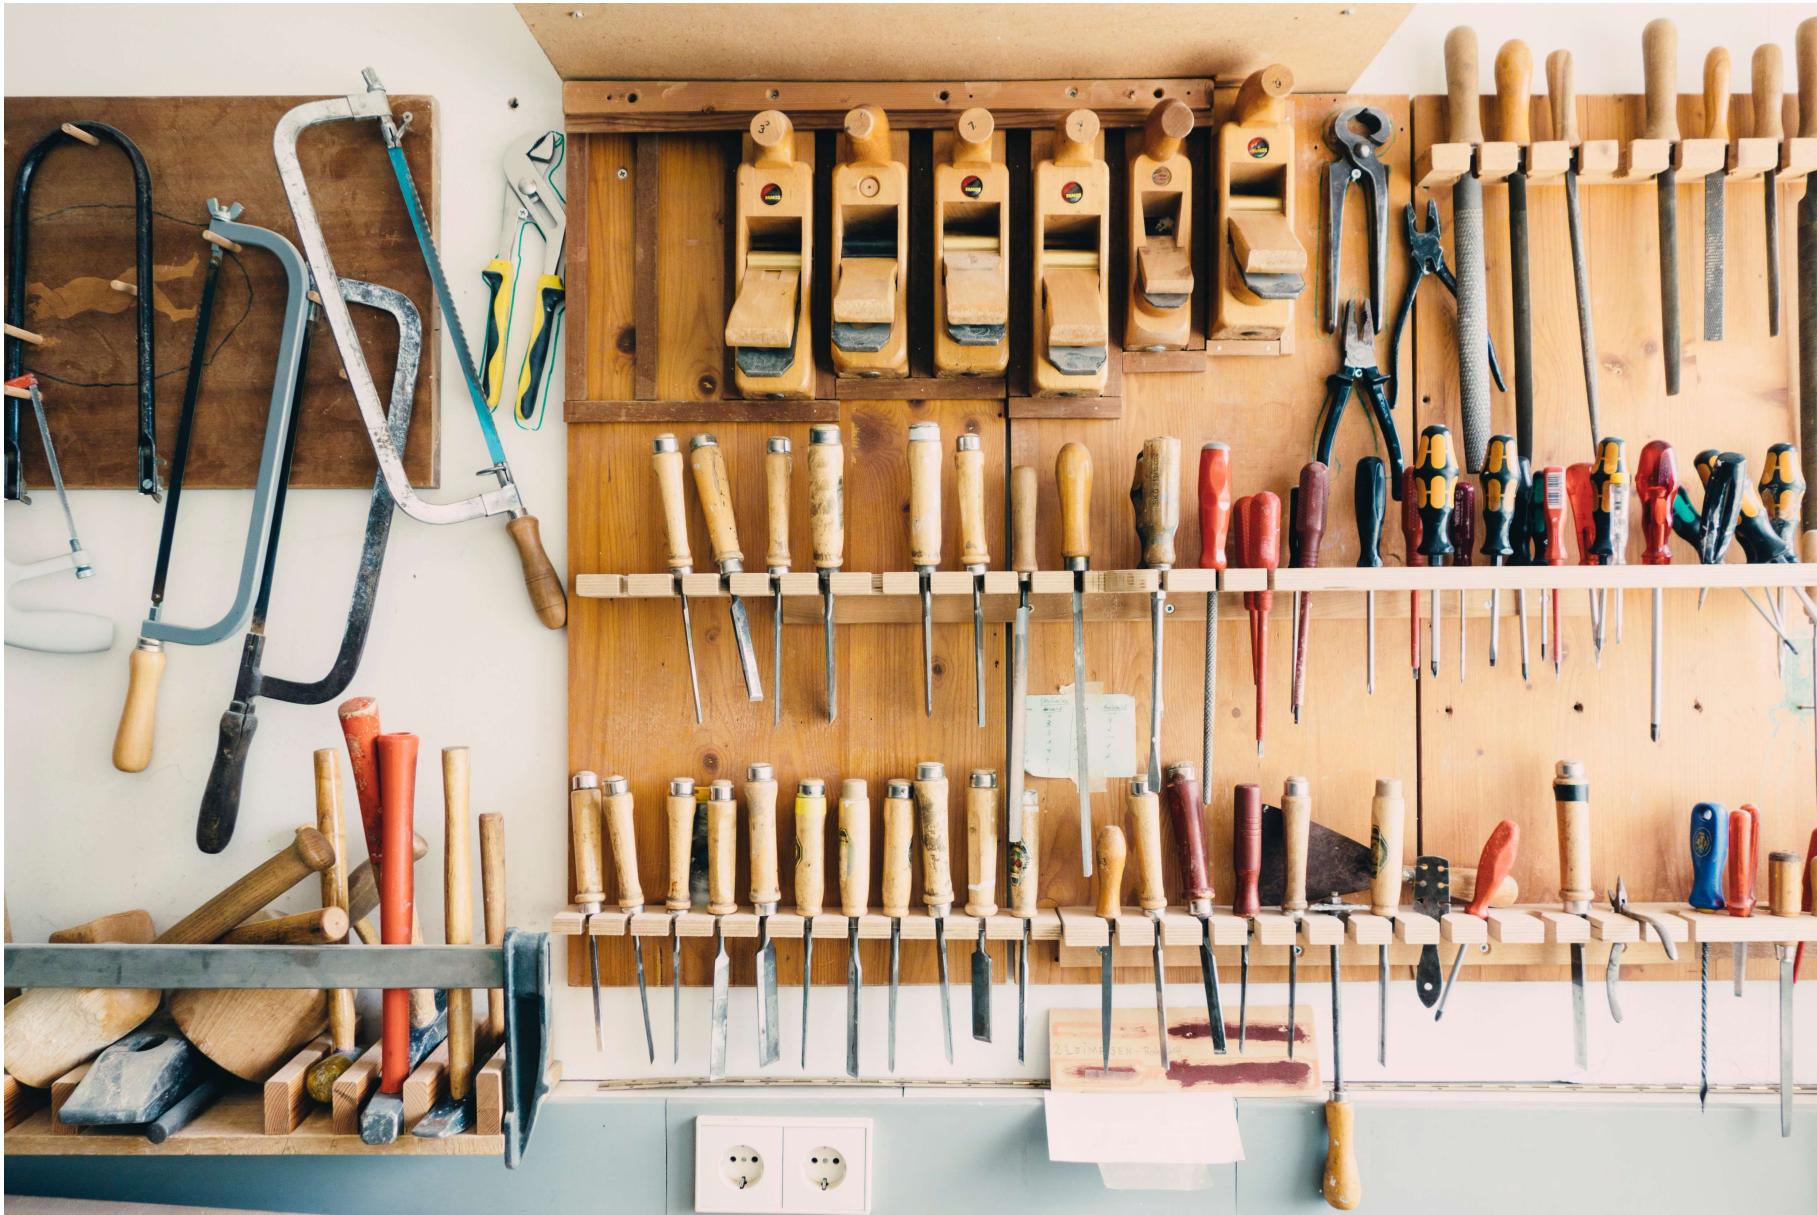
Circuit Breaking



<https://developer.ibm.com/code/wp-content/uploads/sites/118/2017/08/Screen-Shot-2017-08-22-at-12.50.46-AM-768x342.png>



Workshop - Circuit Breaking





Traffic Mirroring

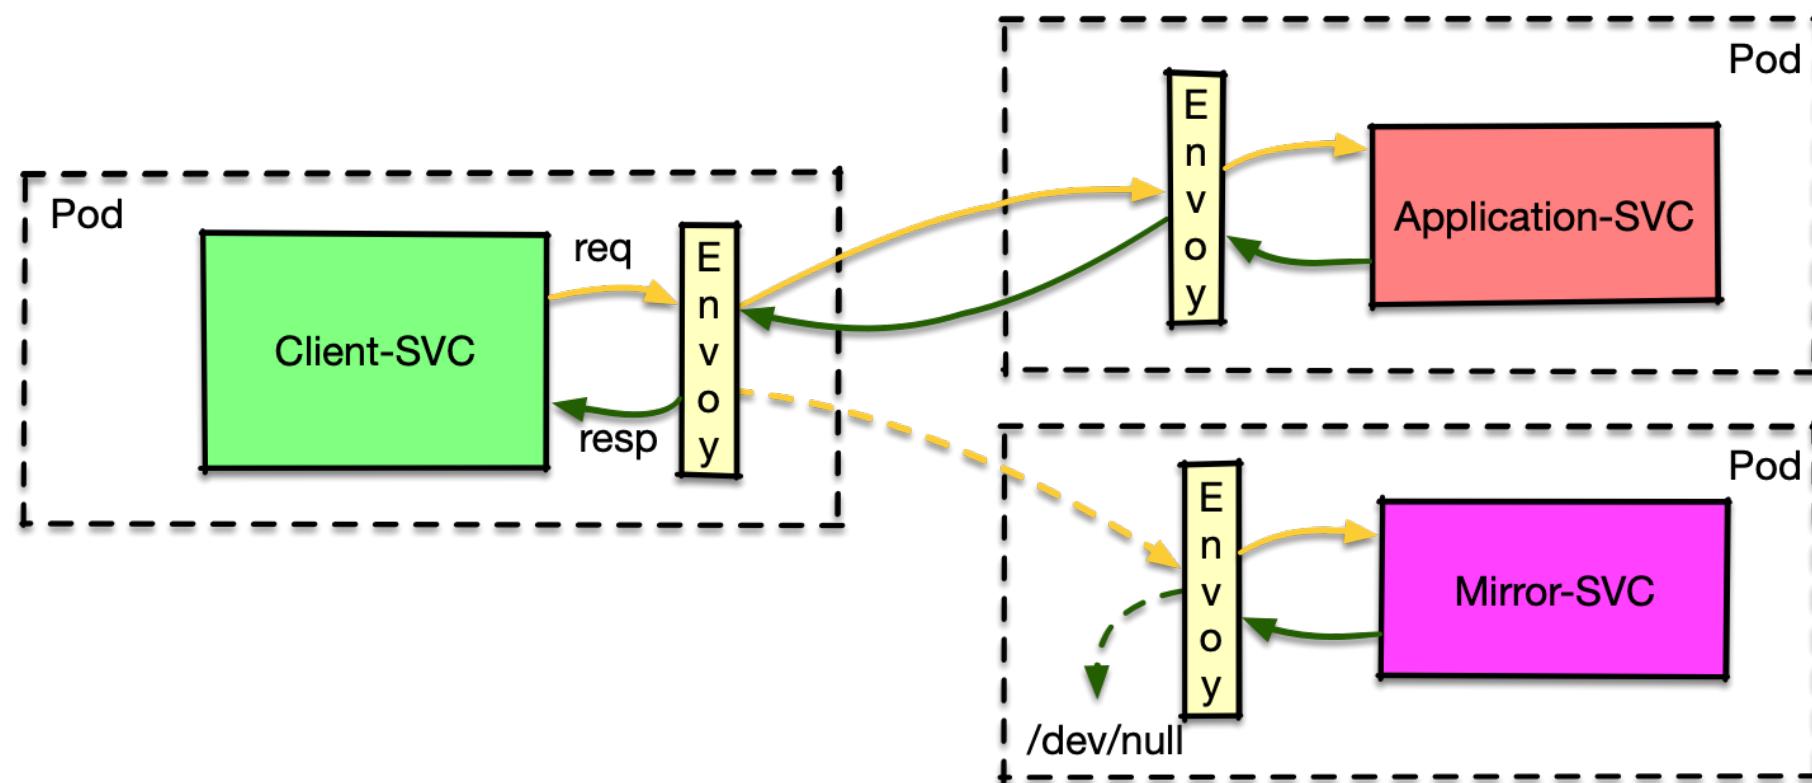


Traffic Mirroring

Mirroring sends a copy of live traffic to a mirrored service. The mirrored traffic happens out of band of the critical request path for the primary service.



Traffic Mirroring



https://miro.medium.com/max/1438/1*9_3WRLFmo3dcOvVMmdbkVw.png



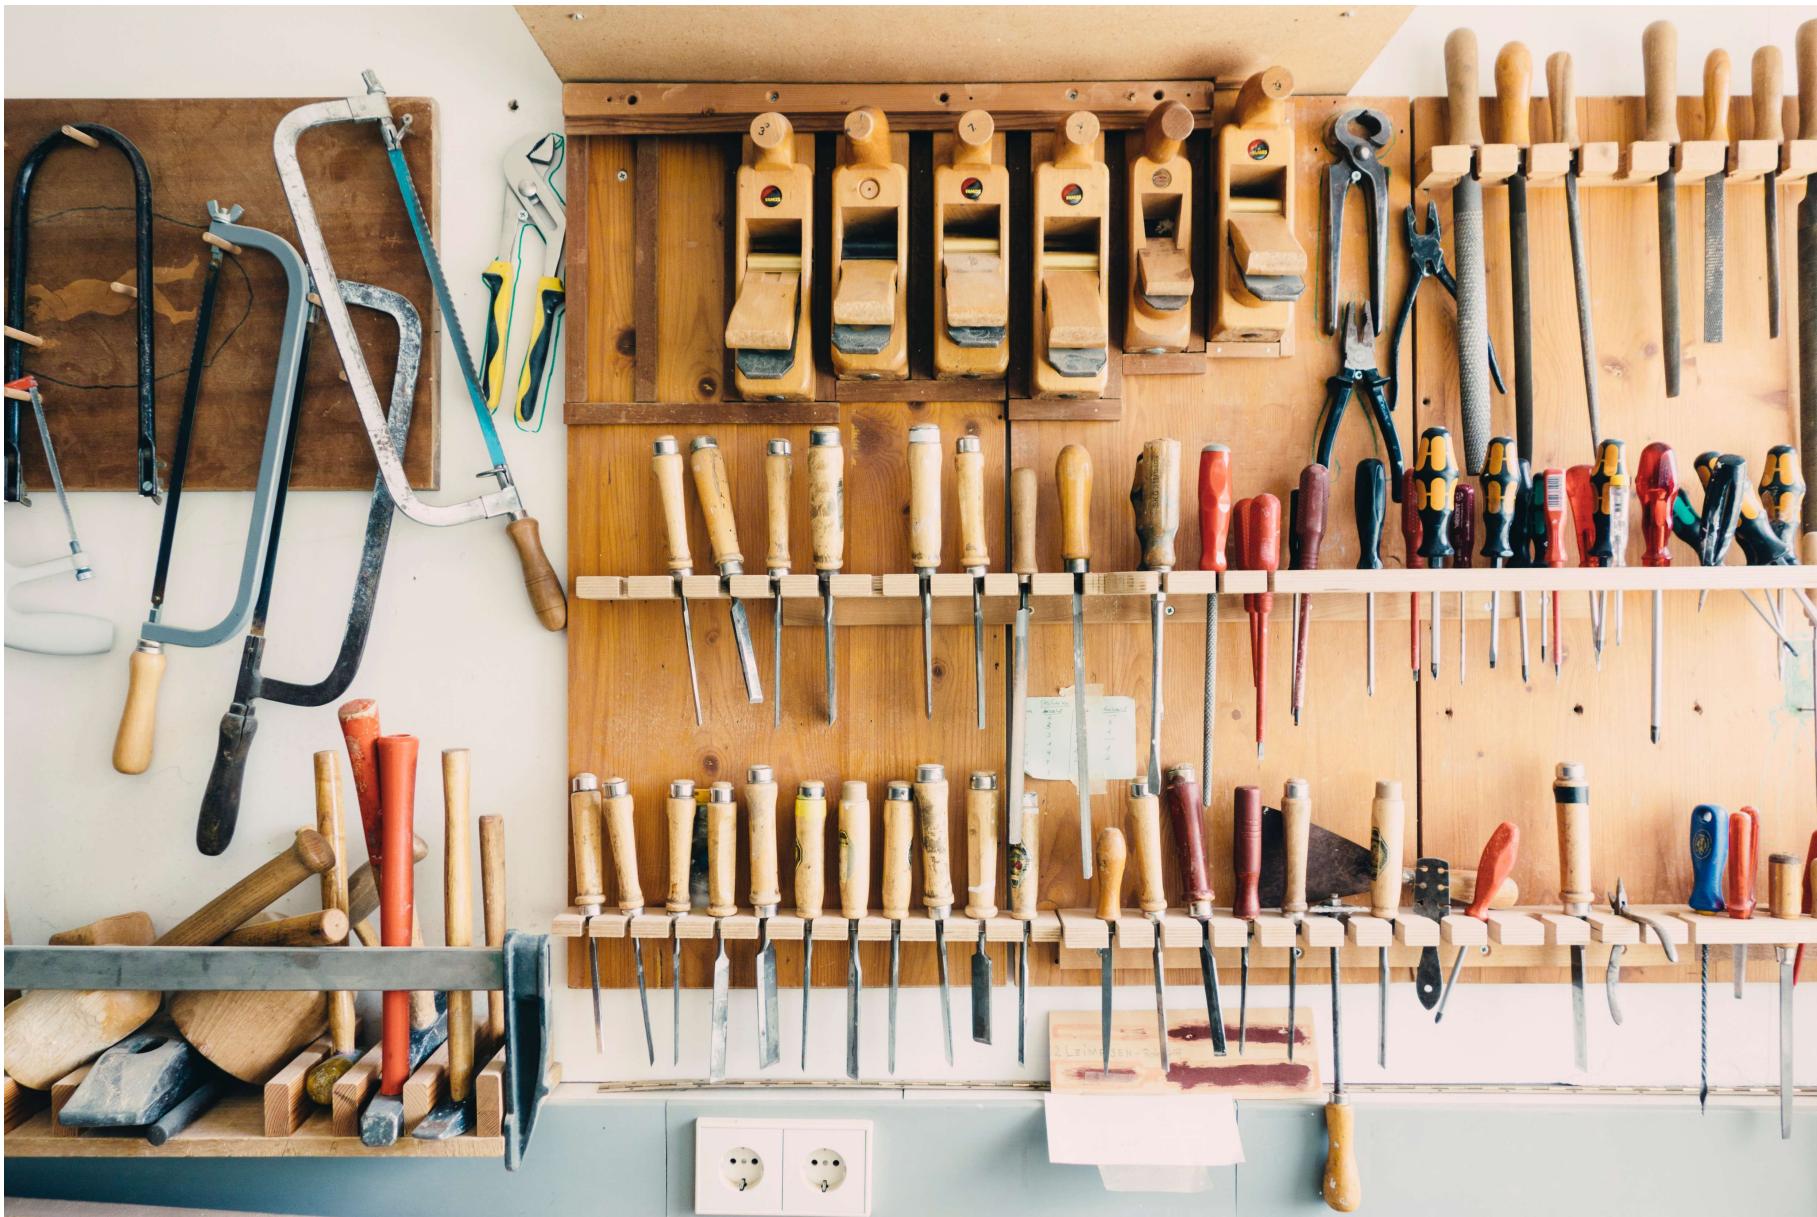
Traffic Mirroring

Mirroring traffic can be achieved using a VirtualService:

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: httpbin
spec:
  hosts:
  - httpbin
  http:
  - route:
    - destination:
        host: httpbin
        subset: v1
        weight: 100
    mirror:
        host: httpbin
        subset: v2
  mirror_percentage:
    value: 100
```



Workshop - Traffic Mirroring





Deployment Models

Replicated control planes, federated meshes and more...



Deployment Tools

- Helm (deprecated)
- istioctl
- Operator (experimental)
- Managed add-ons



istioctl Install

```
istioctl install --set profile=demo
```

	default	demo	minimal	external	empty	preview
Core components						
istio-egressgateway		✓				
istio-ingressgateway	✓	✓			✓	
istiod	✓	✓	✓			✓



IstioOperator Resource

An `IstioOperator` resource describes an Istio deployment.

```
apiVersion: install.istio.io/v1alpha1
kind: IstioOperator
metadata:
  name: example-istiocontrolplane
  namespace: istio-system
spec:
  profile: demo
```



IstioOperator Resource

Can be used to customise Istio components.

```
apiVersion: install.istio.io/v1alpha1
kind: IstioOperator
metadata:
  name: istio
spec:
  components:
    ingressGateways:
    - enabled: true
      name: istio-ingressgateway
    k8s:
      service:
        type: ClusterIP
      serviceAnnotations:
        # stand-alone NEG annotation
        cloud.google.com/neg: '{"exposed_ports":{"443":{},"80":{}}}'
```



istioctl Install With IstioOperator

```
istioctl install --filename istio_operator.yaml
```



istioctl Manifest

Show generated manifests.

```
istioctl manifest generate --set profile=demo > istio_manifests.yaml
```

```
istioctl manifest generate --filename istio_operator.yaml > istio_manife
```



Istio Operator

Manage Istio deployment with an operator.

```
istioctl operator init
```

```
kubectl apply -f istio_operator.yaml
```



GKE

```
gcloud beta container clusters create ...  
  --addons=Istio \  
  --istio-config=auth=MTLS_PERMISSIVE
```

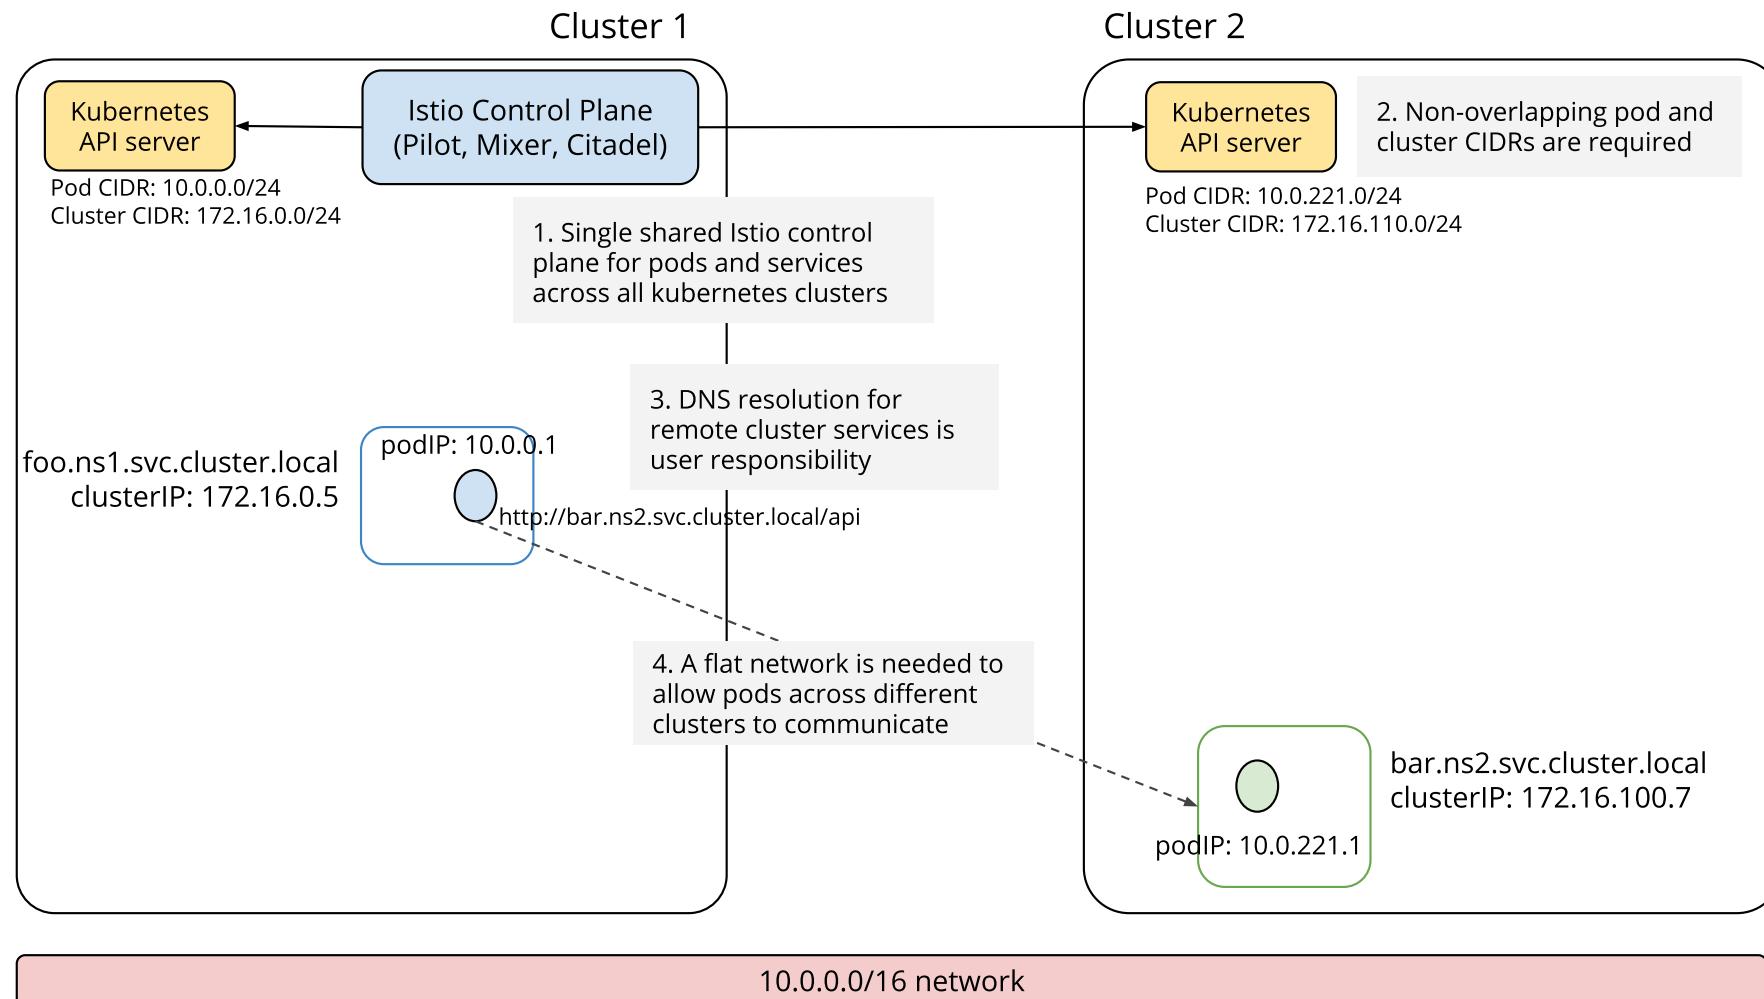


Deployment Models

- Single or multiple cluster
- Single or multiple network
- Single or multiple control plane
- Single or multiple mesh



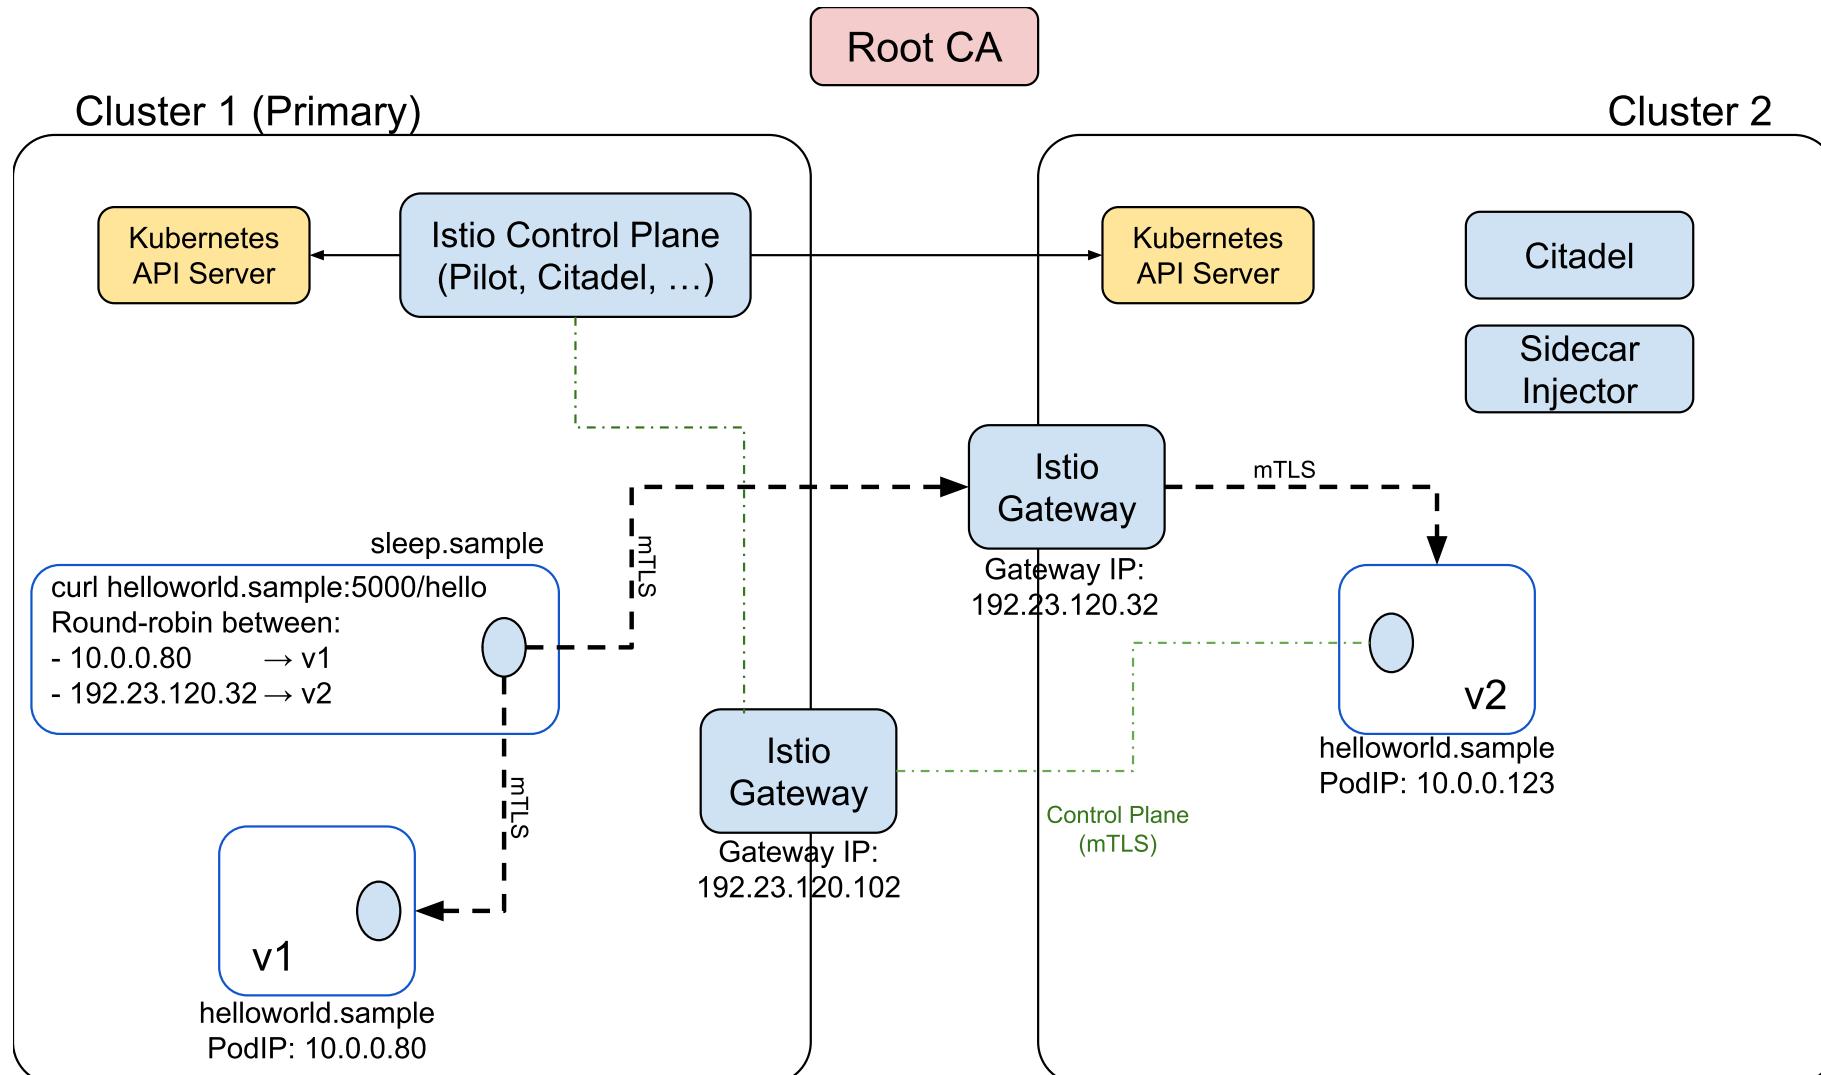
Multicloud, shared control plane, single network



<https://istio.io/docs/setup/install/multicloud/shared/>

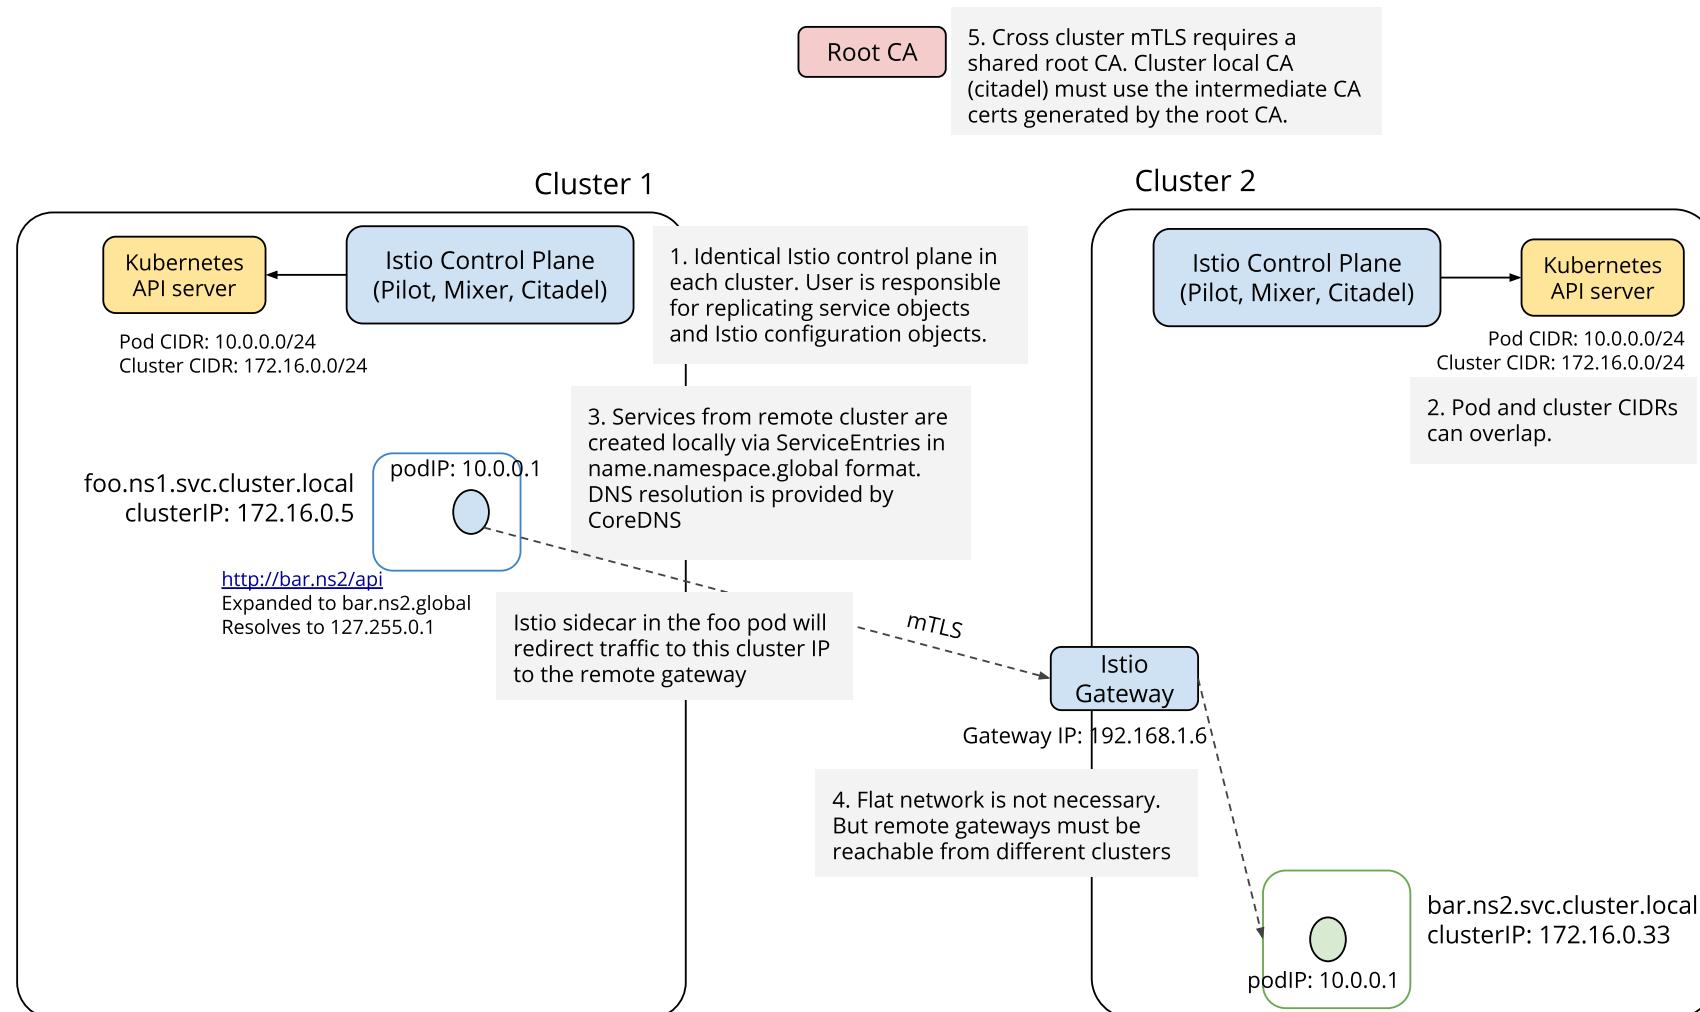


Multicloud, shared control plane, multiple networks





Multicloud, replicated control planes, multiple networks

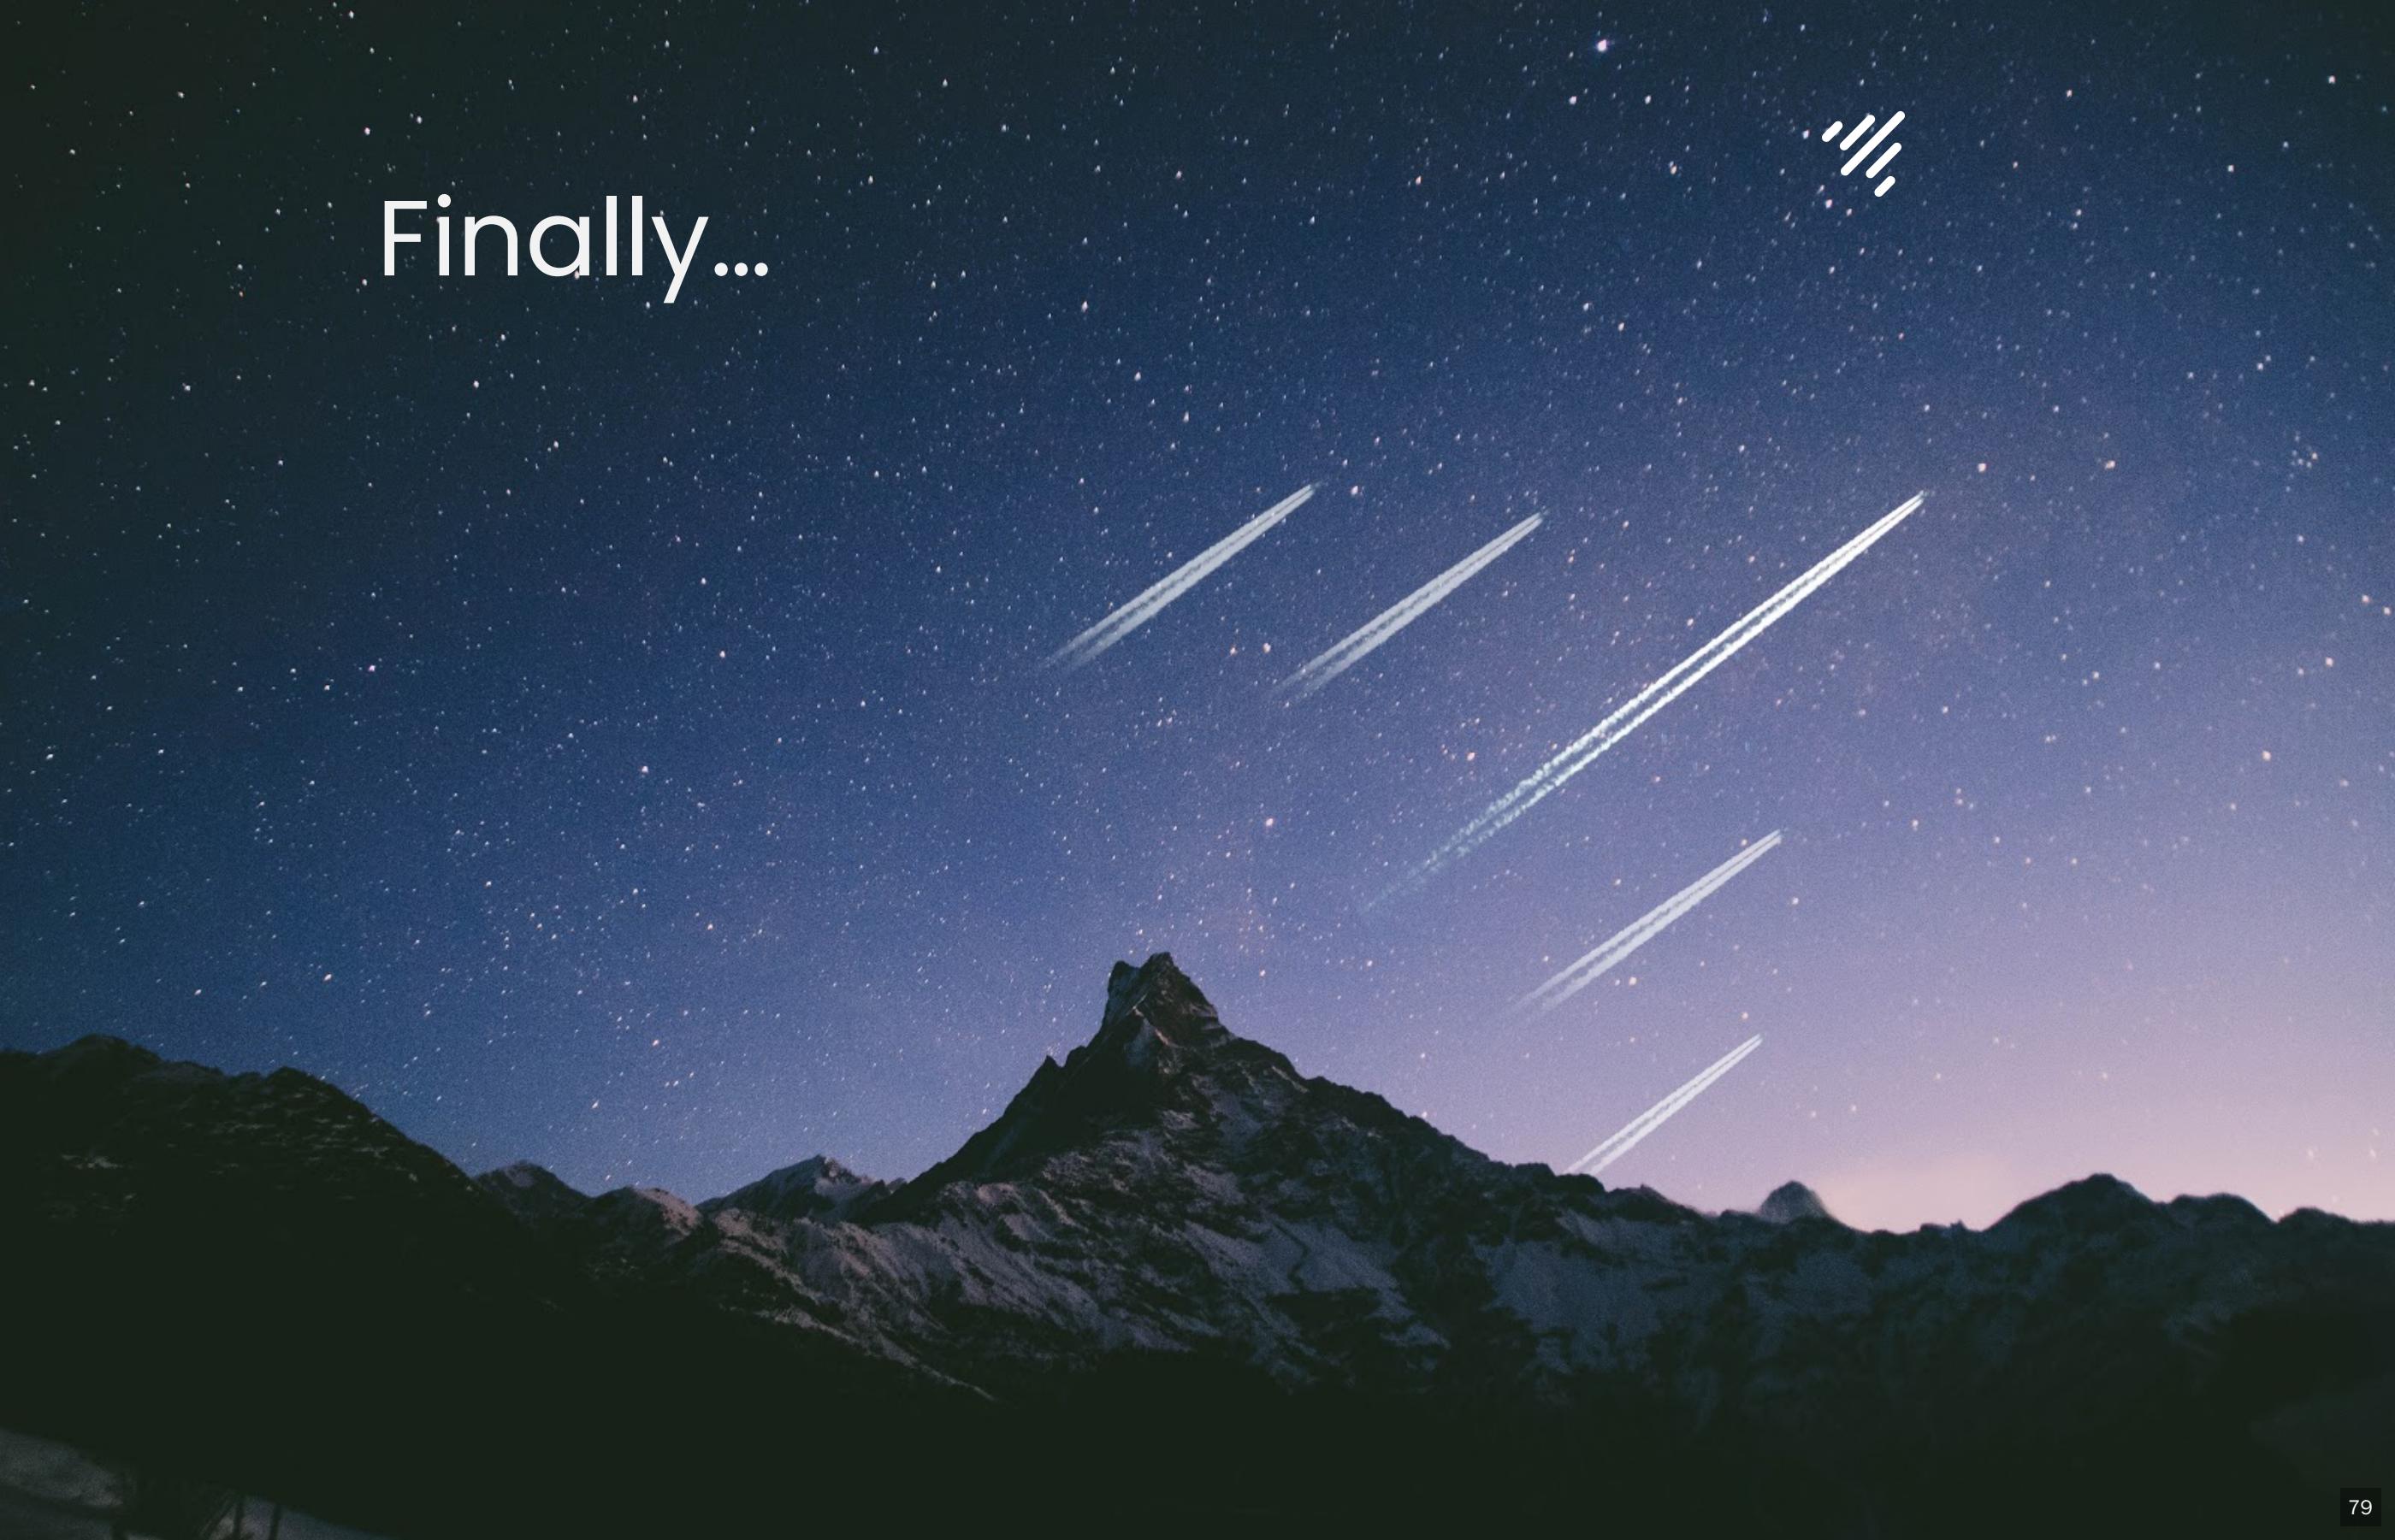




Shared vs Multiple Control planes

	Shared Control plane	Multiple Control planes
	<ul style="list-style-type: none">- Easy configuration- Centralized control	<ul style="list-style-type: none">- Distributed / decentralized controlplane- No SPoF- Scalable
	<ul style="list-style-type: none">- Single point of failure- Latency considerations- Scalability considerations	<ul style="list-style-type: none">- Harder configuration- Larger footprint

Finally...





Training

- Kubernetes in Practice (Day One and Day Two)
- Kubernetes for Application Developers (Day One and Day Two)
- Istio in Practice (Day One and Day Two)
- Operational Wargaming
- Kubernetes from Scratch
- Extending Kubernetes with Operators
- Serverless on Kubernetes with Knative



Consulting

On-site / Remote / Project-based

Production Readiness Review

Cloud Native Discovery Workshop



Jetstack Subscription

Jetstack Subscription is designed to help you verify, support and continually optimise your Kubernetes production environment.

-  Break Fix Support →
-  Customer Reliability Engineering →
-  Kubernetes Reference Architecture →
-  Automated Cluster Compliance Checking →
-  Online Training and Wargaming →

[Register Interest](#)



Jetstack is trusted by





The End

